

# **EFFECTIVE PLATFORM BUILDING WITH KUBERNETES.**

Wojciech Barczyński - [SMACC.io](#) | [Hypatos.ai](#)  
16 November 2018

# WOJCIECH BARCZYŃSKI

- Lead Software Developer & System Engineer
- K8S:  
2.5 years @ startups
- Before:  
Openstack, SAP R&D
- + Visiting lecturer and Trainer



Github: [wojciech12](#) | Linked: [IN](#) | HP: [wbarczynski.pl](#) | T: [@wbarczynski](#)

## STORY

- Lyke - [12.2016 - 07.2017]
- SMACC - [10.2017 - present]

My Role: Leading the change, implementing it

# WHY WE LIKE KUBERNETES?



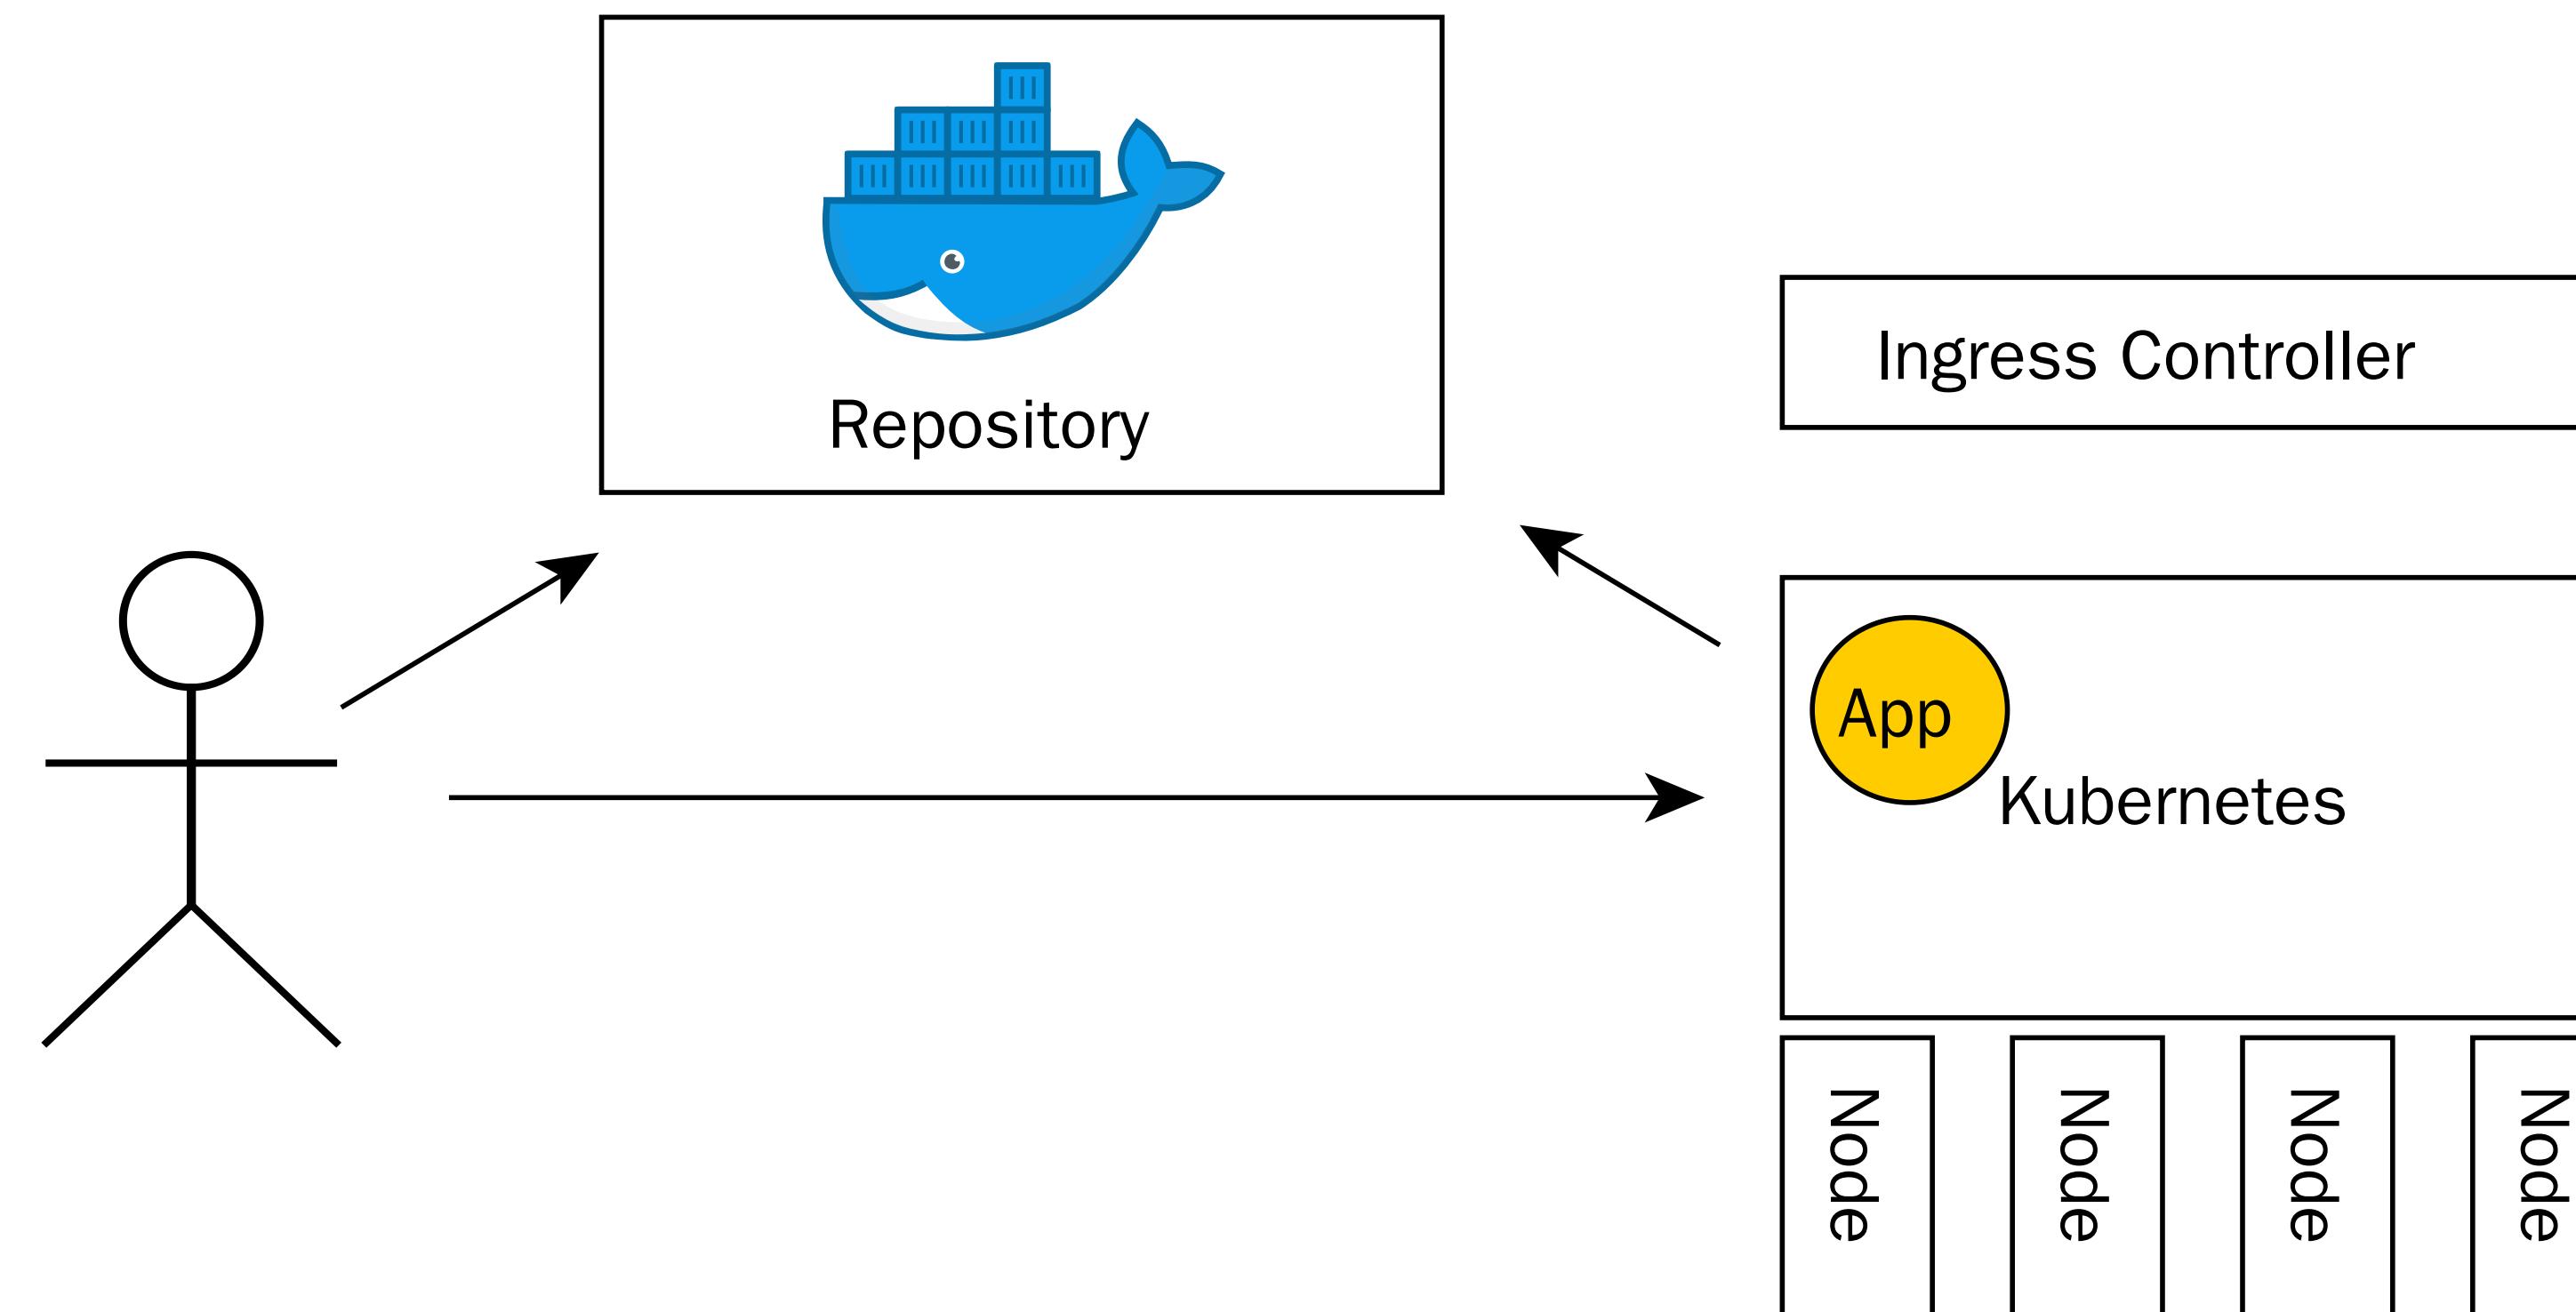
## FOR ME

- Data center as a black box
- Lingua Franca
- Batteries included
- Learn-as-you-go experience
- Independent from IaaS provider

Notice: not a silver bullet

# BLACK BOX

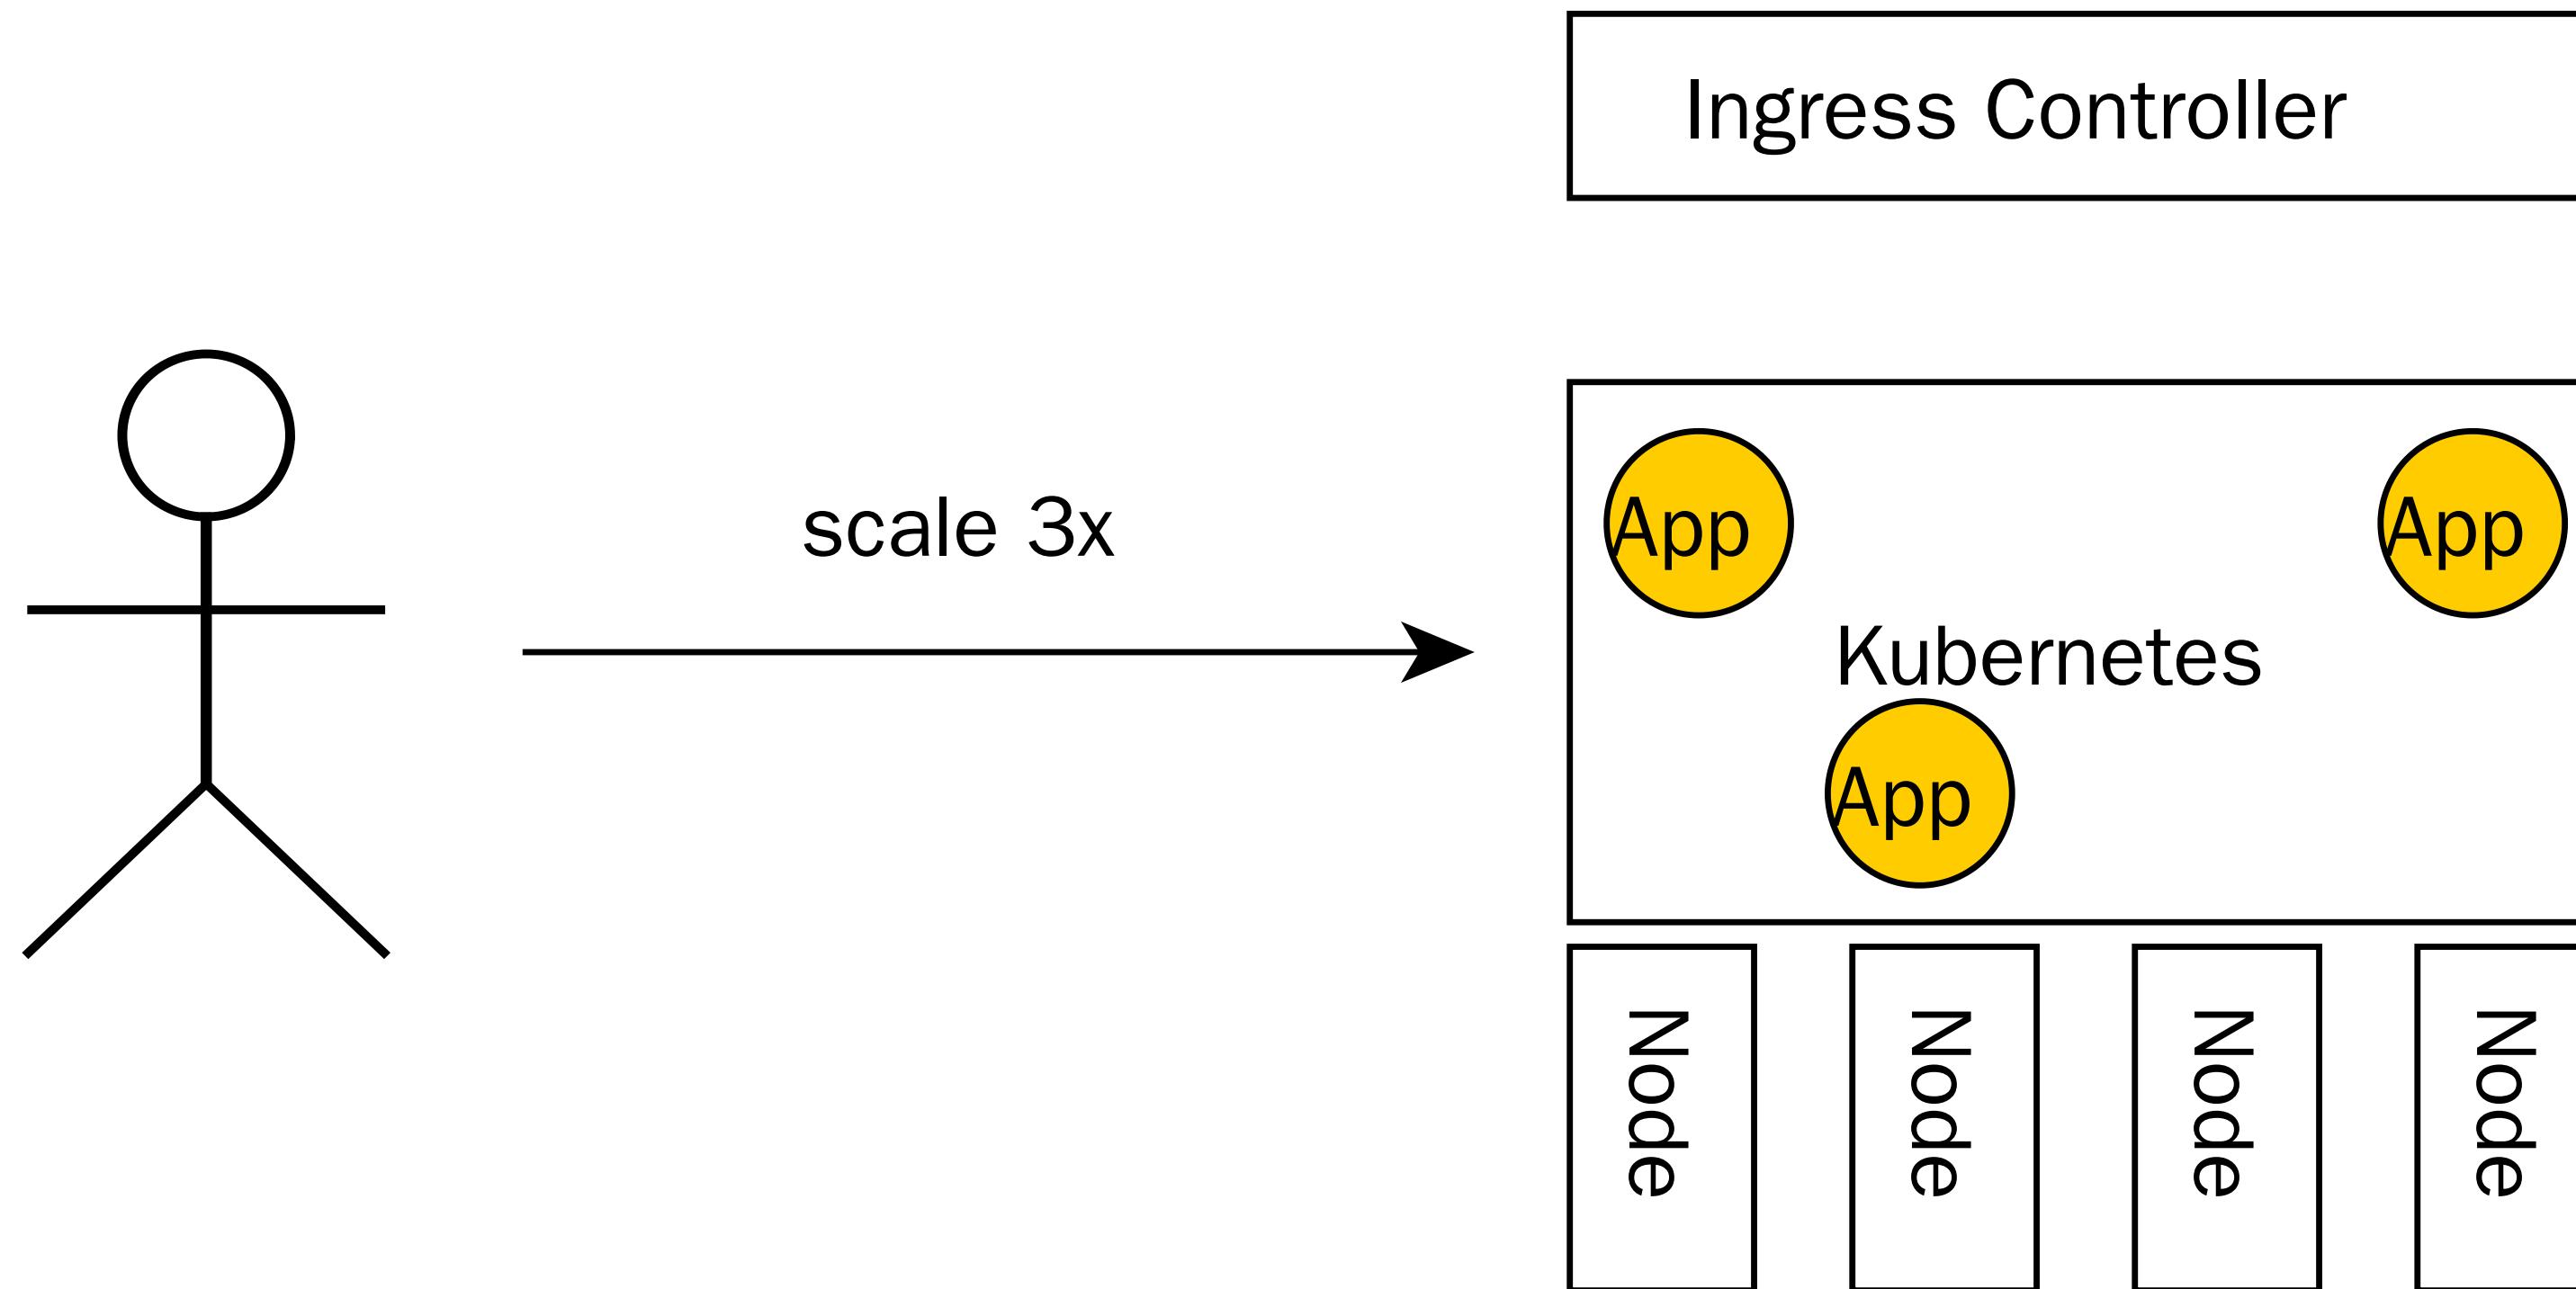
Deploy!



```
make docker_push; kubectl create -f app-srv-dpl.yaml
```

# BLACK BOX

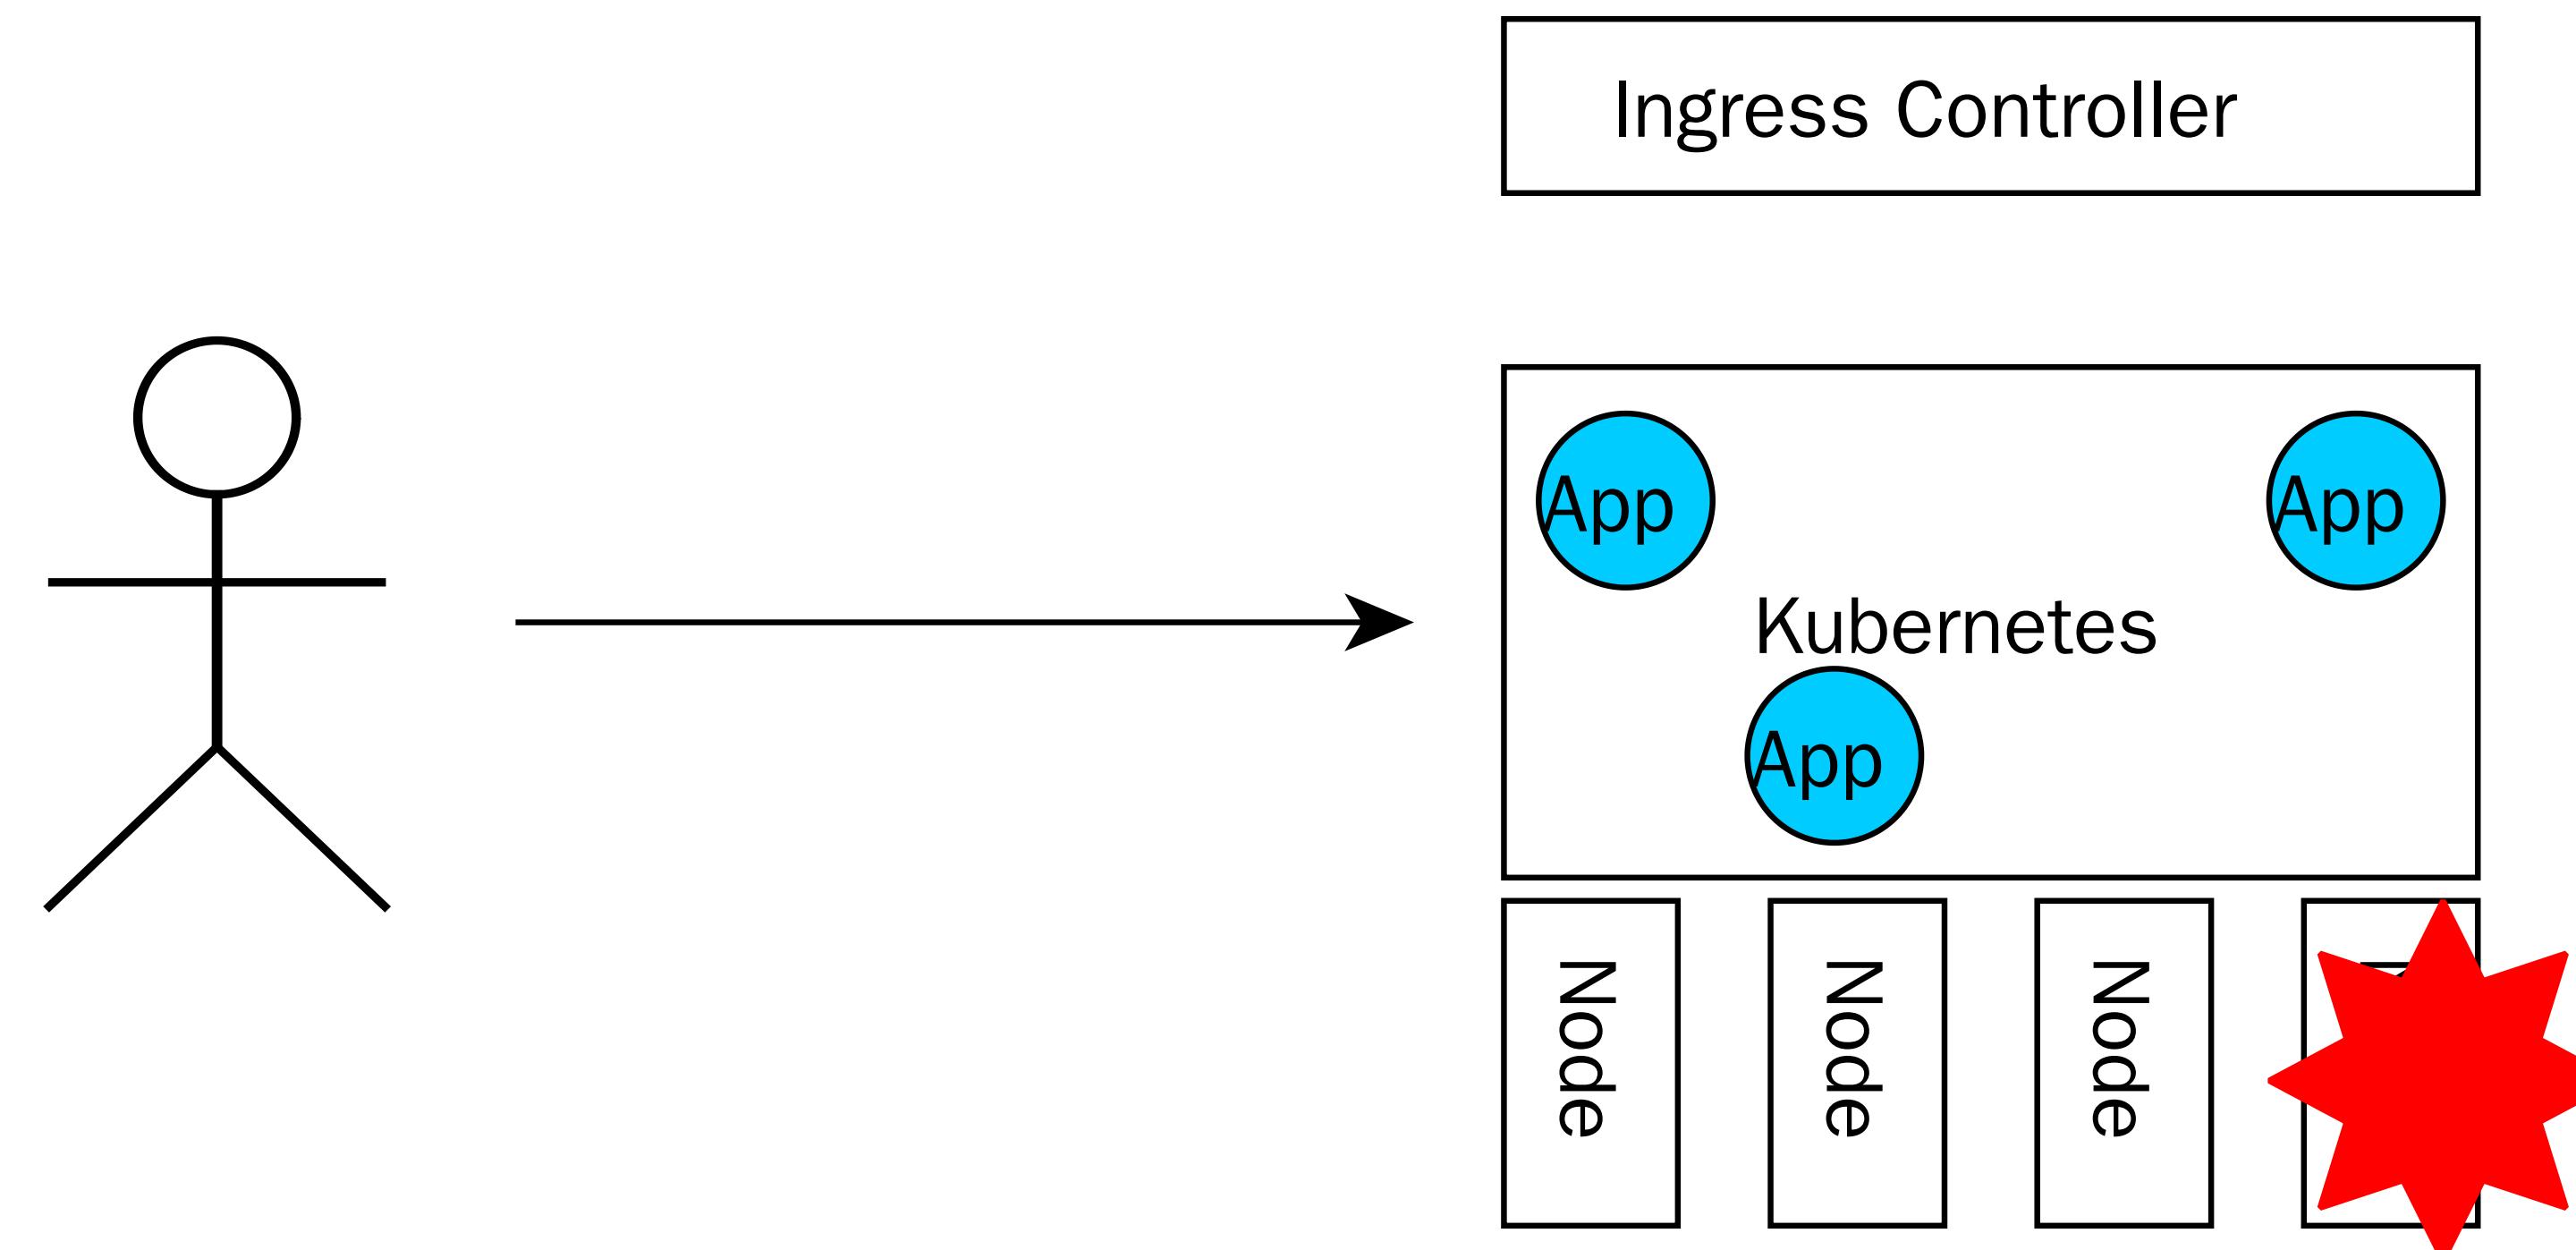
Scale up! Scale down!



`kubectl --replicas=3 -f app-srv-dpl.yaml`

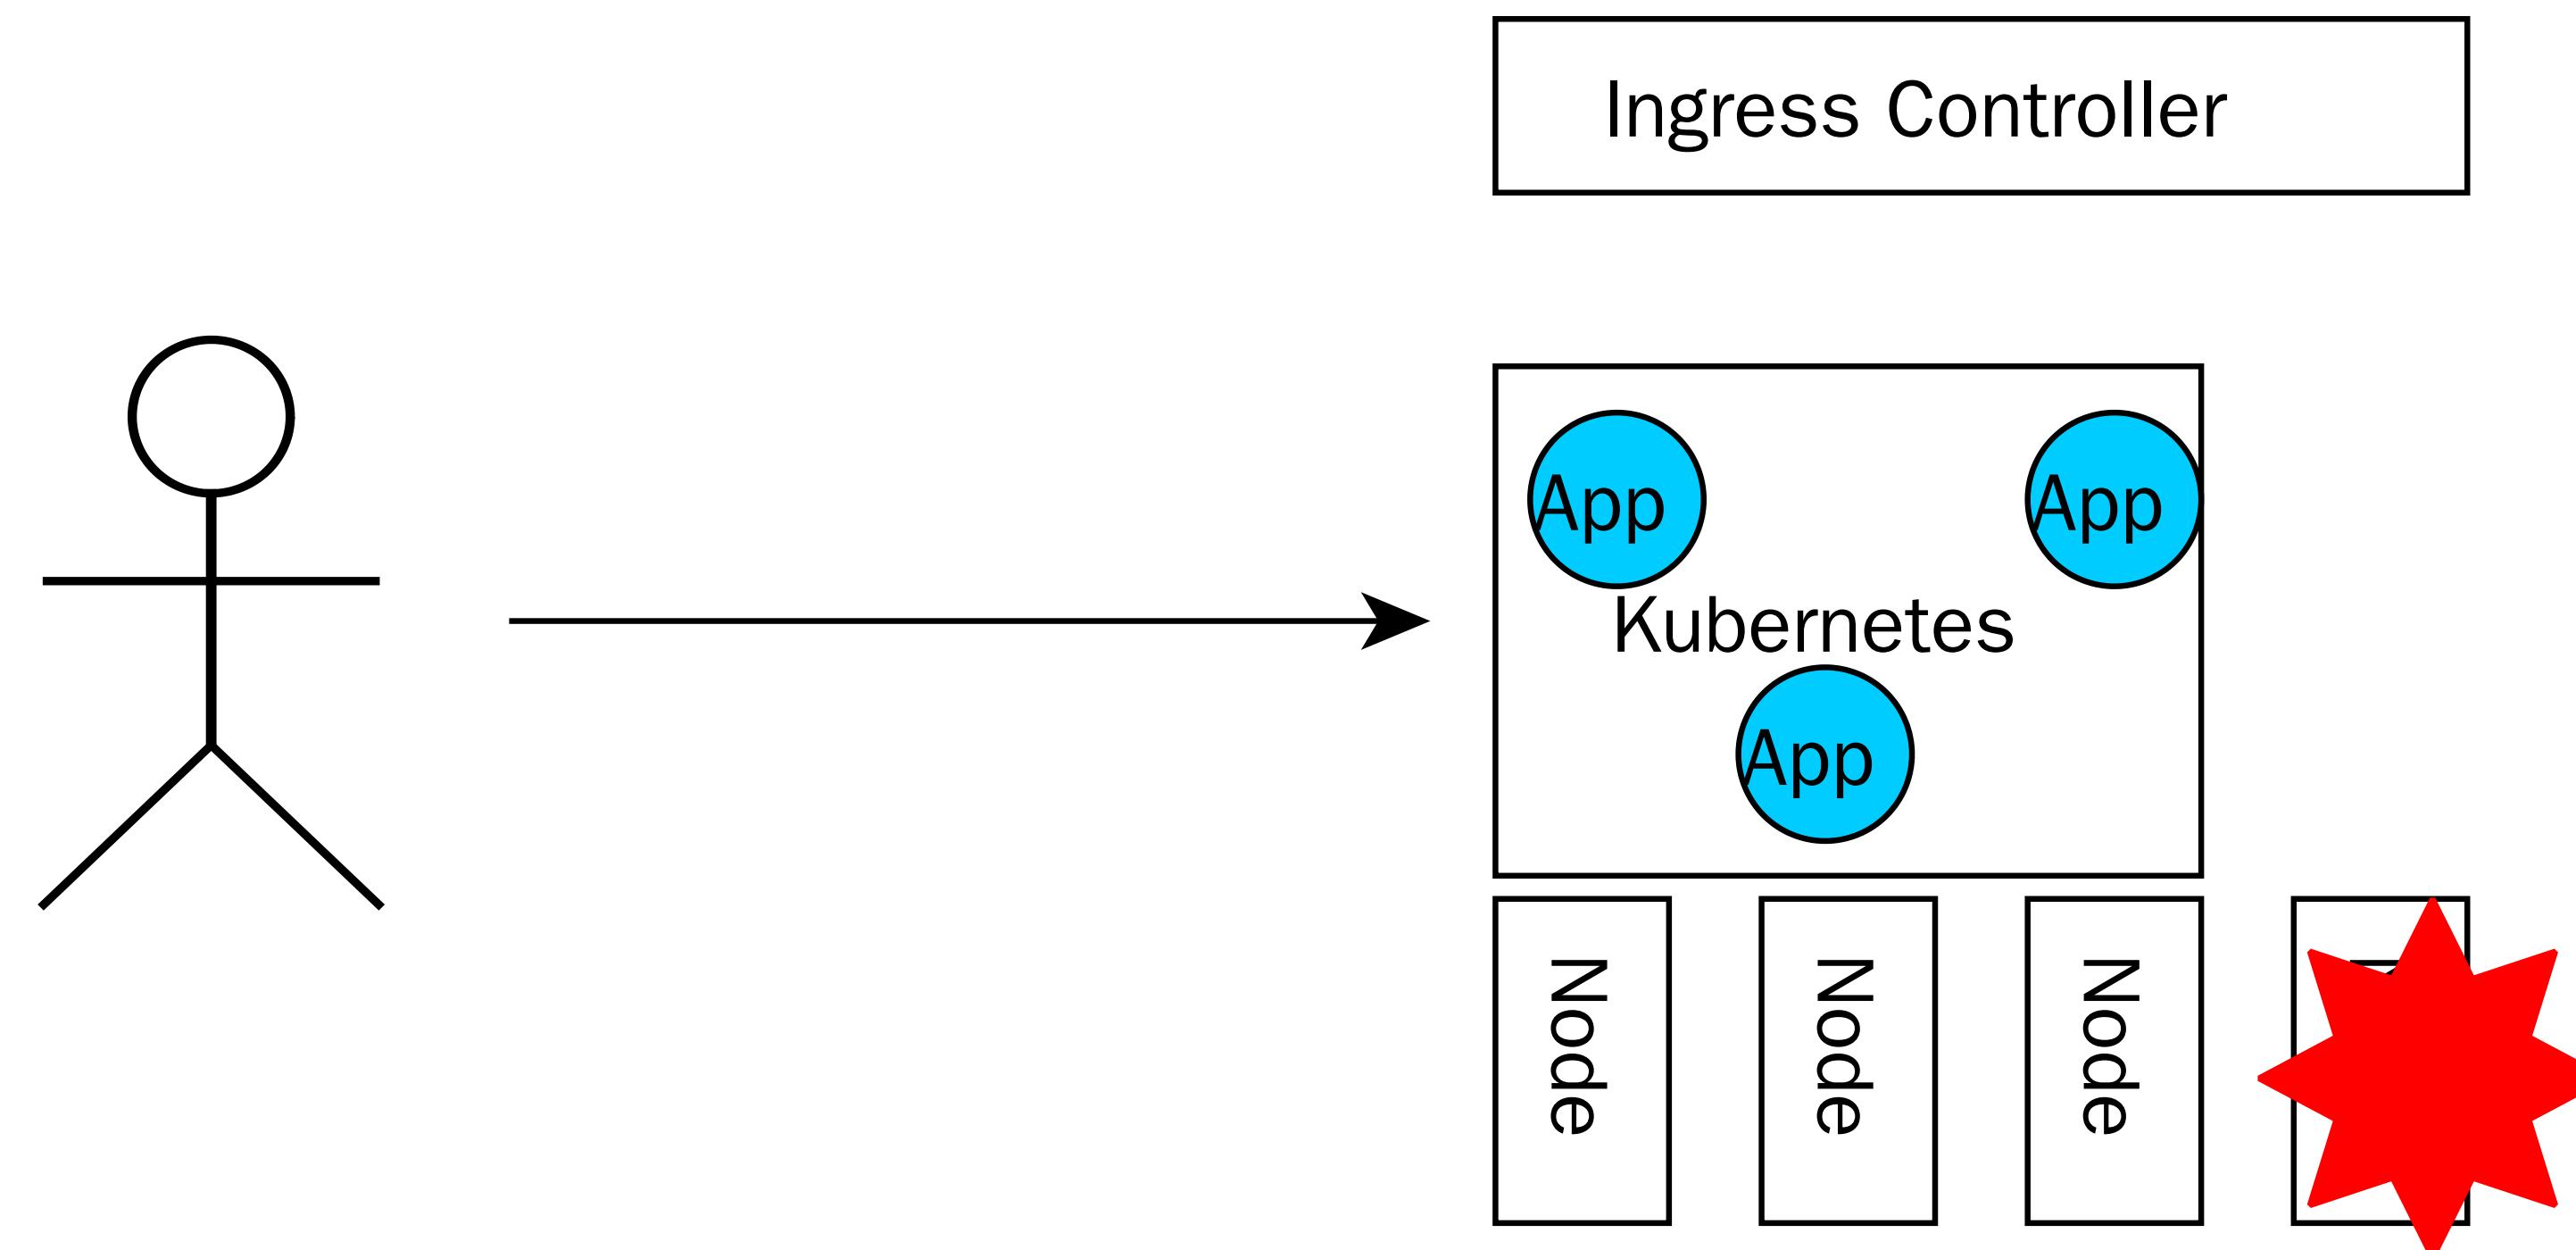
# BLACK BOX

## Resistance and Migrations



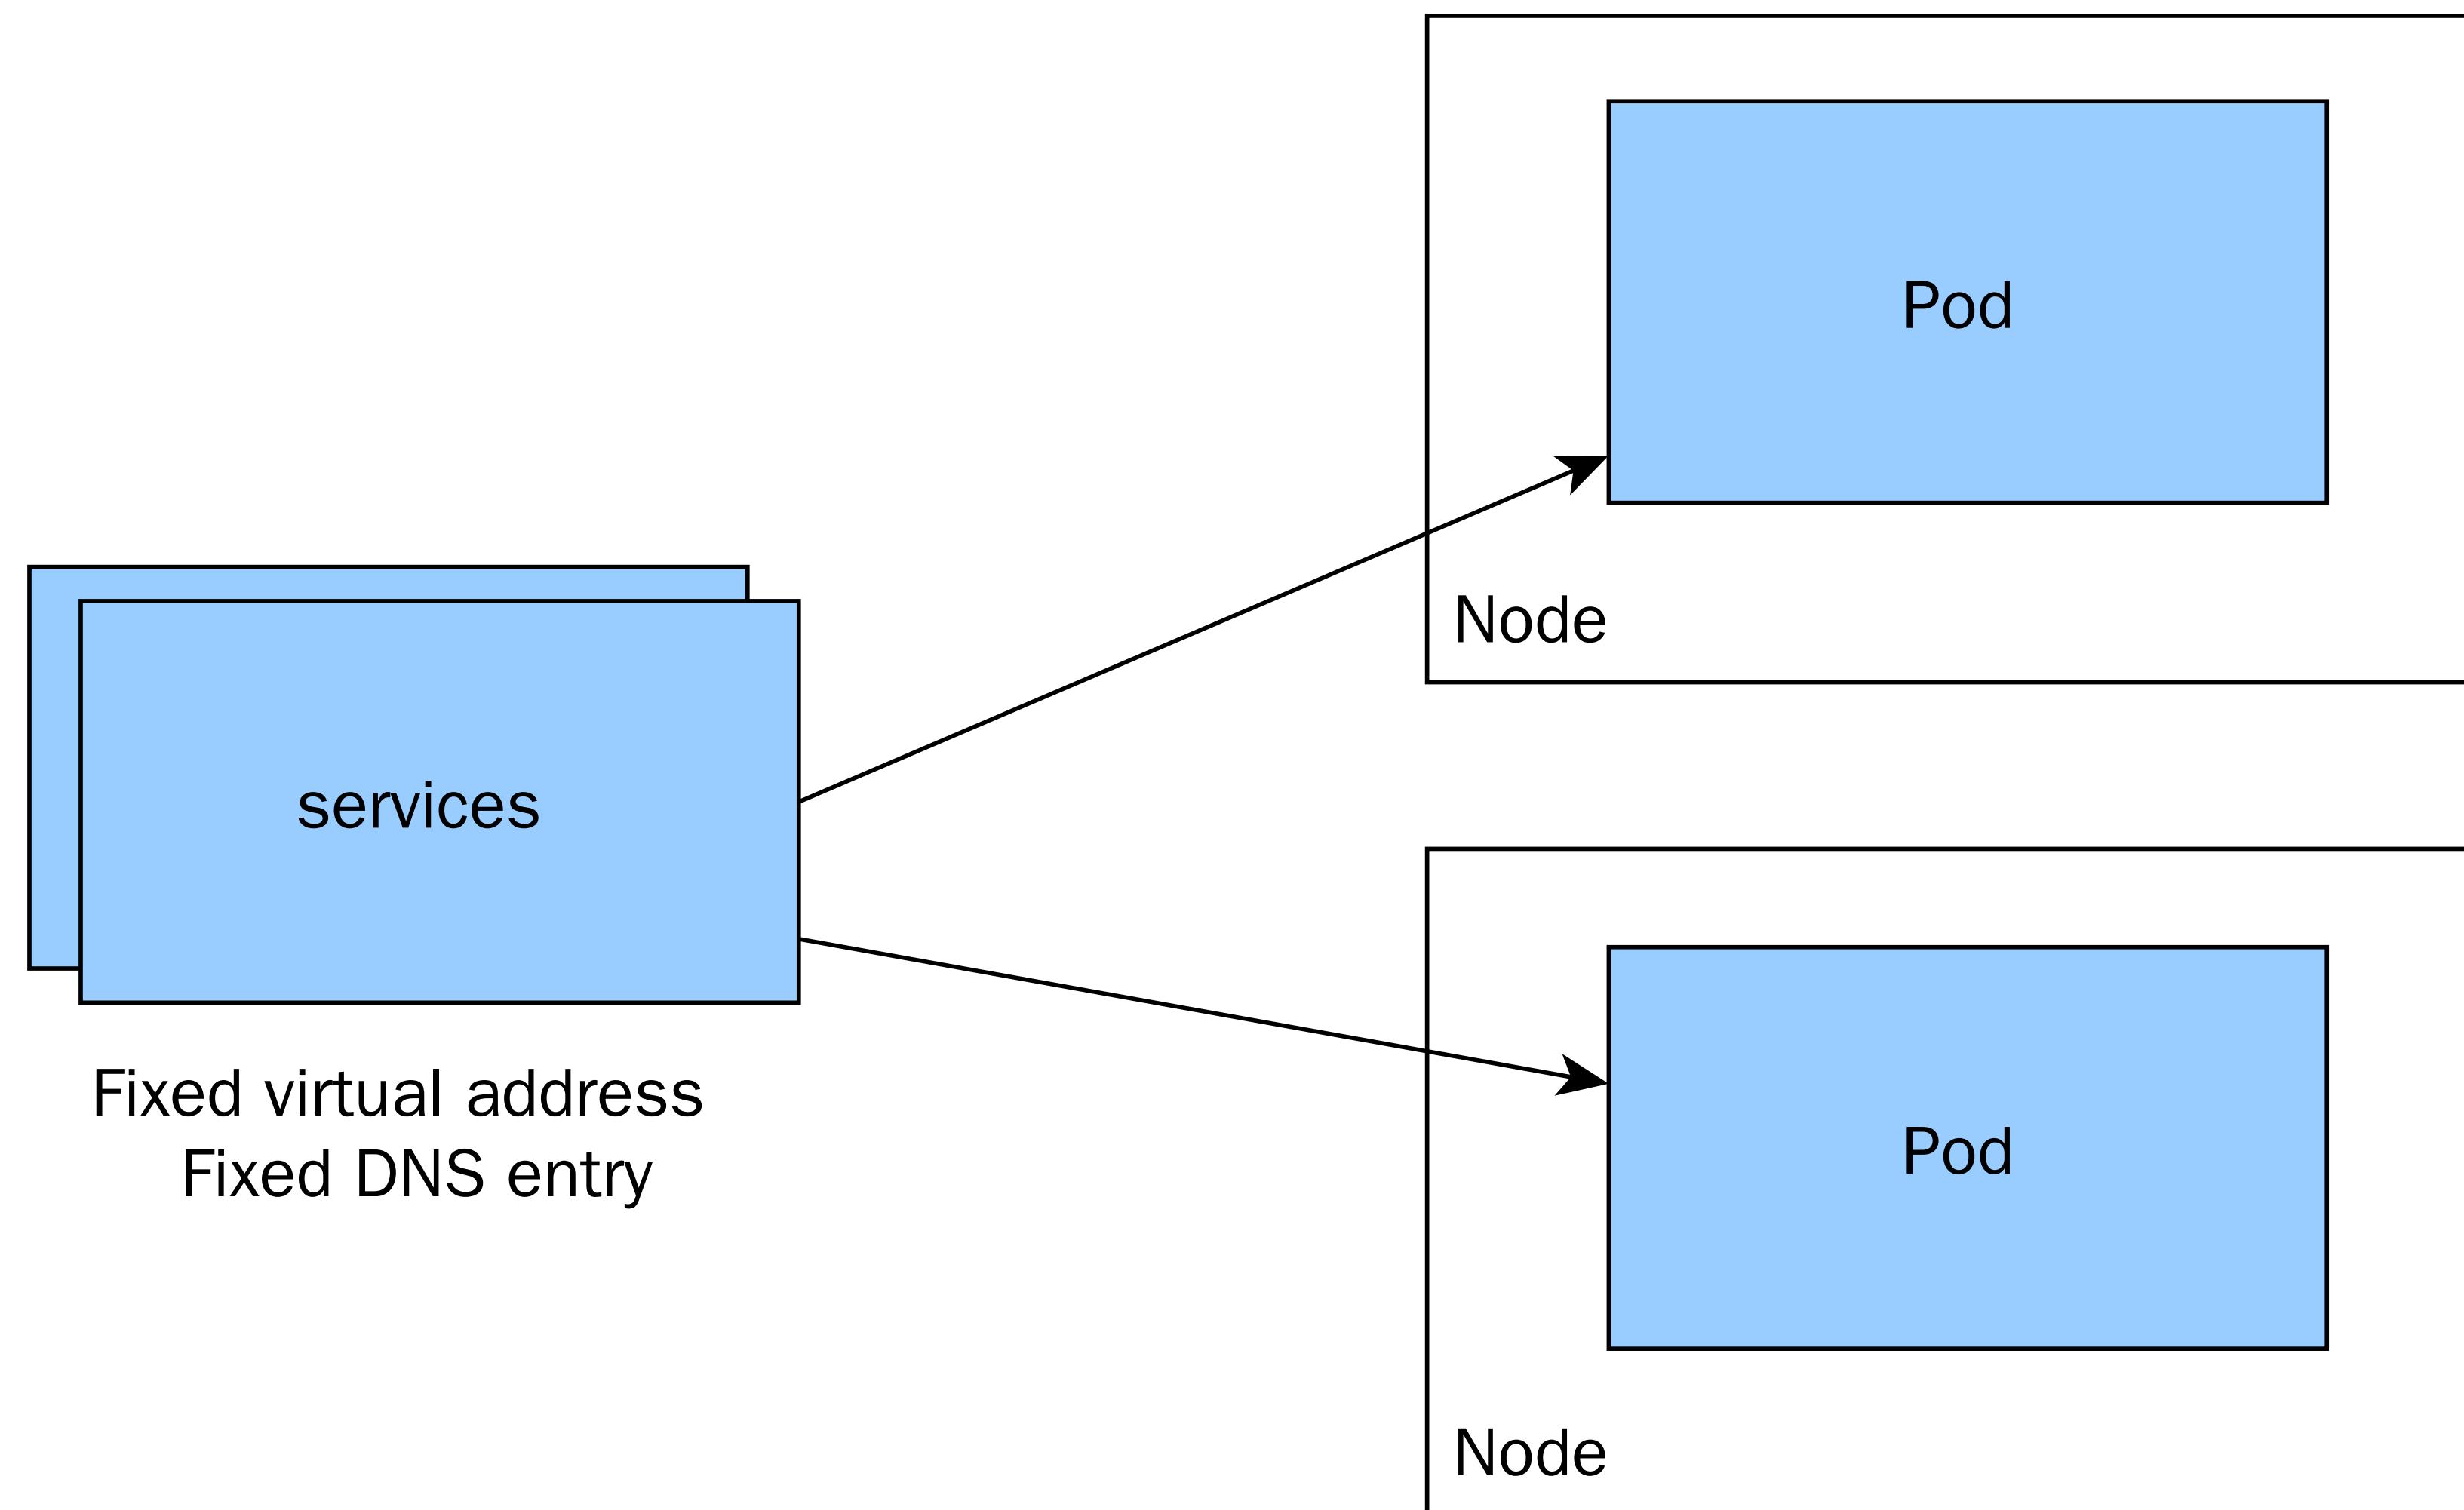
# BLACK BOX

## Resistance and Migrations

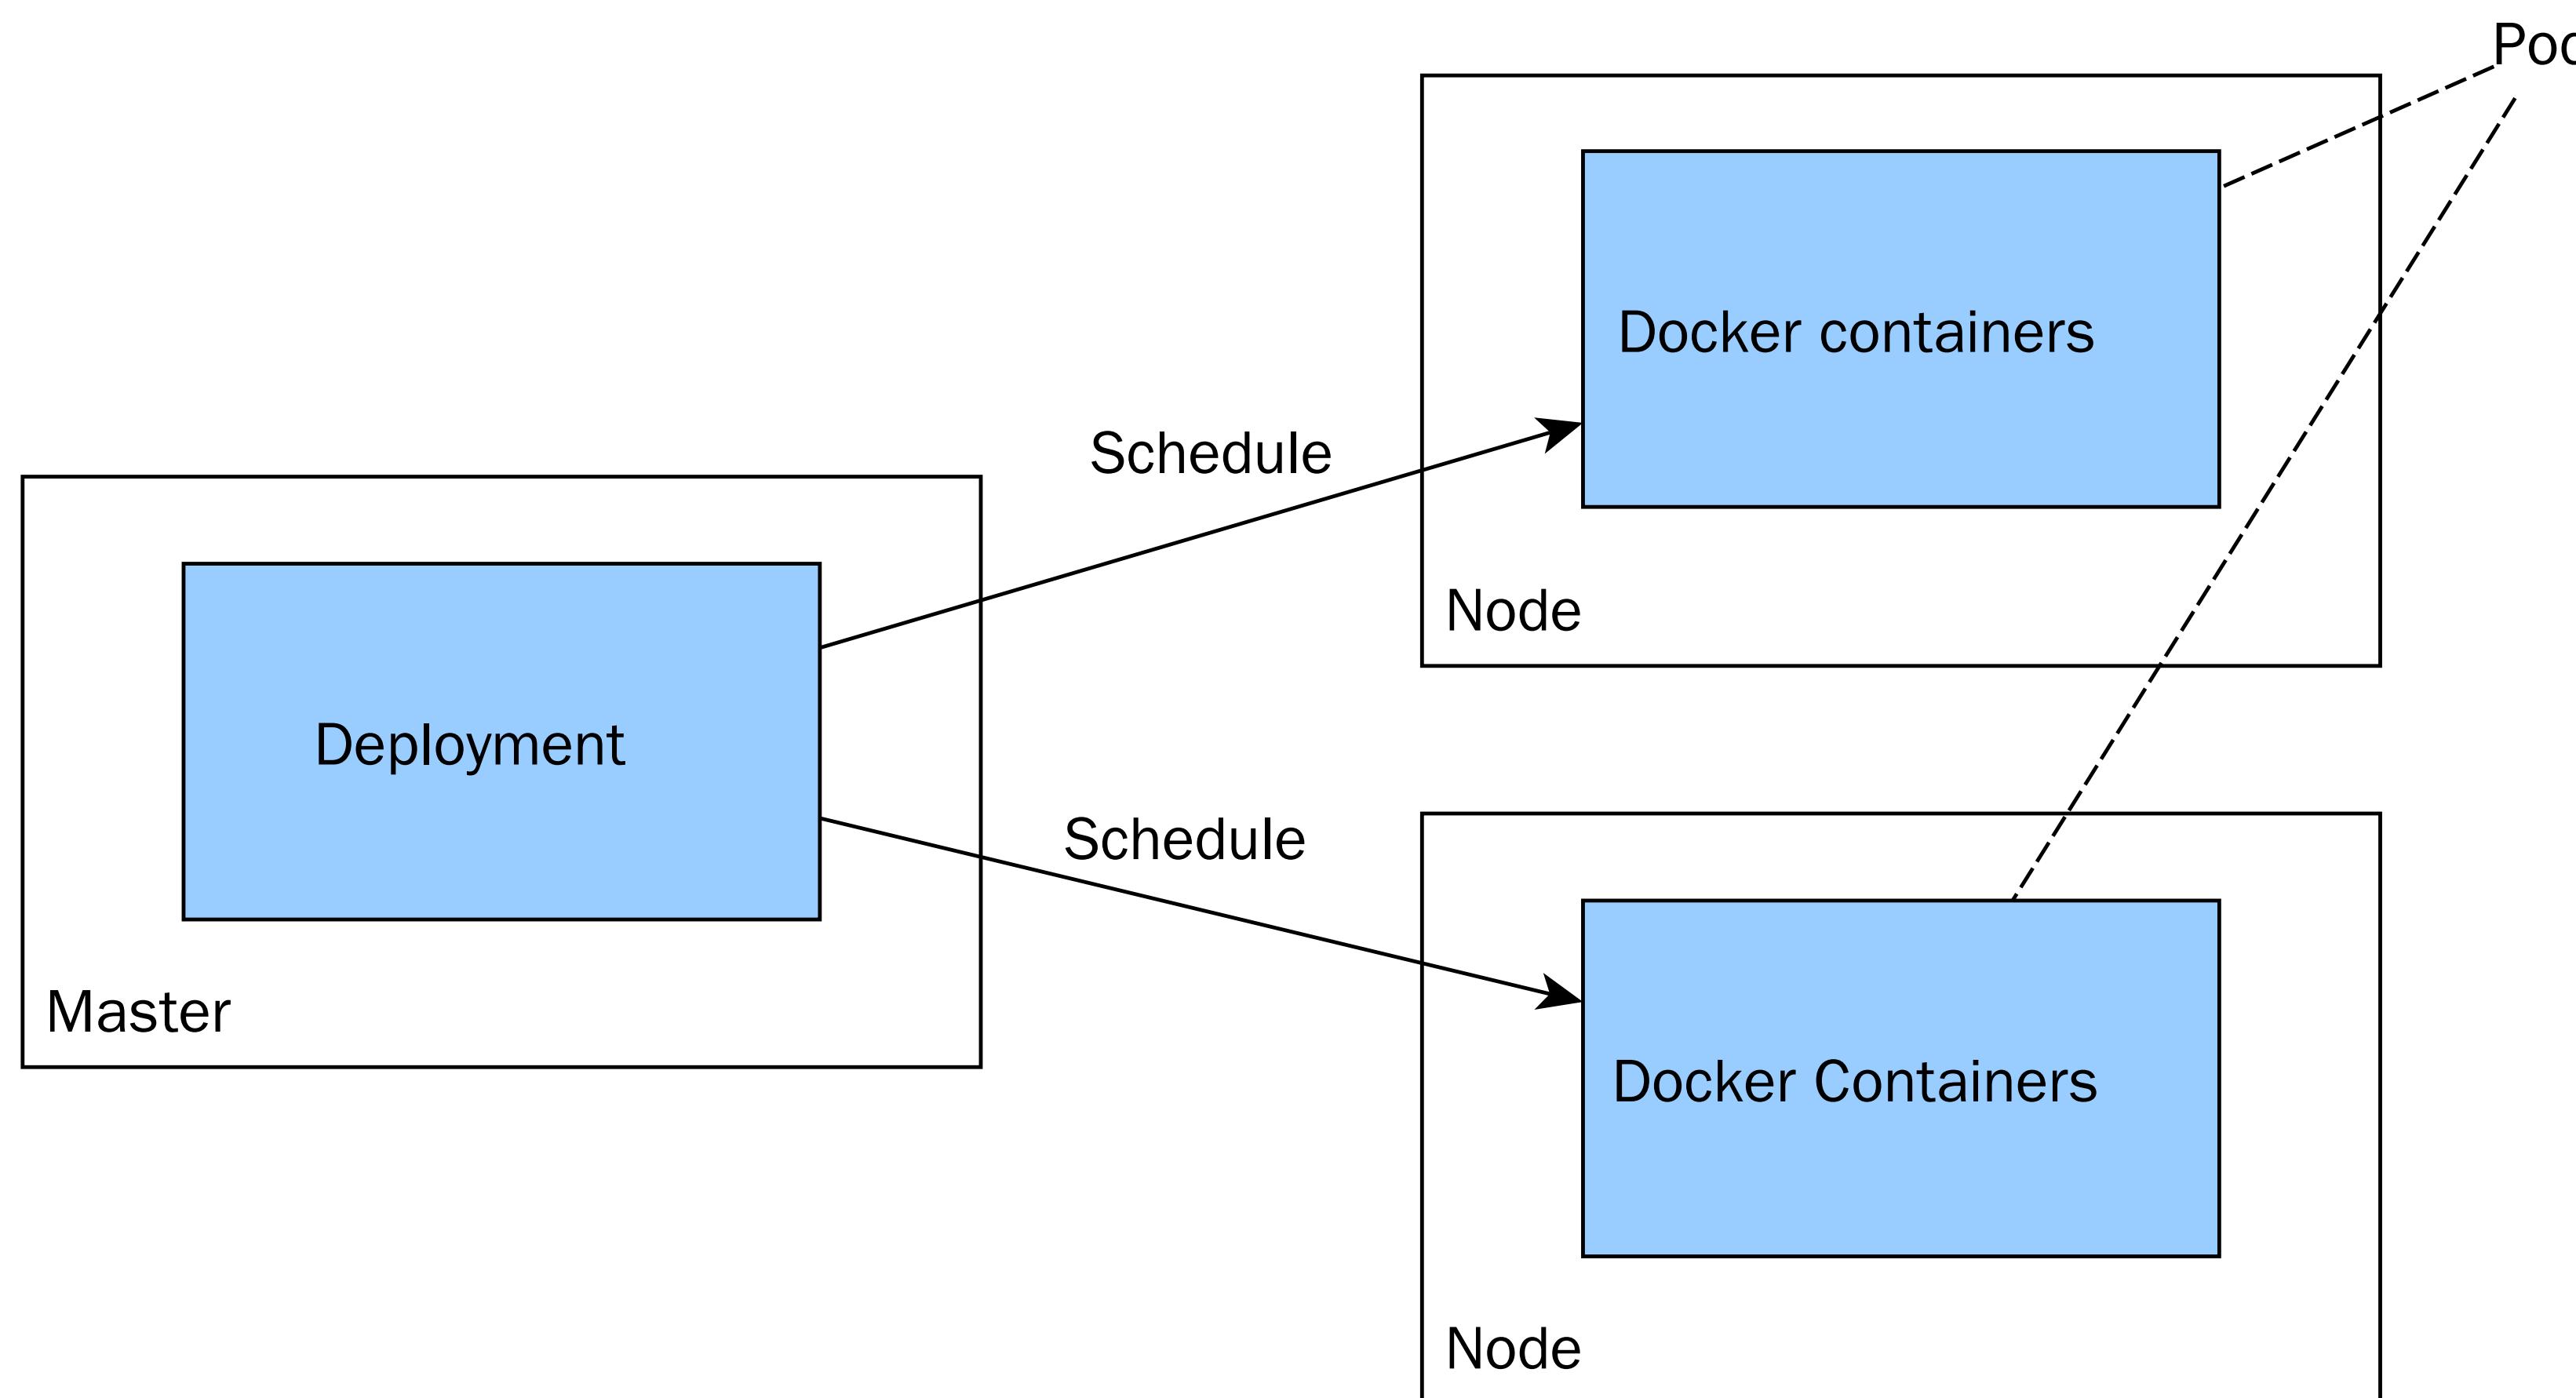


and much more: load balancing

# LINGUA FRANCA



# LINGUA FRANCA



# LIGUA FRANCA

**Pattern**

api.smacc.io/v1/users      users-v1

---

api.smacc.io/v2/users      users-v2

---

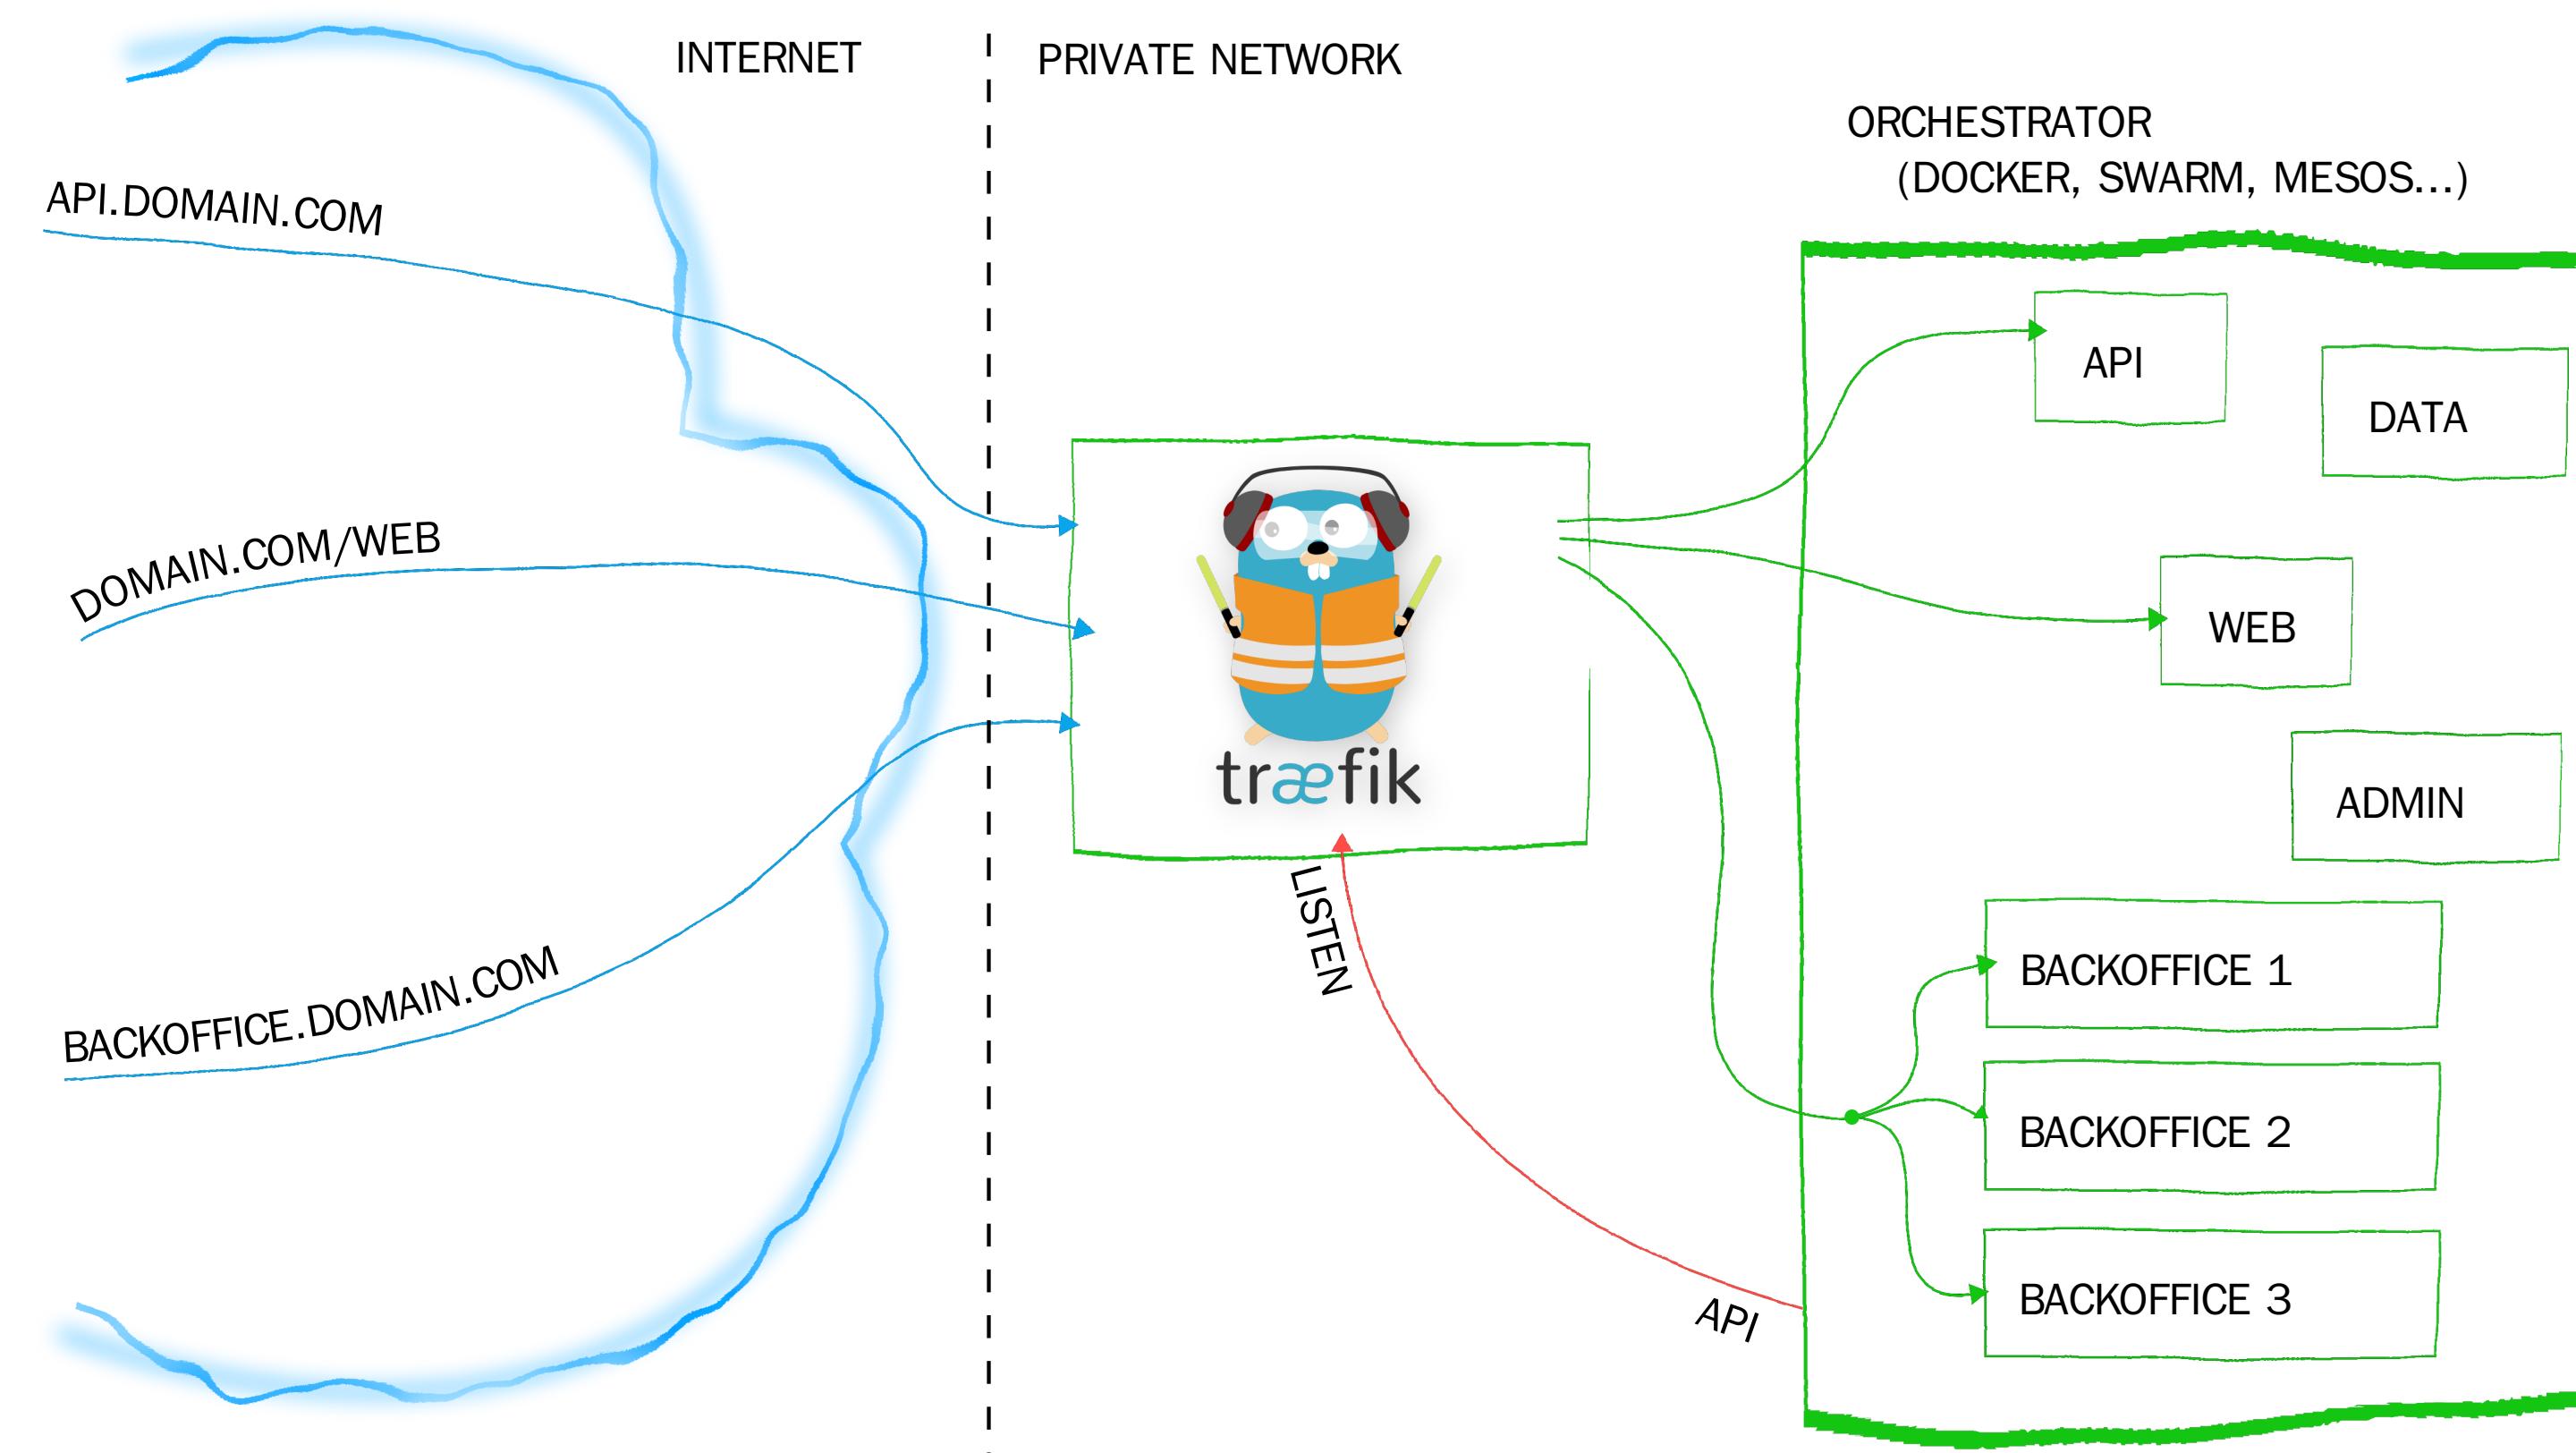
smacc.io

web

**Target App Service**

# LINGUA FRANCA

## Ingress Controller:



Drop-in: traefik, nginx, haproxy, envoy...

# BATTERIES

Service Discovery:

- service injected to DNS:

```
curl http://users/list
```

- labels:

name=value

- annotations:

prometheus.io/scrape: "true"

# SERVICE DISCOVERY

- Loosely couple components
- Auto-wiring with logging and monitoring
- Integration

# Learn-as-you-go experience

Monitoring? Ingress?

1. Drop-in: traefik, nginx, ...
2. oooo ruuuuns!
3. Make a hell of mistakes
4. Get the right one or It is OK for now

# Learn-as-you-go experience

Where to start?

→ With service and deployment

# Learn-as-you-go experience

memcached in k8s?

→ statefulset

## Learn-as-you-go experience

My pod takes all memory

→ resource & limits

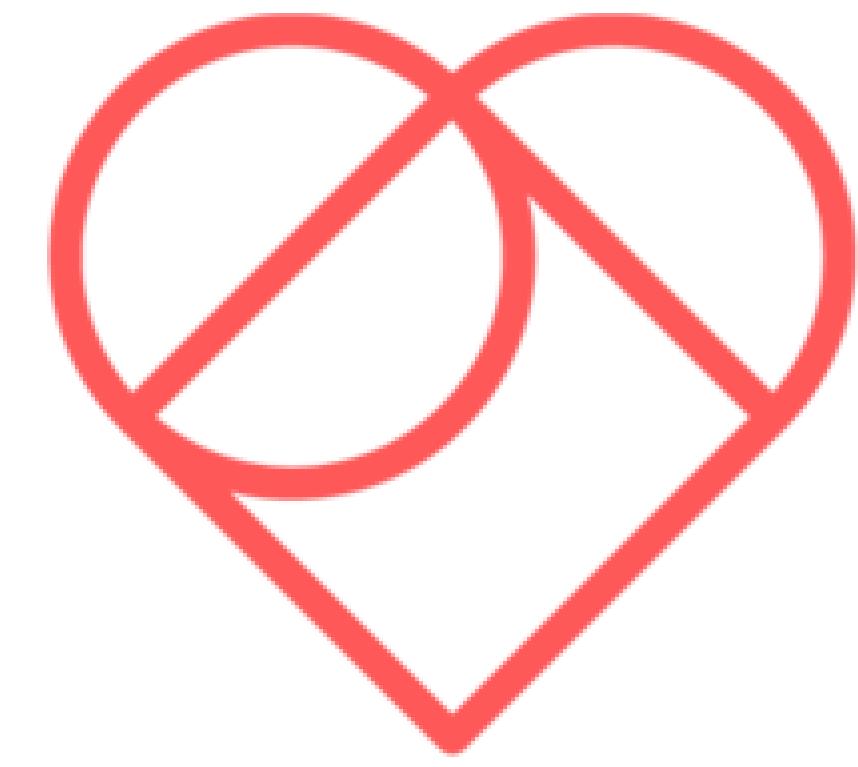
My ML-based app has slow start.

Users noticed downtime.

# LONG LIVE GIT!

- All yaml in your service git
- Integration with e.g.:
  - monitoring, alarming
  - ingress-controller
- Forget about infrastructure... almost

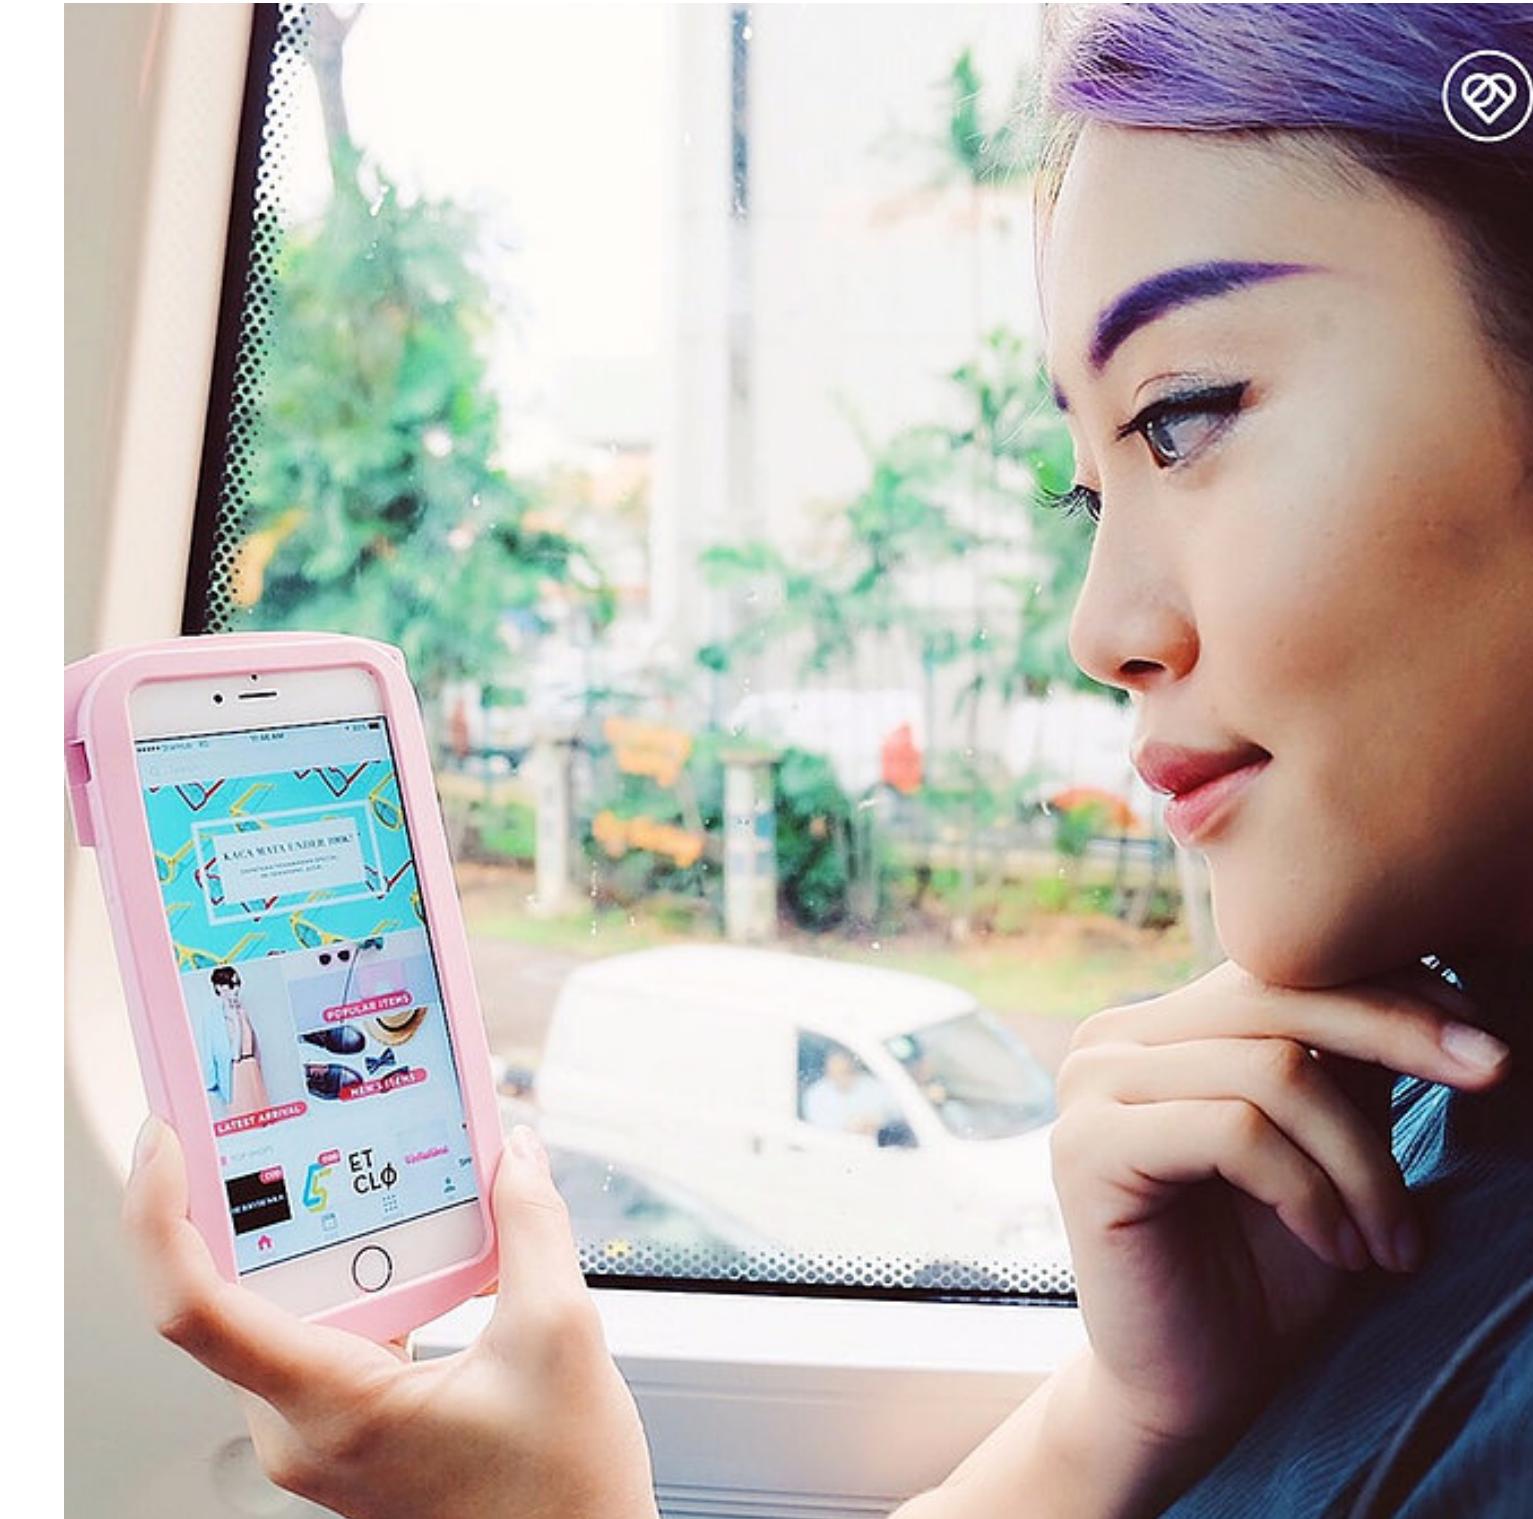
DevOps Culture Dream!



LYKE

# LYKE

- E-commerce Indonesia
- Mobile-only
- 50k+ users
- 2M downloads
- Top 10 Fashion Apps  
Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

# GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,  
UI Designer,  
Mobile Dev,  
and tester - one body



## CHALLENGES

- 50+ VMs in Amazon, 1 VM - 1 App
- 65% Idle machine \$\$\$
- Puppet with manual deployment process
- Fear, Forgotten components
- Performance issues

## APPROACH

1. Simplify Infrastructure
2. Change the Development Practices  
(12factor + kubernetes)
3. Change the Work Organization

see: Conway's law

## SIMPLIFY

1. Kubernetes with Google Kubernetes Engine
2. Terraform for all new
3. Database-as-a-service

## SIMPLIFY

1. Prometheus, AlertManager, and Grafana
2. Elasticsearch-Fluentd-Kibana
3. Google Identity-Aware-Proxy  
to protect all dev dashboards
4. 3rd party SaaS: statuscake and opsgenie

One person efford

## CONTINUOUS DEPLOYMENT

- branch-based:
  - master
  - staging
  - production
- repo independent

## TRAVISCI

1. Tests
2. Build docker
3. Deploy to Google Container Registry
4. Deploy to k8s only new docker
5. no config/secrets applied

# GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \ - kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```

# Makefile

```
SERVICE_NAME=v-connector
GCP_DOCKER_REGISTRY=eu.gcr.io

test: test_short test_integration

run_local:

docker_build: docker_push

kube_create_config:

kube_apply:

kube_deploy:
```

Copy&Paste from the project to project

# 1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

## 2. GET BACK ALL THE KNOWLEDGE

Extract from:

- Puppet, ... ➔ Dockerfile
- Running Instances ➔ Dockerfile, README.rst
- Nagios, ... ➔ README.rst + *checks/*

## **3. INTRODUCE RUN\_LOCAL**

- make `run_local`
- A nice section on how to run in README.rst
- with `docker-compose`

The most crucial point.

## 4. GET TO KUBERNETES

- make kube\_create\_config
- make kube\_apply
- Generate the yaml files if your envs differ

secrets from gopass (password manager)

## 5. CONTINUOUS DEPLOYMENT

Travis:

- the same Makefile as dev use

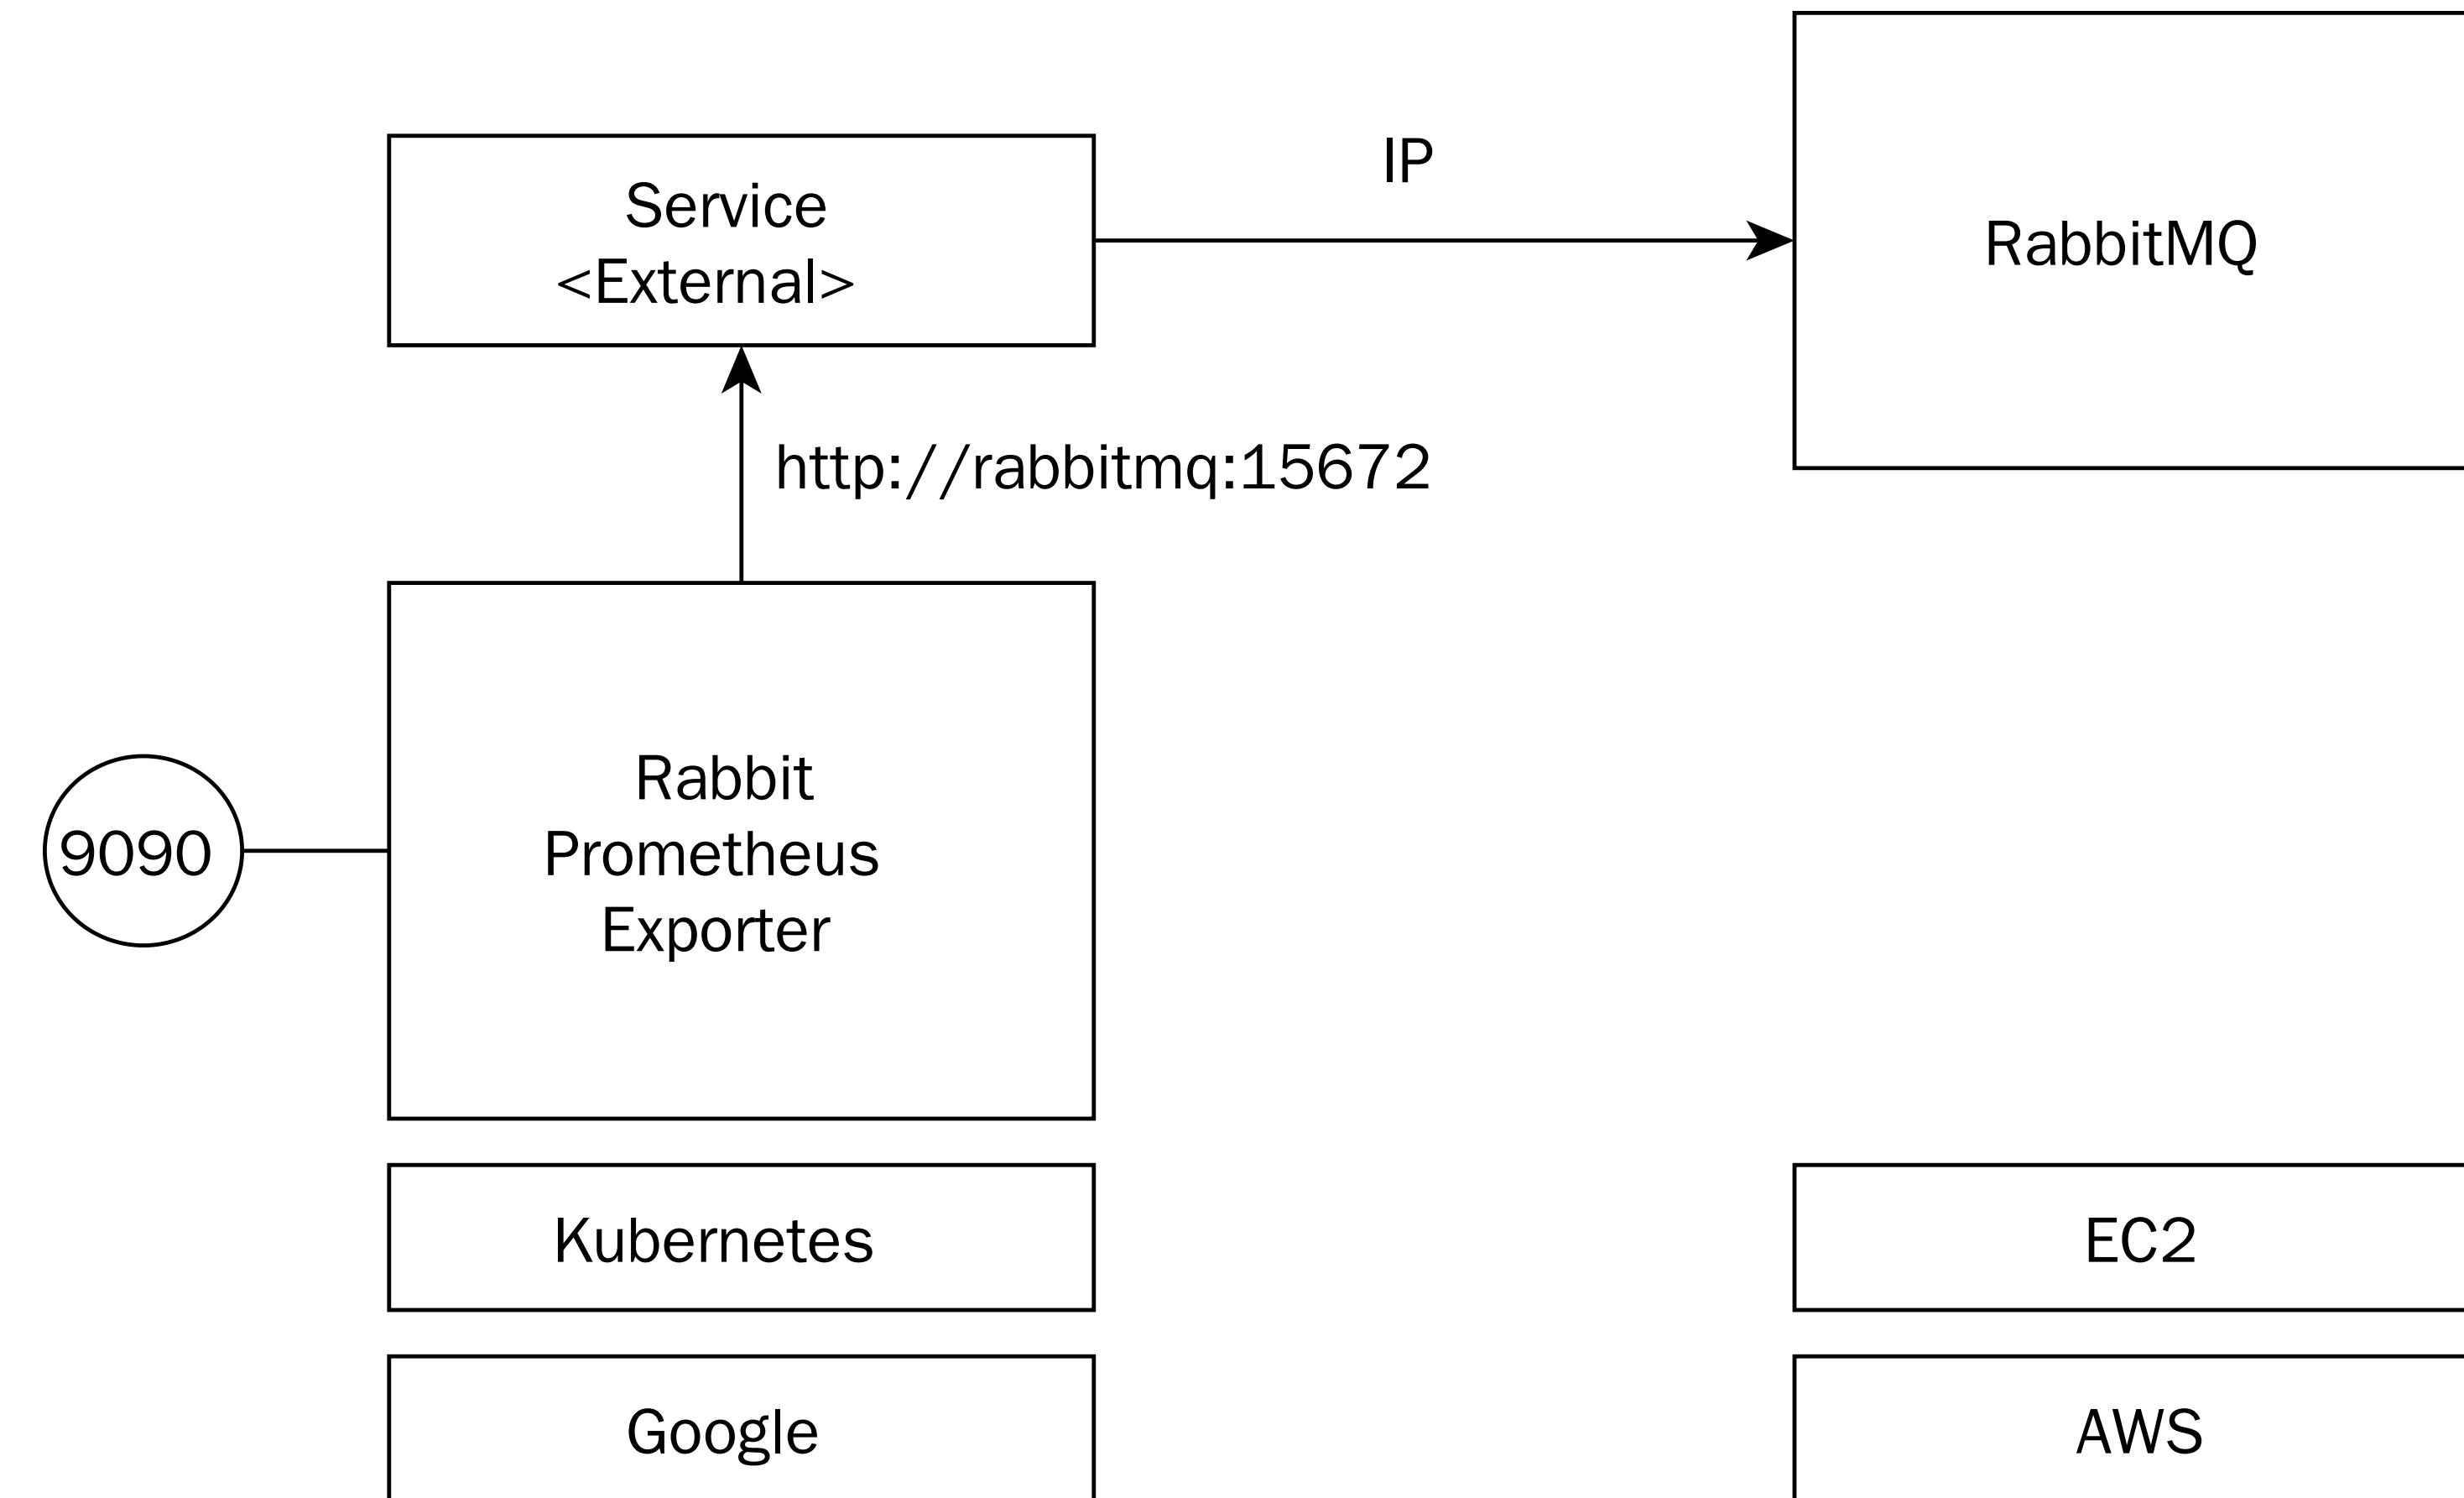
## 6. KEEP IT RUNNING

Bridge the new with old:

- Use External Services in Kubernetes
- Expose k8s in the Legacy [1]

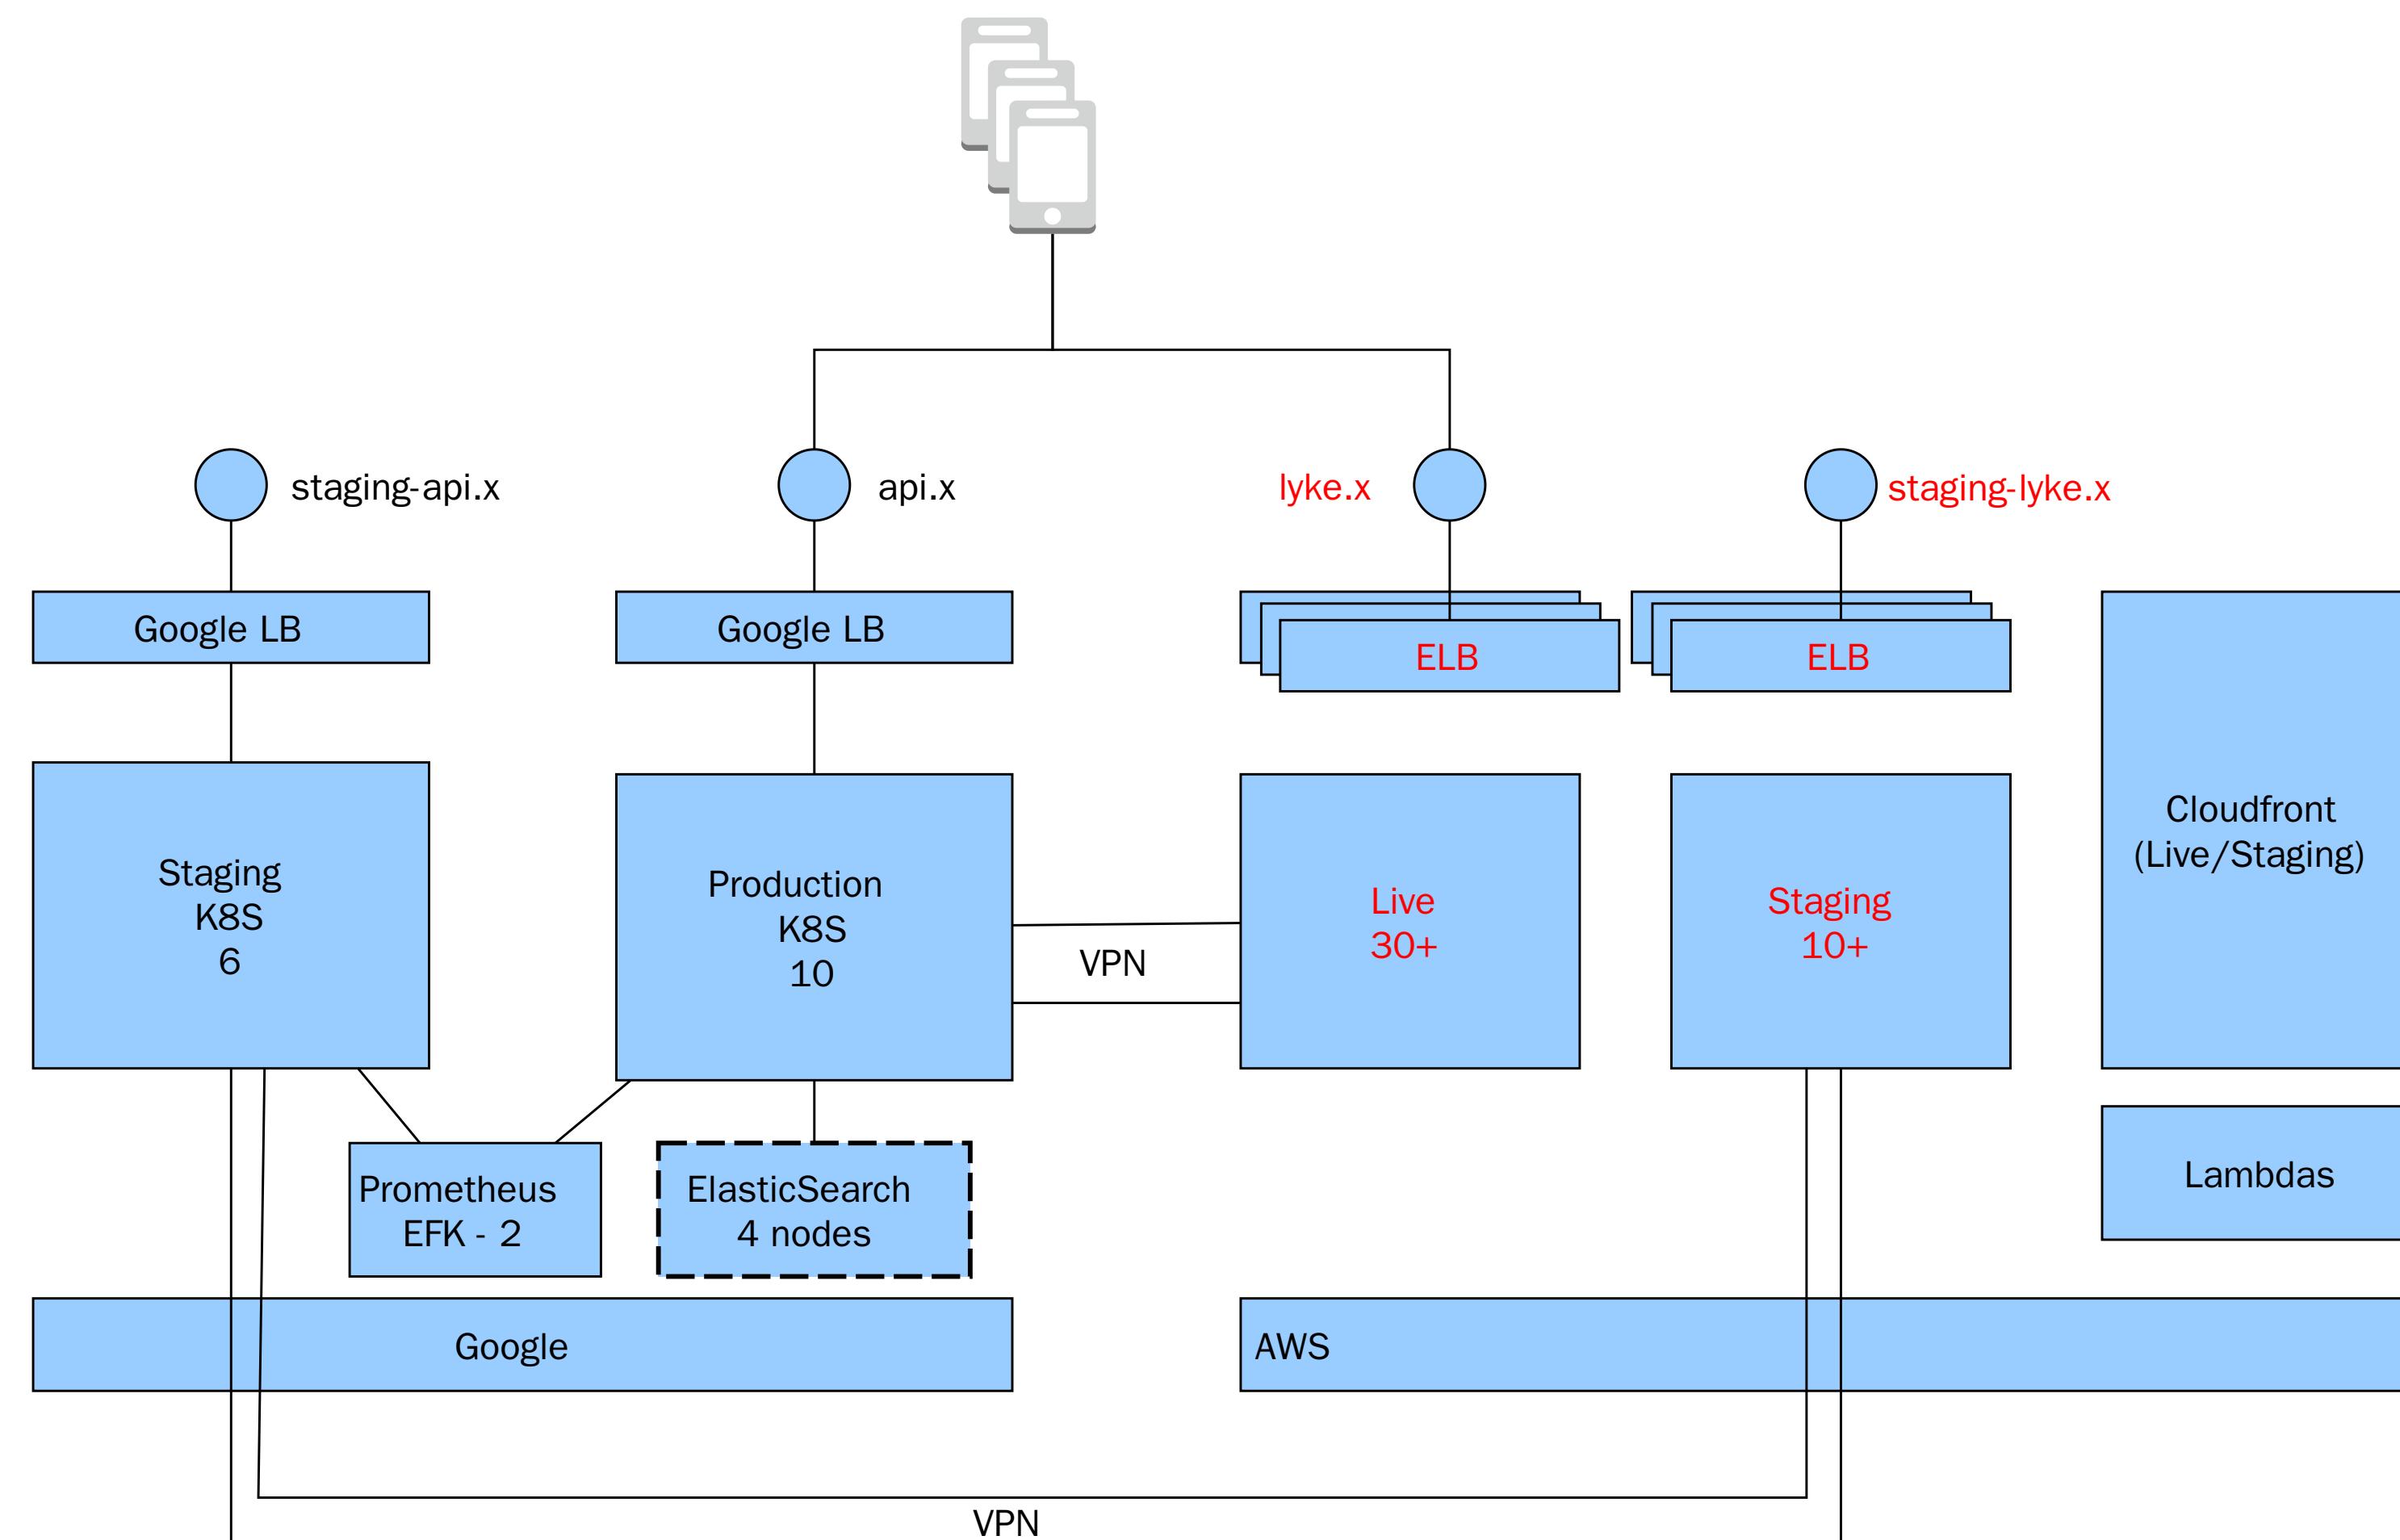
[1] hard coded IP:PORT, considered: K8S events to Consul

# Bridge the new with old



Monitor legacy with new stack

# Architecture During Migration



## 7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test  
TARGET_URL=api.example.com/users make smoke_test
```

## 8. SERVICE SELF-AWARE

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

## 9. MOVE TO MICRO-SERVICES

Offload Legacy:

- Keep the lights on
- New functionality to micro-services

## WHAT WORKED

1. Copy&Paste Makefile and k8s yaml
2. Separate deployments a good transition strategy

## WHAT DID NOT WORK

1. Too many PoC, cut to 2 weeks max
2. Doing it with too small steps
3. Push back to k8s yaml [\*]
4. Alert rules too hard to write

[\*] With coaching, I thought, it is OK

## DO DIFFERENT

1. Move data day one
2. Make devs know it is a transition stage
3. Teach earlier about resources
4. EFK could wait
5. All-hands for a paid-XXX% weekend for migration

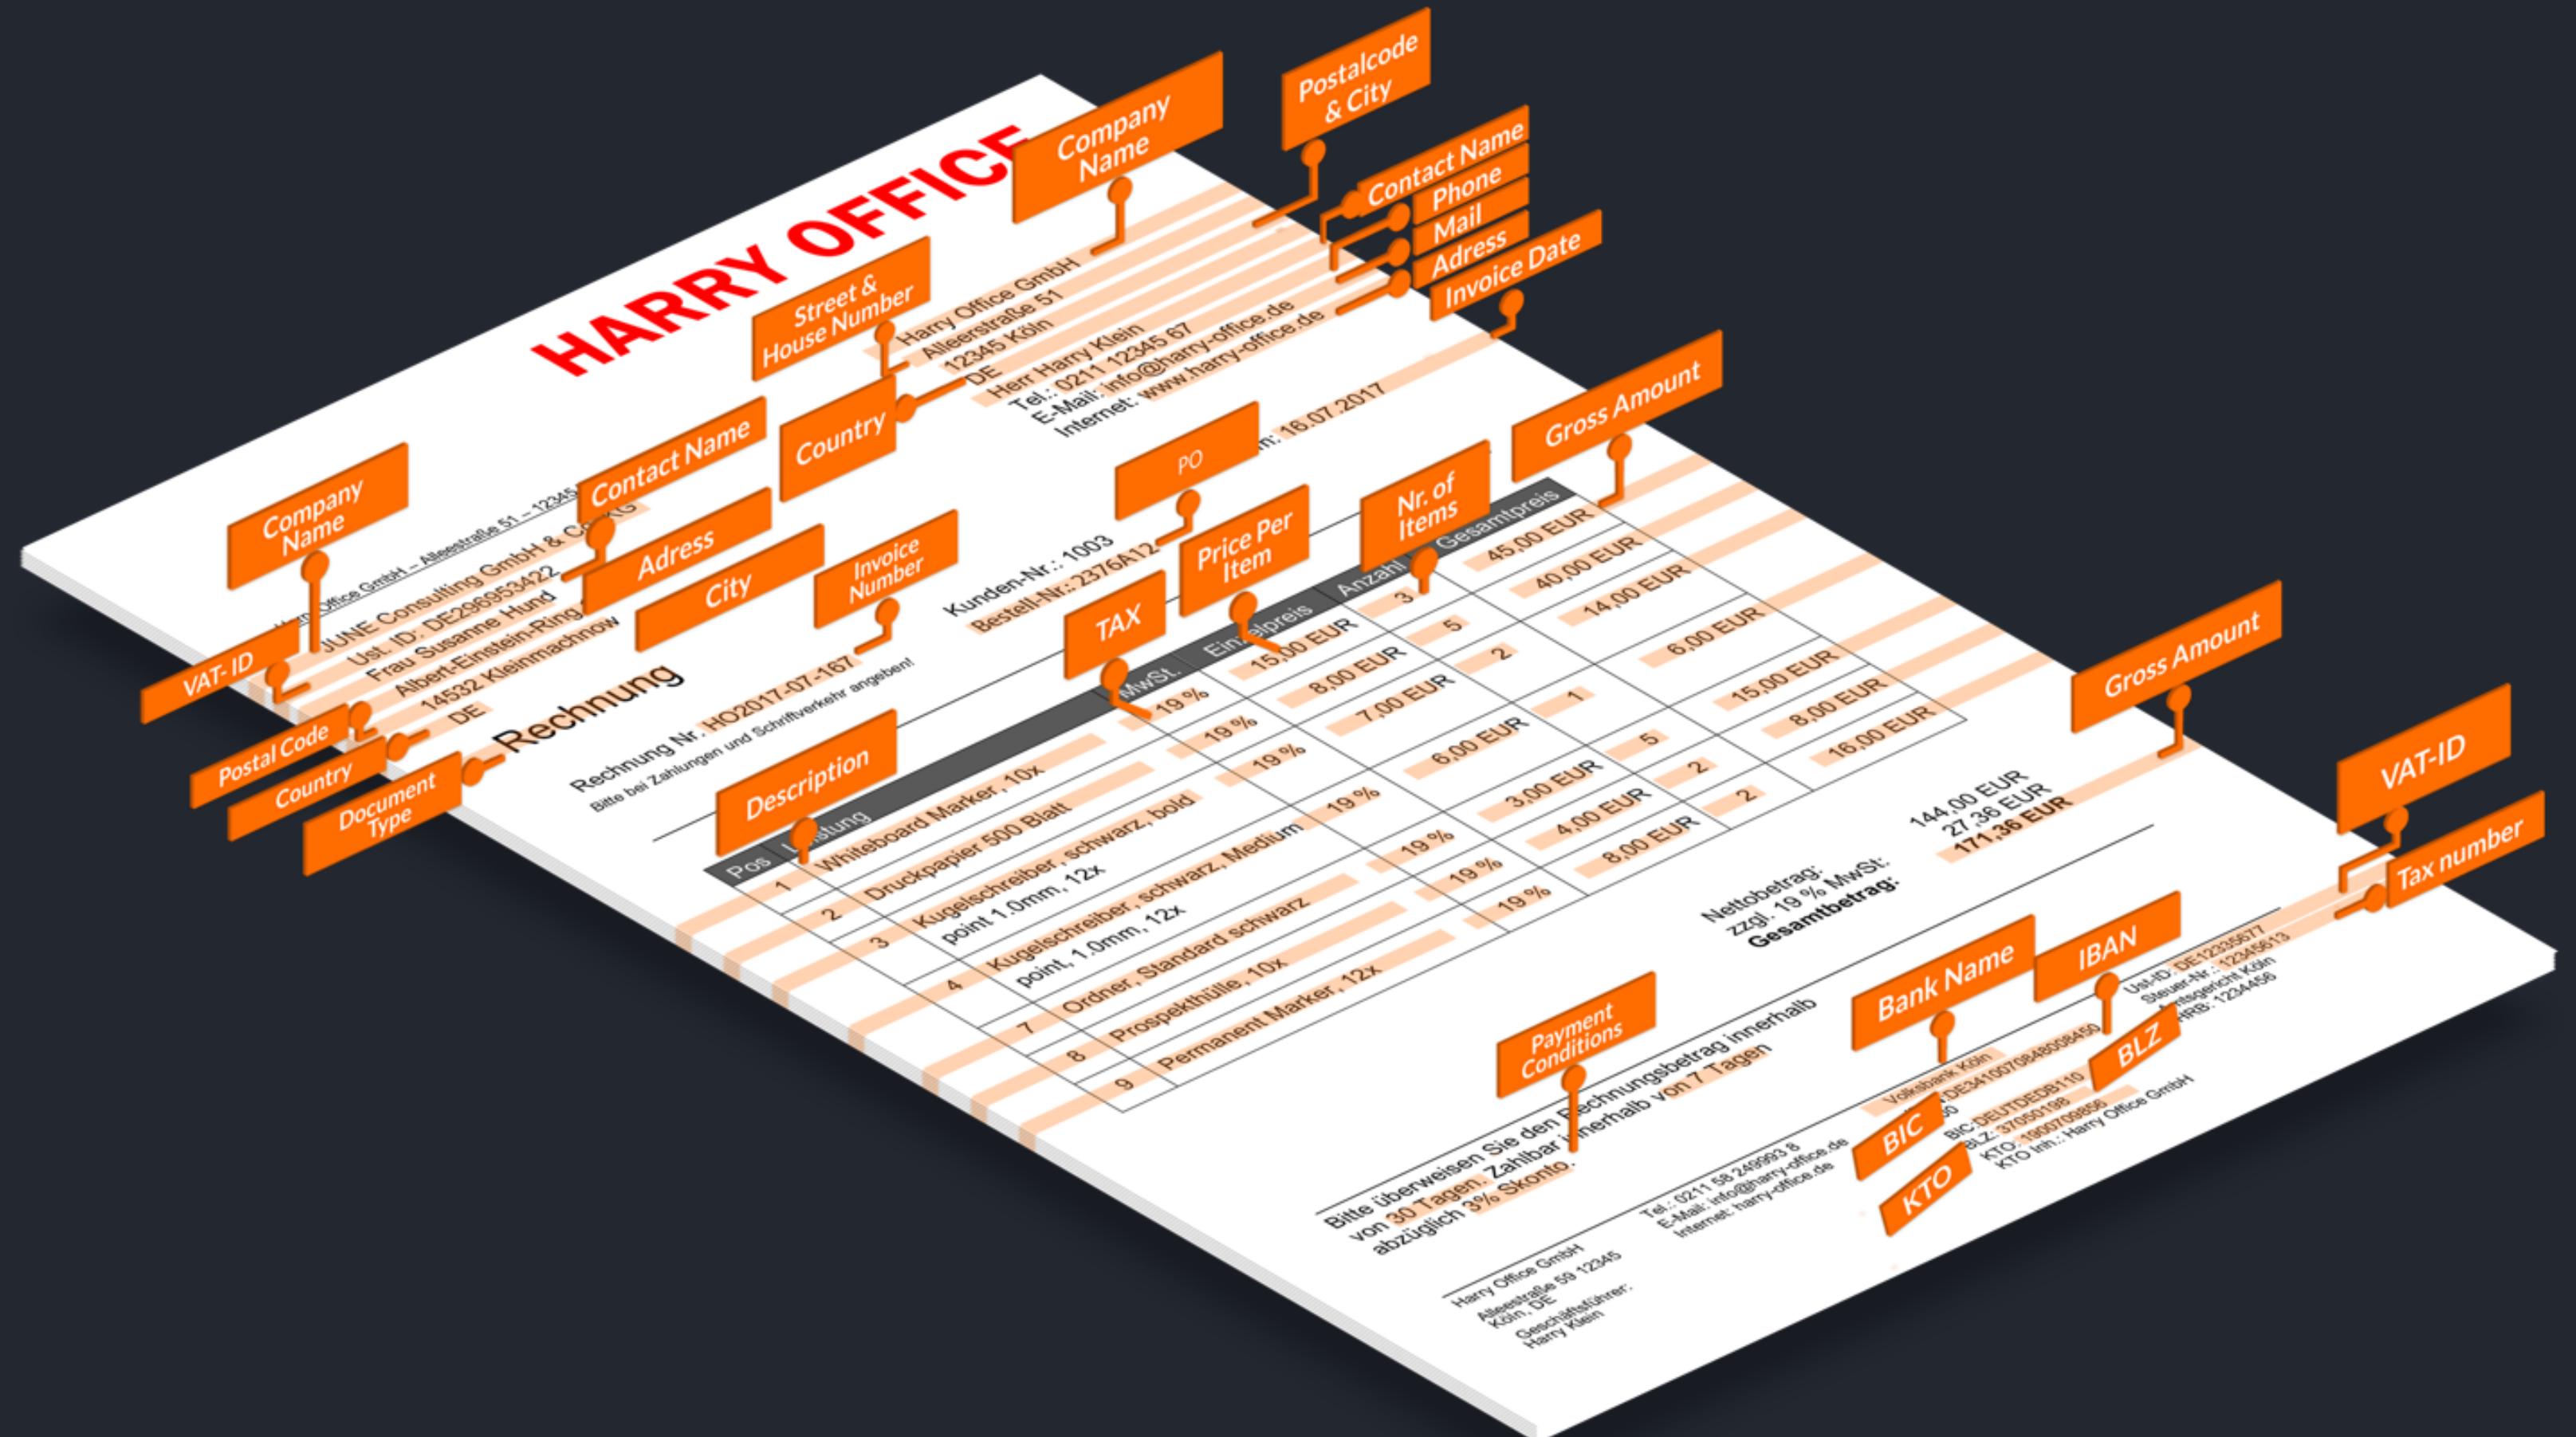
# SMACC

---

## Hypatos

Hypatos

# Problem SMACC solves



SMACC

## Global clients

Deutsche Bank



BMB  
巴三巴  
GRUPPE

Deloitte.



accenture >

McKinsey&Company

Hypatos

SMACC

# STORY

1. Legacy on AWS and AWS ECS :/
2. Self-hosted K8S on ProfitBricks (PB)
3. ooo... K8S can use Azure Active Directory :D

# STORY

4. Get to Microsoft ScaleUp
5. Welcome Azure!
6. Luckily. AKS
7. Easy migration from PB to AKS

# AZURE KUBERNETES SERVICE

- Independent from IaaS
- Plug and play
- Integration with GPU
- Our OnPrem = Our OnCloud

# SIMPLICITY

- az aks CLI for setting k8s - README.rst
- Terraform for everything else
- Secrets: 1Password and gopass.pw

Terraform also sets our AWS

# DIFFERENCE ☠

- Two teams in Berlin and Warsaw
- Me in Warsaw

# NEW EXPERIENCE

Devs do not:

- like TravisCI, Makefiles, Yaml
- feel it is too much hasle

Transition from PB to AKS was painful

# SOLUTION

- make everything lighter
- c&p without modifications
- hide the k8s, remove magic
- deploy on *tag*

Similar to the [Kelsey Hightower approach](#)

# Repo .travis.yml

```
language: go
go:
- '1.16'
services:
- docker
install:
- curl -sL https://$GITHUB_TOKEN@raw.githubusercontent.com
- if [ -f "tools/travis/install.sh" ]; then bash tools/travis/install.sh; fi
script:
- dep ensure
- make lint
- make test
- if [ -z "${TRAVIS_TAG}" ]; then make snapshot; fi;
deploy:
```

# Makefile

```
| - tools
|   | - Makefile
|   | - kube-service.yaml
|   \- kube-deployment.yaml
|
| - Dockerfile
\- Makefile
```

Makefile only tasks for dev

# CONTINUOUS DEPLOYMENT

- Github
- TravisCI
- hub.docker.com
- AKS

# CONTINUOUS DEPLOYMENT

1. git tag and push
2. smacc-platform.git
3. Deploy to staging
4. PR to production

# KUBERNETES

- Pure, generated, kubernetes config
- 2x kubernetes operators

# WHAT WORKED

- Hiding k8s
- Understandable CD process

# **WOULD DO DIFFERENT**

- More sensitive to feedback

# NEXT

- Acceptance tests listen on k8s events
- Deployment tool based on [missy](#)
- Keeping an eye on Istio

# K8S - Linux / App server?

- Out-of-box integration
- Lingua Franca - AWS Service Operator
- Learn as you go
- onPremise = onCloud = OnLocal (e.g., [kube-desktop](#))

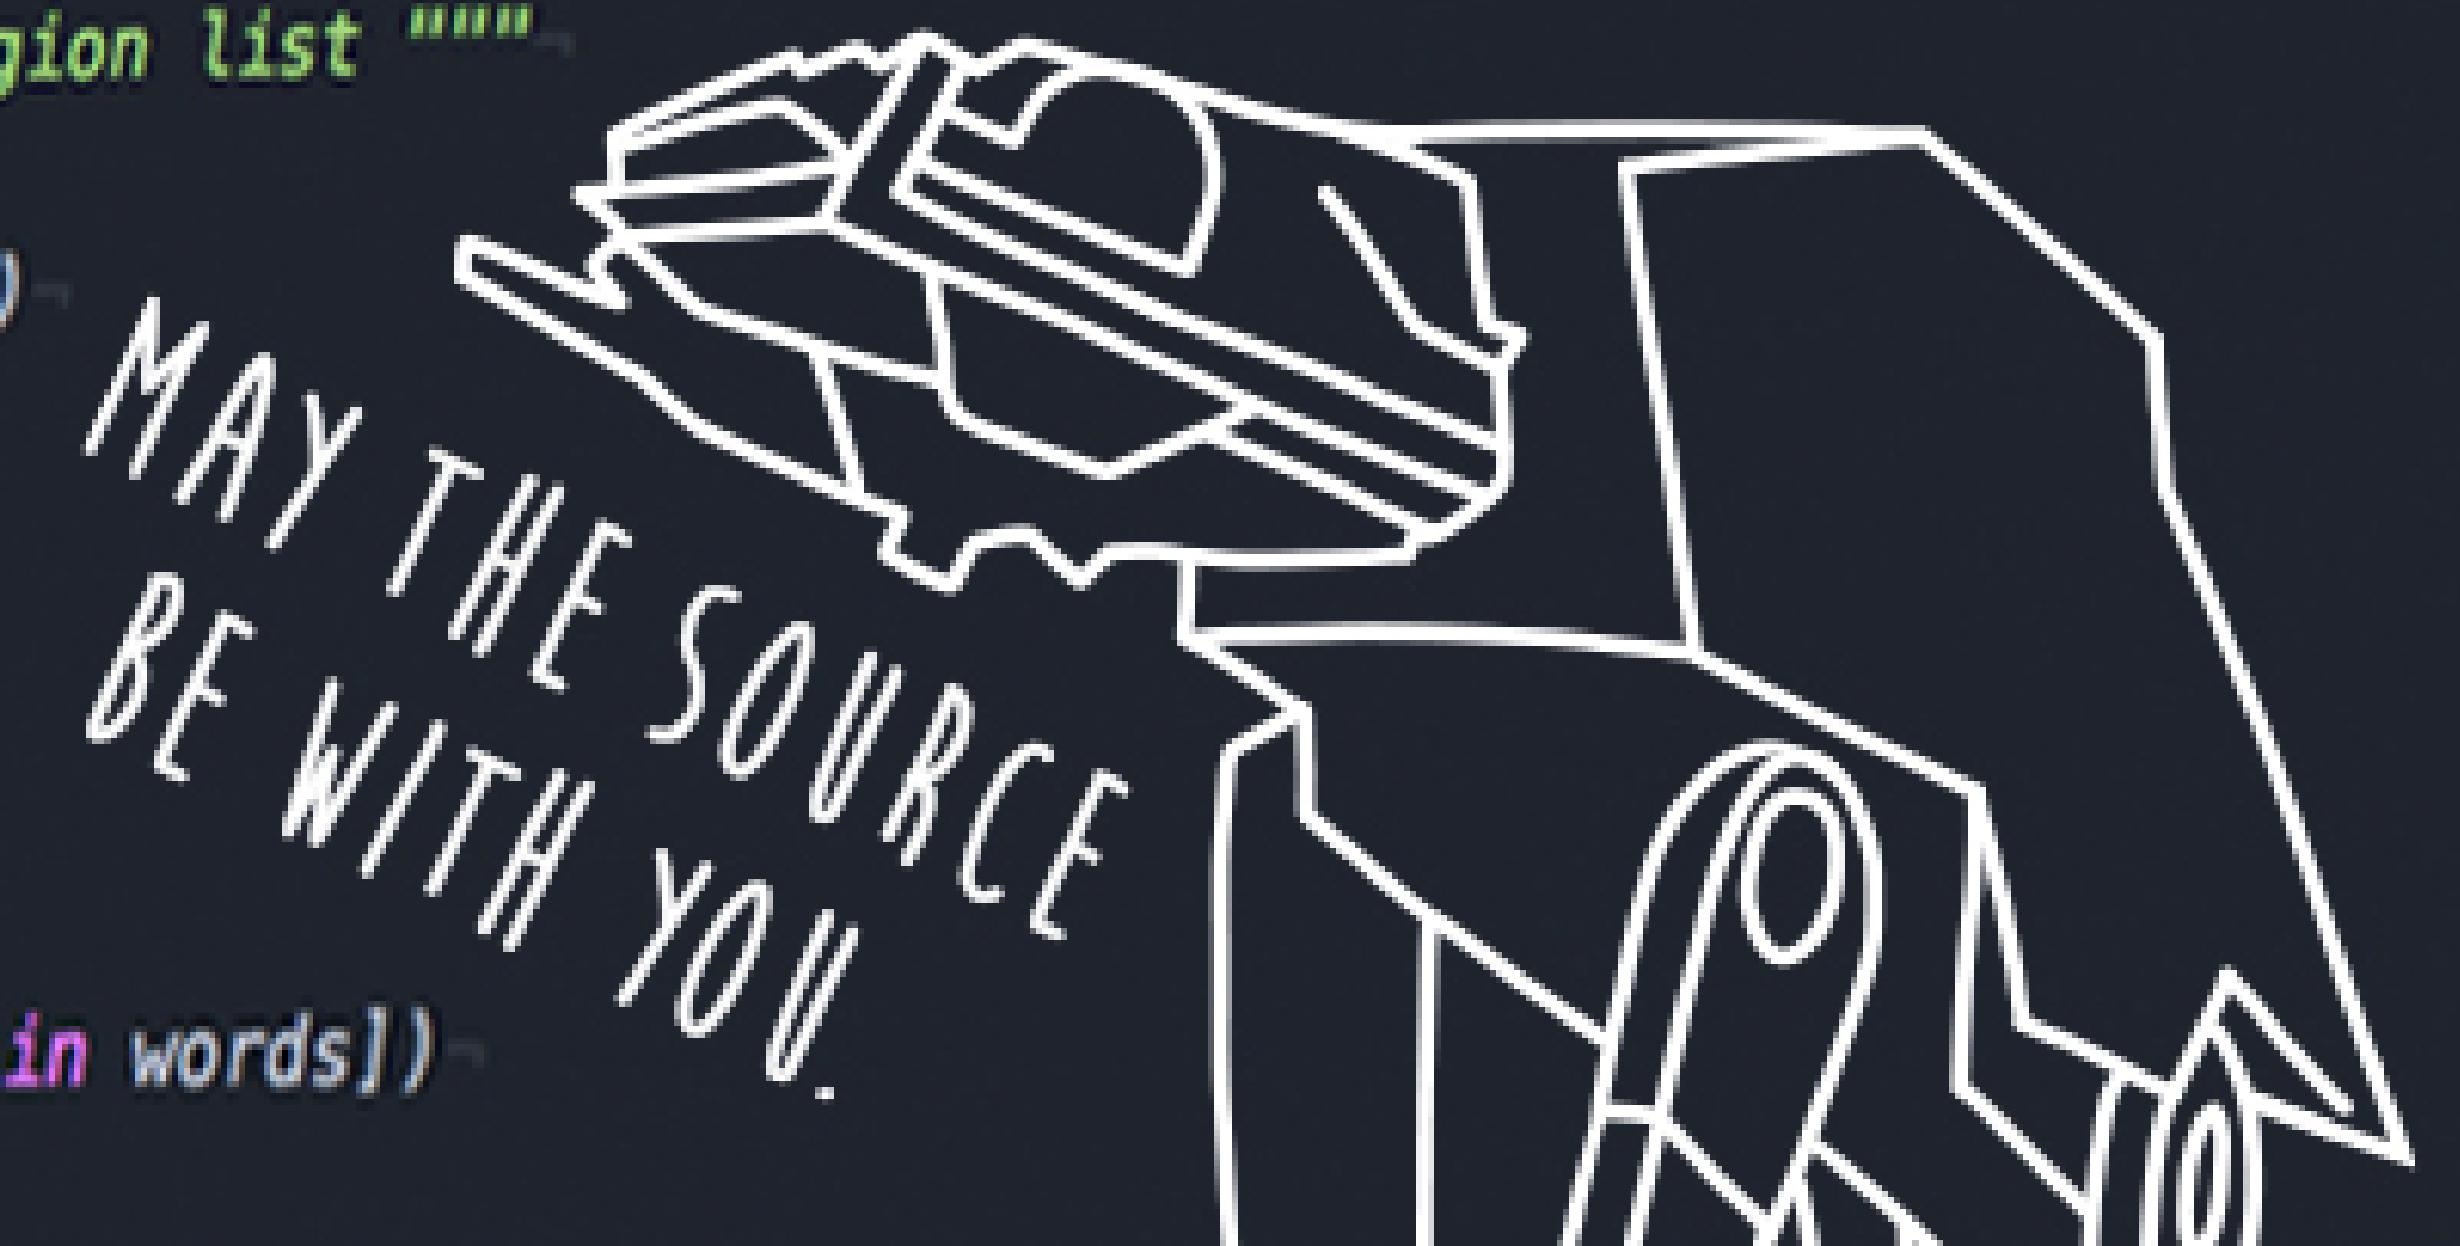
## K8S - Linux / App server?

- Do not terrorize your devs with K8S
- No free lunch... app must be smarter
- On VM vs K8S vs Lambdas?

# THANK YOU. QUESTIONS?

ps. We are hiring.

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

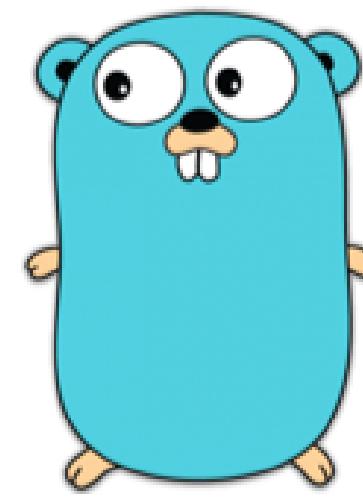


[github.com/wojciech12/talk\\_cloudnative\\_and\\_kube  
rnetes\\_waw](https://github.com/wojciech12/talk_cloudnative_and_kubernetes_waw)

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# SMACC



Go



PYTORCH

TensorFlow™



amazon  
web services™



Azure

λ



# BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# CHANGE THE WORK ORGANIZATION

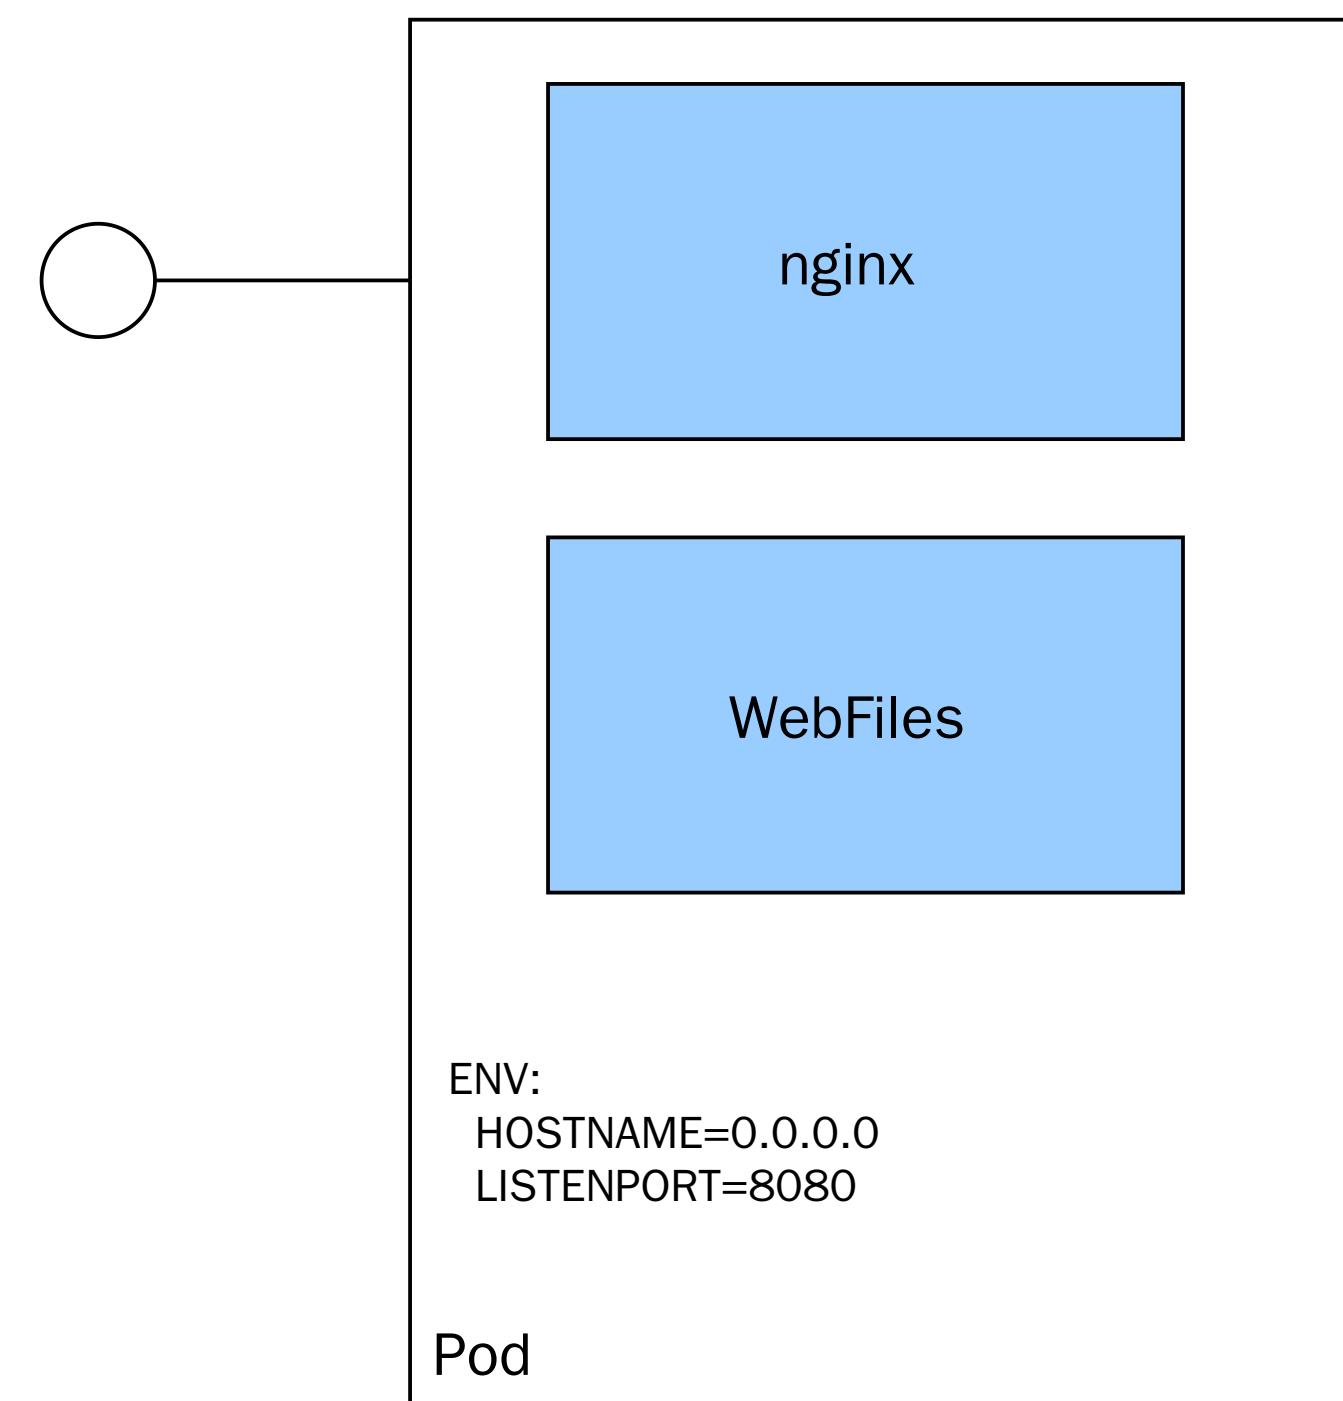
- From Scrum
- To Kanban
- from Kanban to Squads

For the next talk

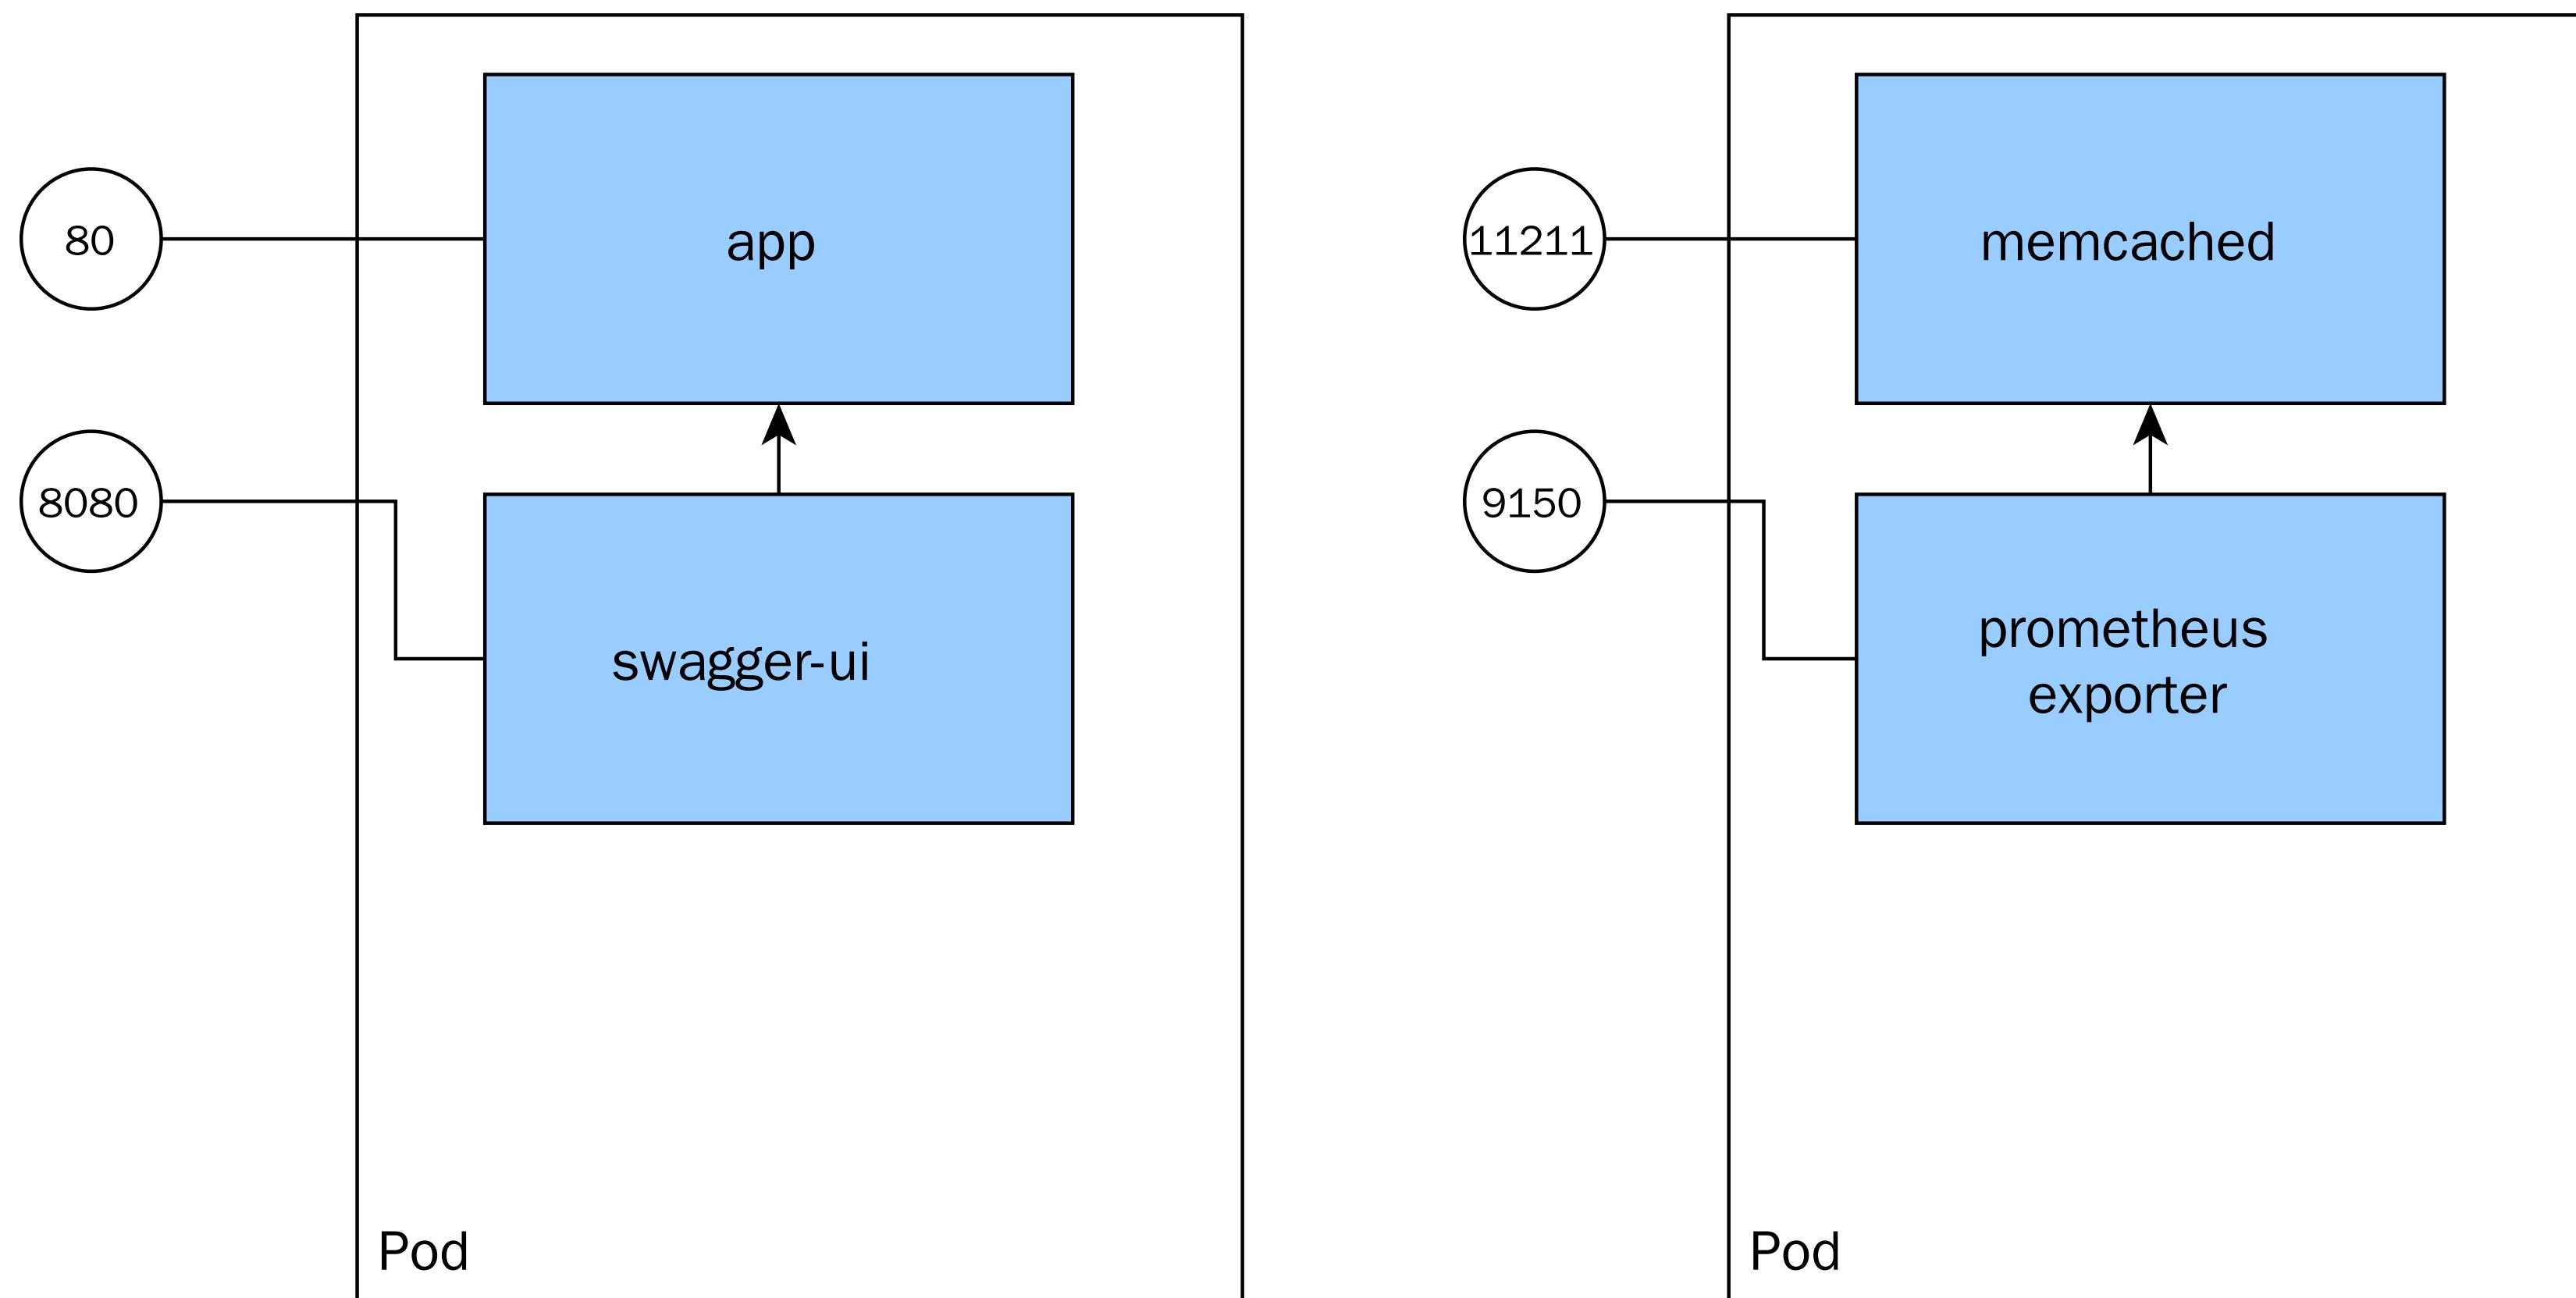
# KUBERNETES CONCEPTS

# PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



# SIDE-CARS



## BASIC CONCEPTS

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running