

EFFECTIVE PLATFORM BUILDING WITH KUBERNETES.

Wojciech Barczyński - [SMACC.io](#) | [Hypatos.ai](#)
16 November 2018

WOJCIECH BARCZYŃSKI

- Lead Software Developer
& System Engineer
- K8S:
2.5 years @ startups
- Before:
Openstack, SAP R&D
- + Visiting lecturer and Trainer



Github: [wojciech12](#) | Linked: [IN](#) | HP: [wbarczynski.pl](#) | T: [@wbarczynski](#)

STORY

- Lyke - [12.2016 - 07.2017]
- SMACC - [10.2017 - present]

My Role: Leading the change, implementing it

WHY WE LIKE KUBERNETES?



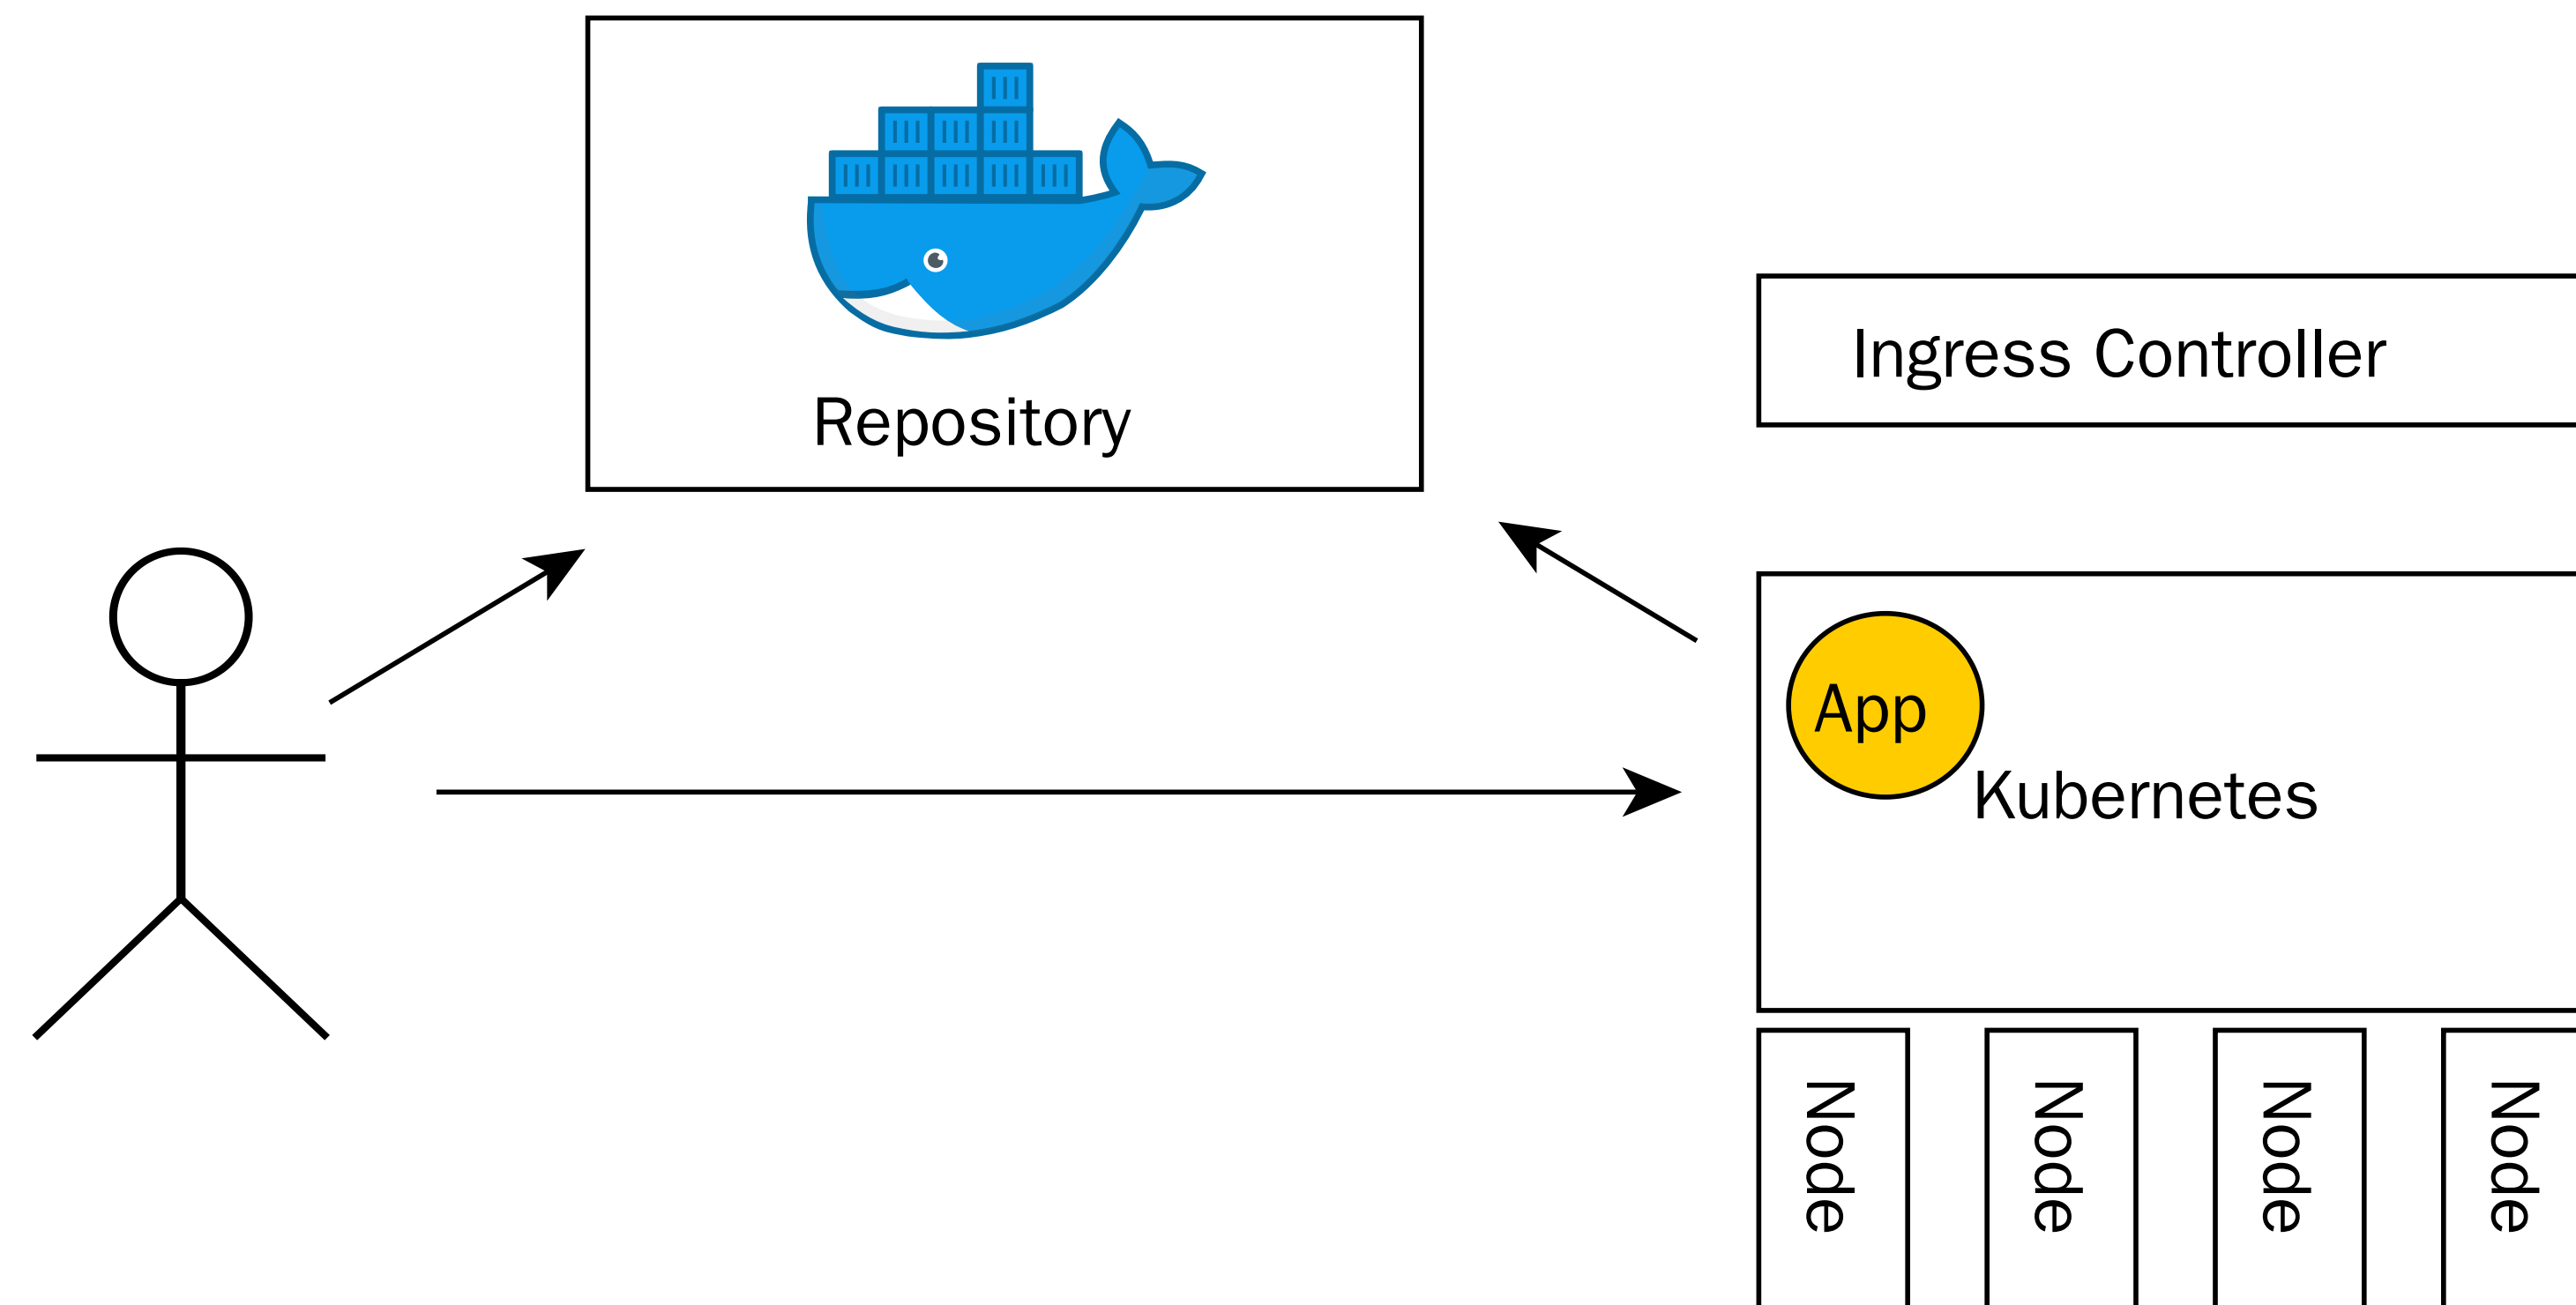
FOR ME

- Data center as a black box
- Lingua Franca
- Batteries included
- Learn-as-you-go experience
- Independent from IaaS provider

Notice: not a silver bullet

BLACK BOX

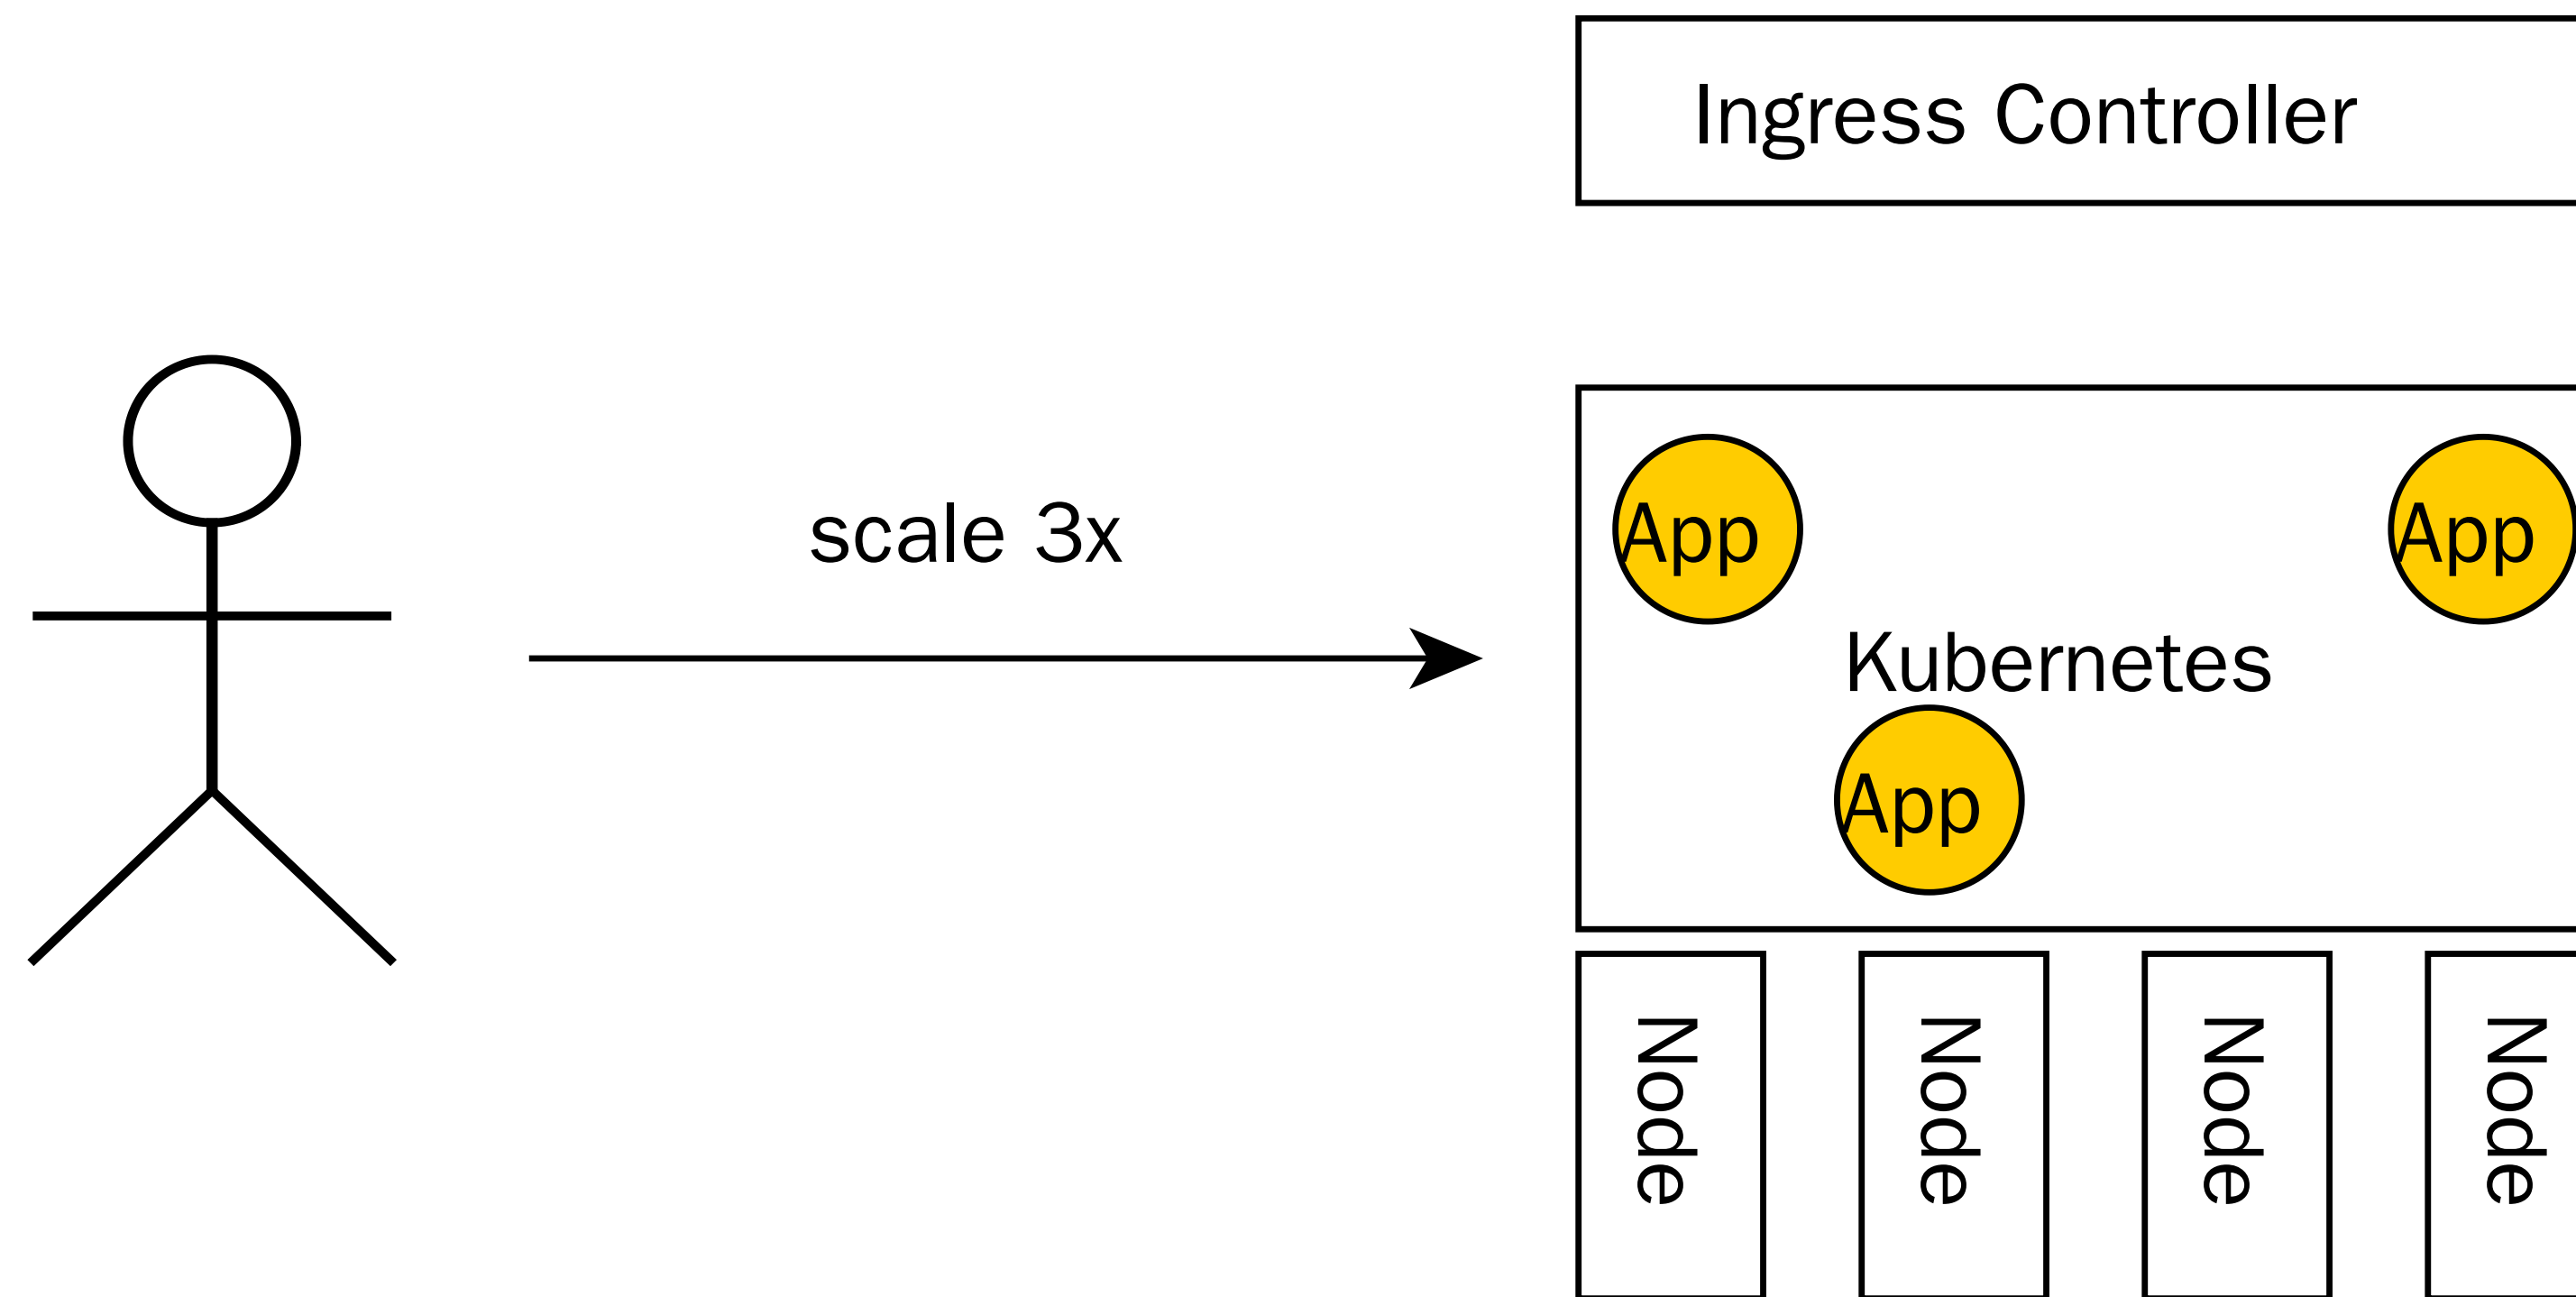
Deploy!



```
make docker_push; kubectl create -f app-srv-dpl.yaml
```

BLACK BOX

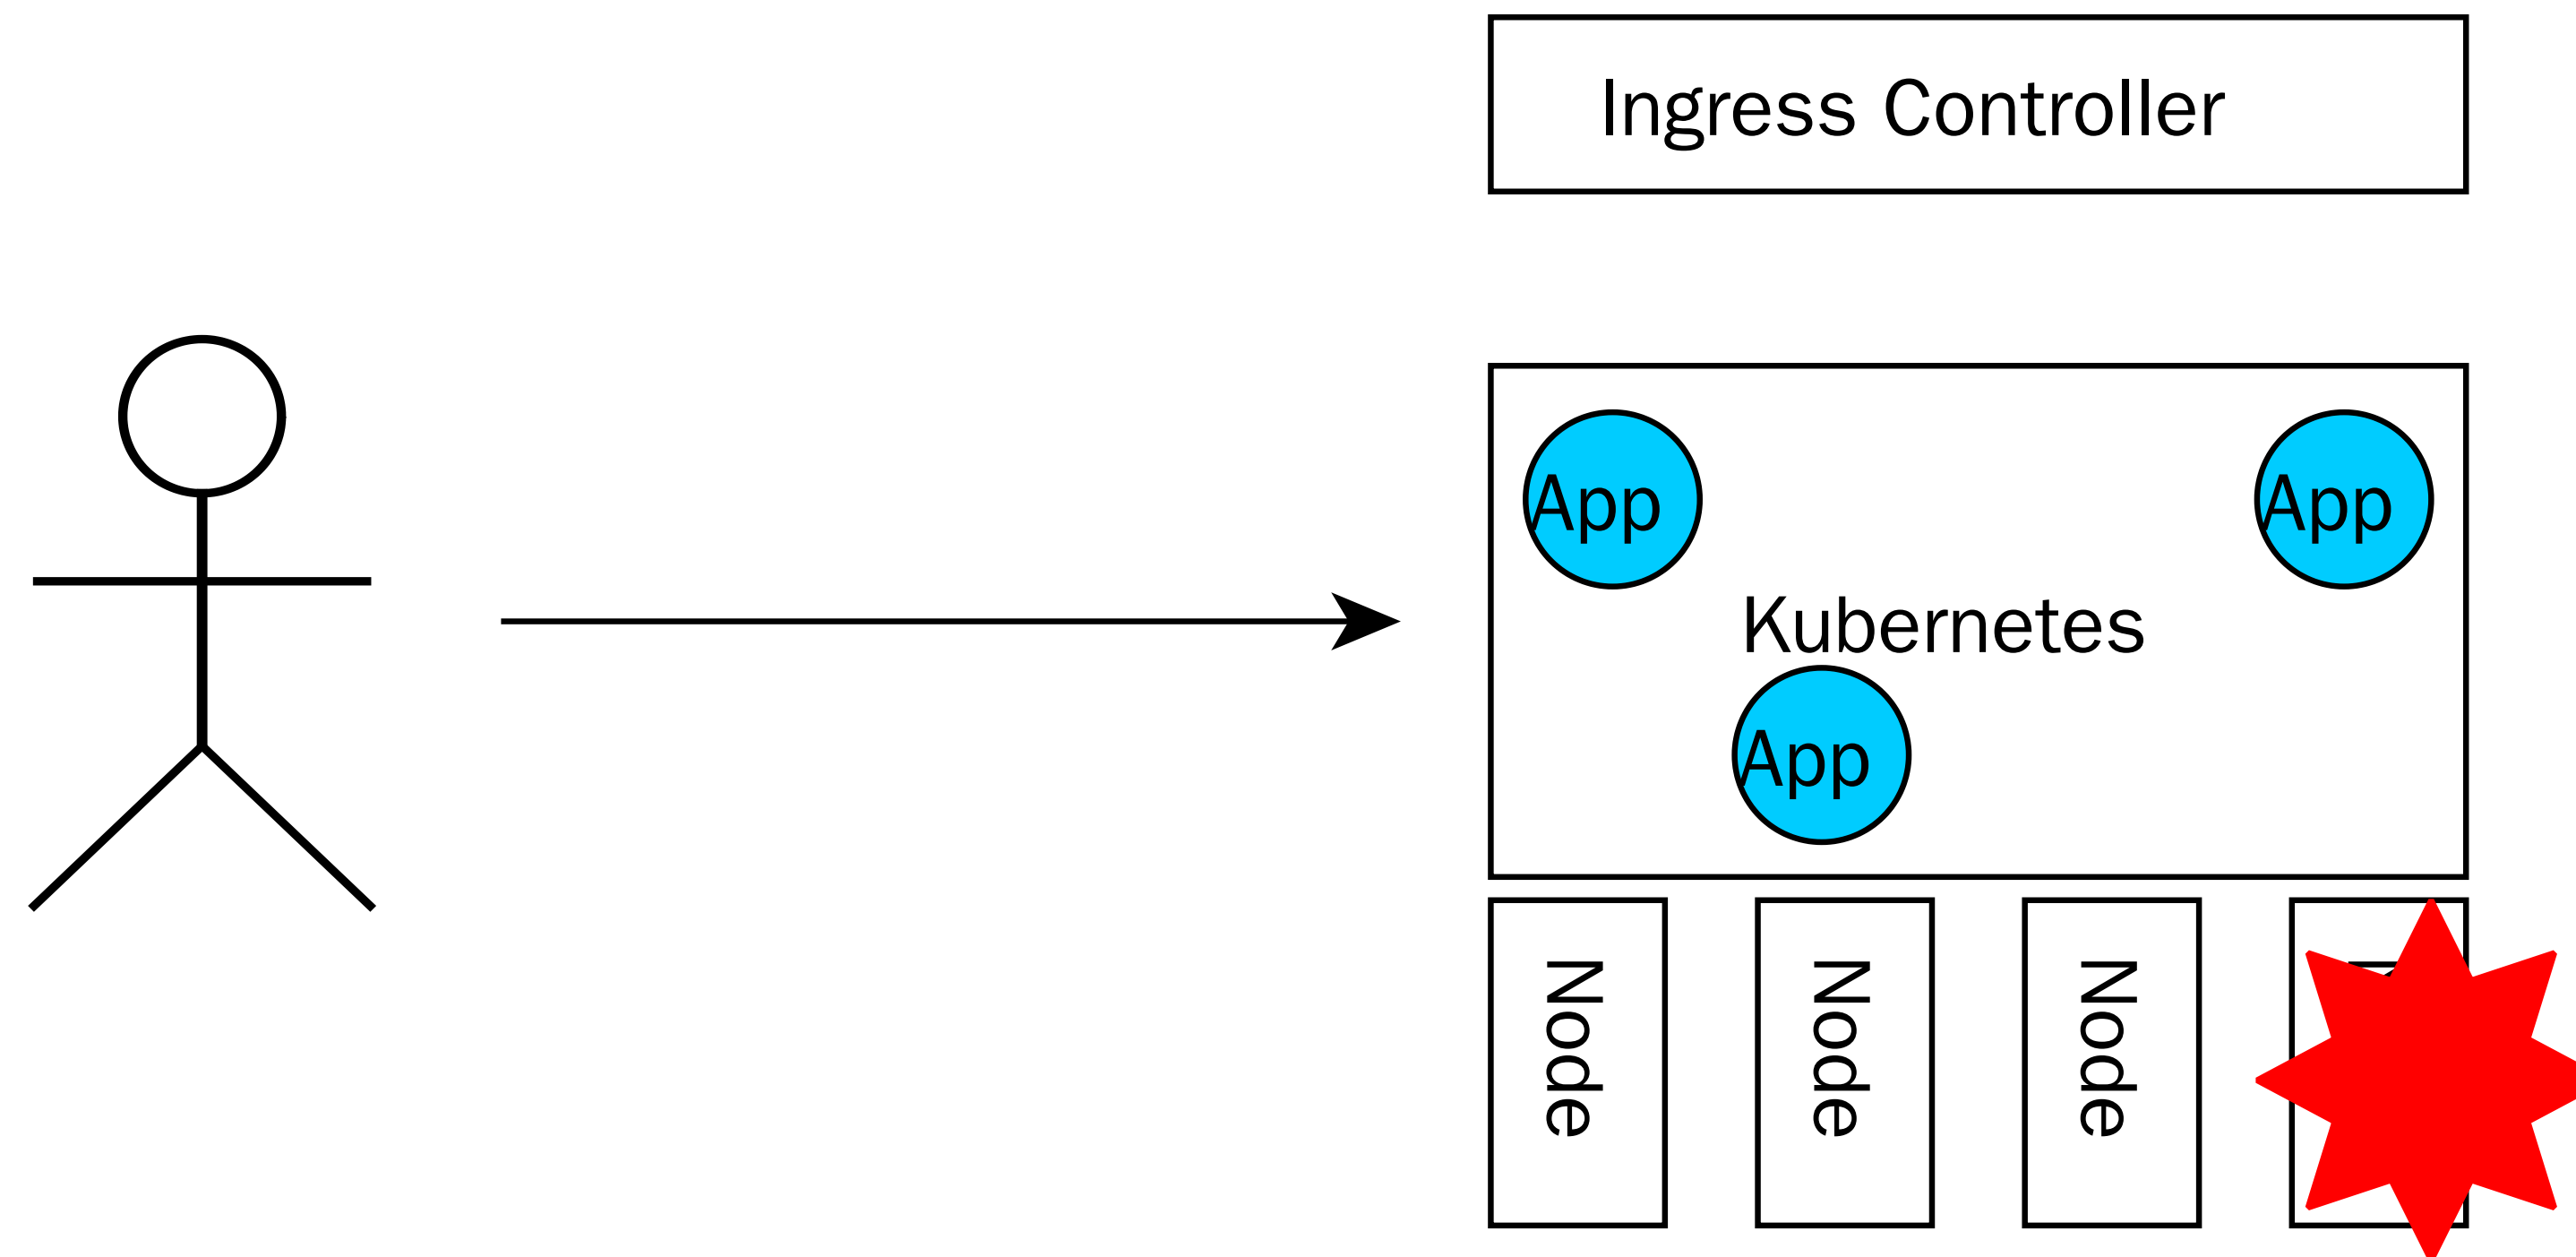
Scale up! Scale down!



```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

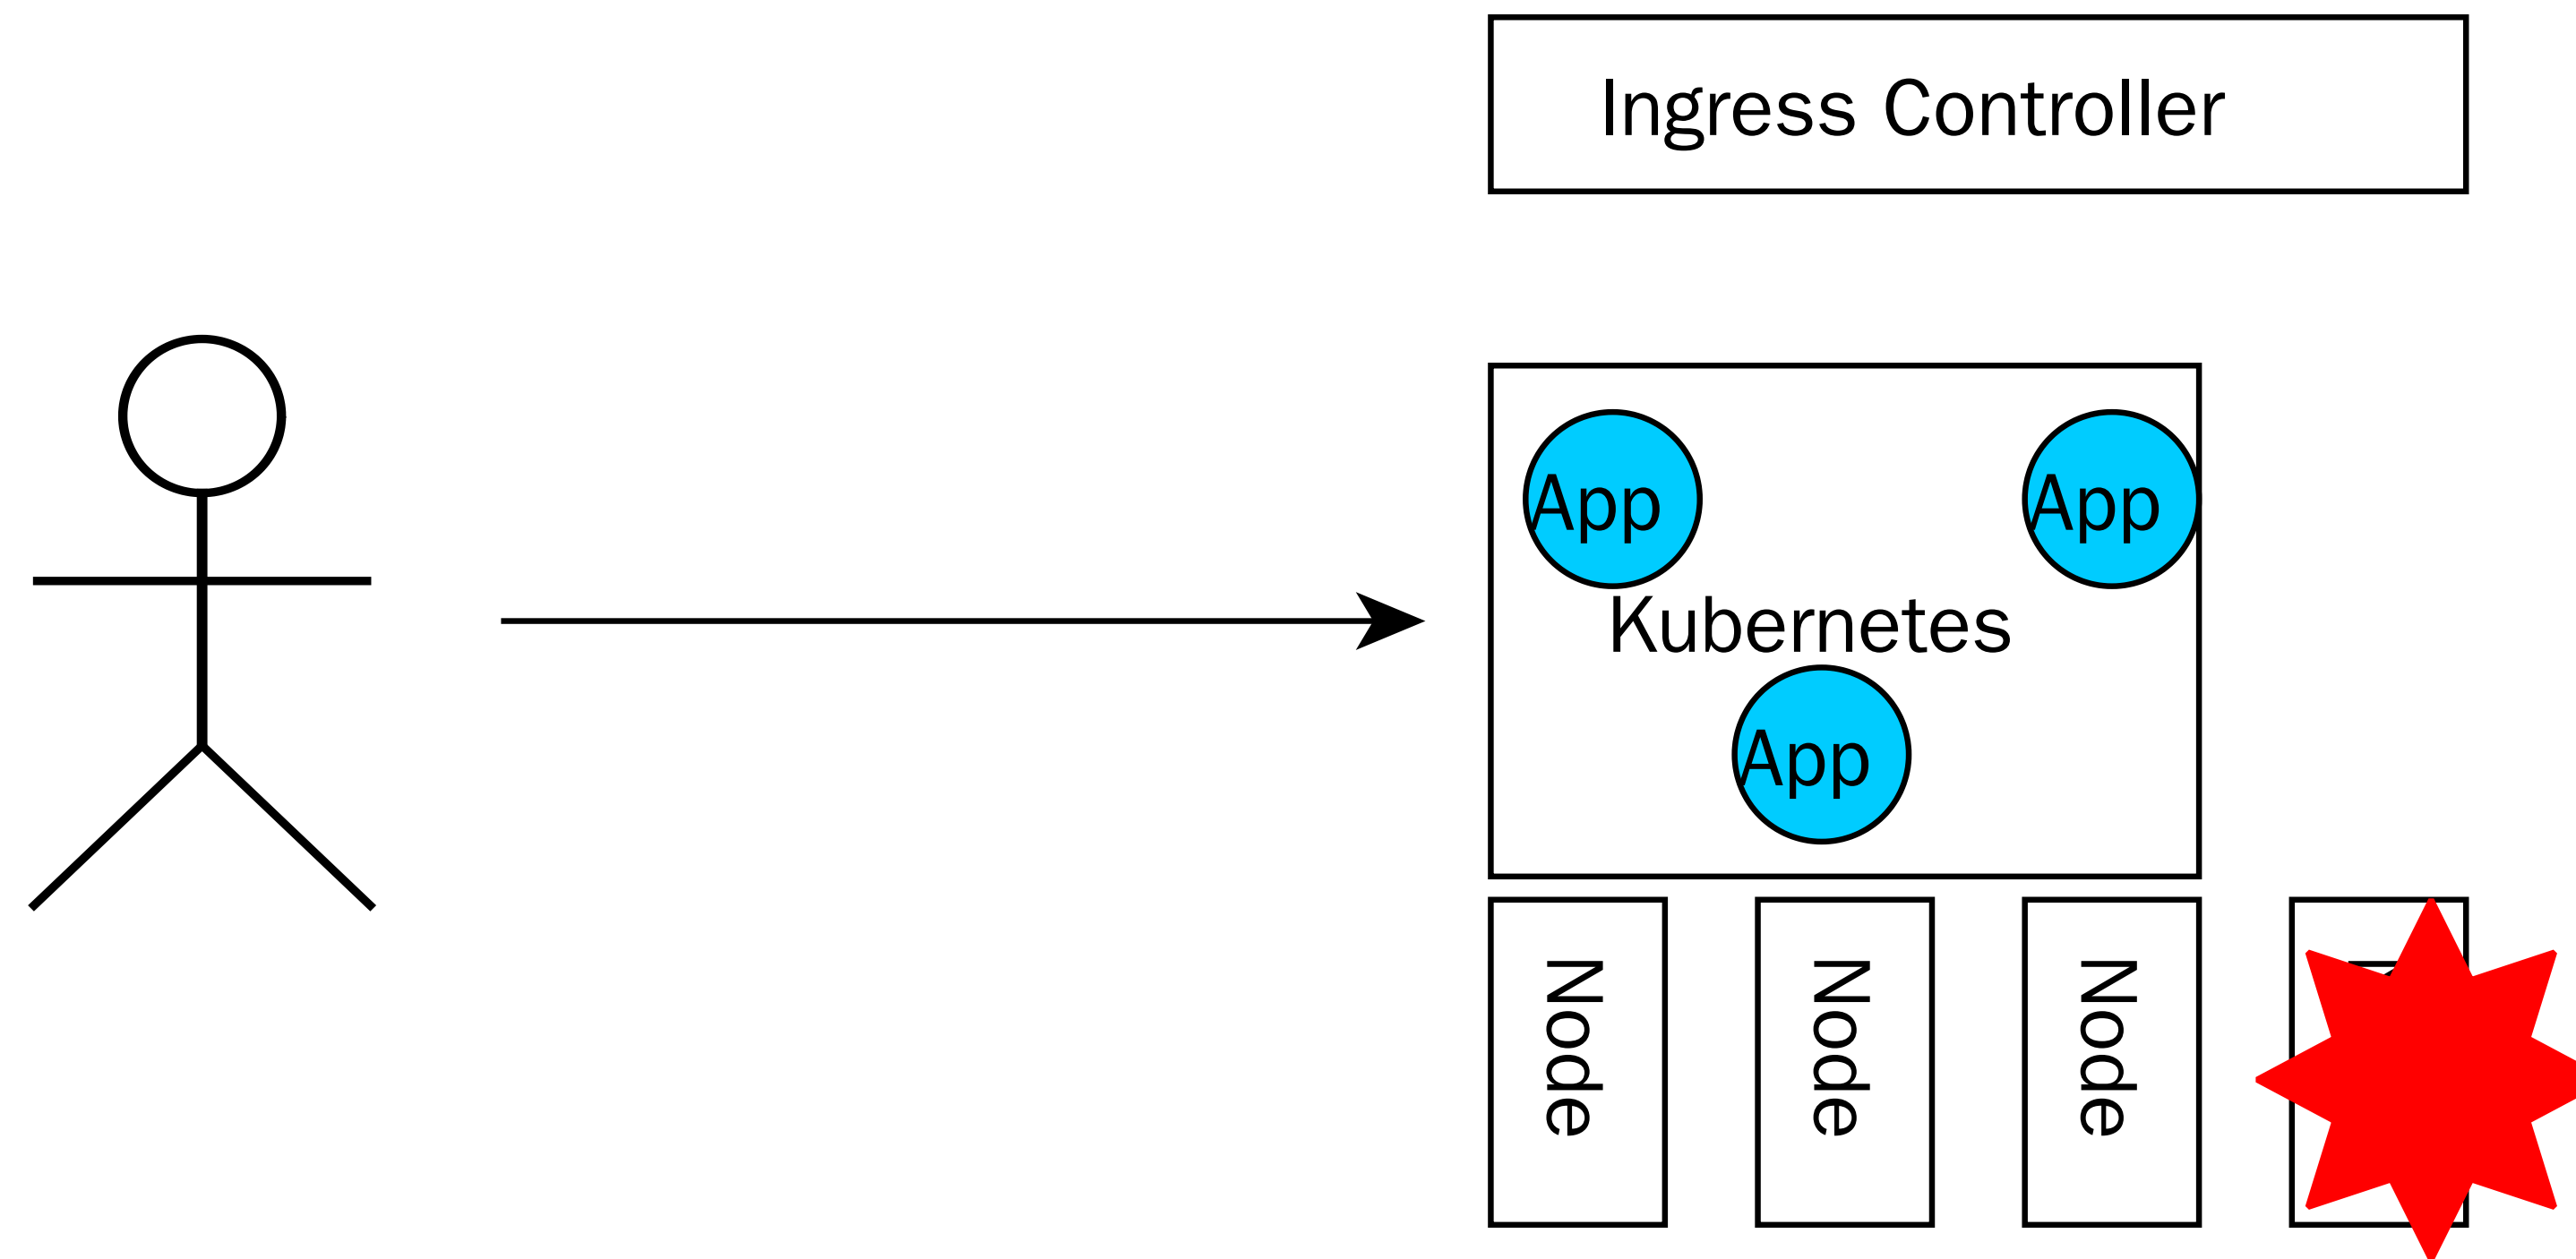
BLACK BOX

Resistance and Migrations



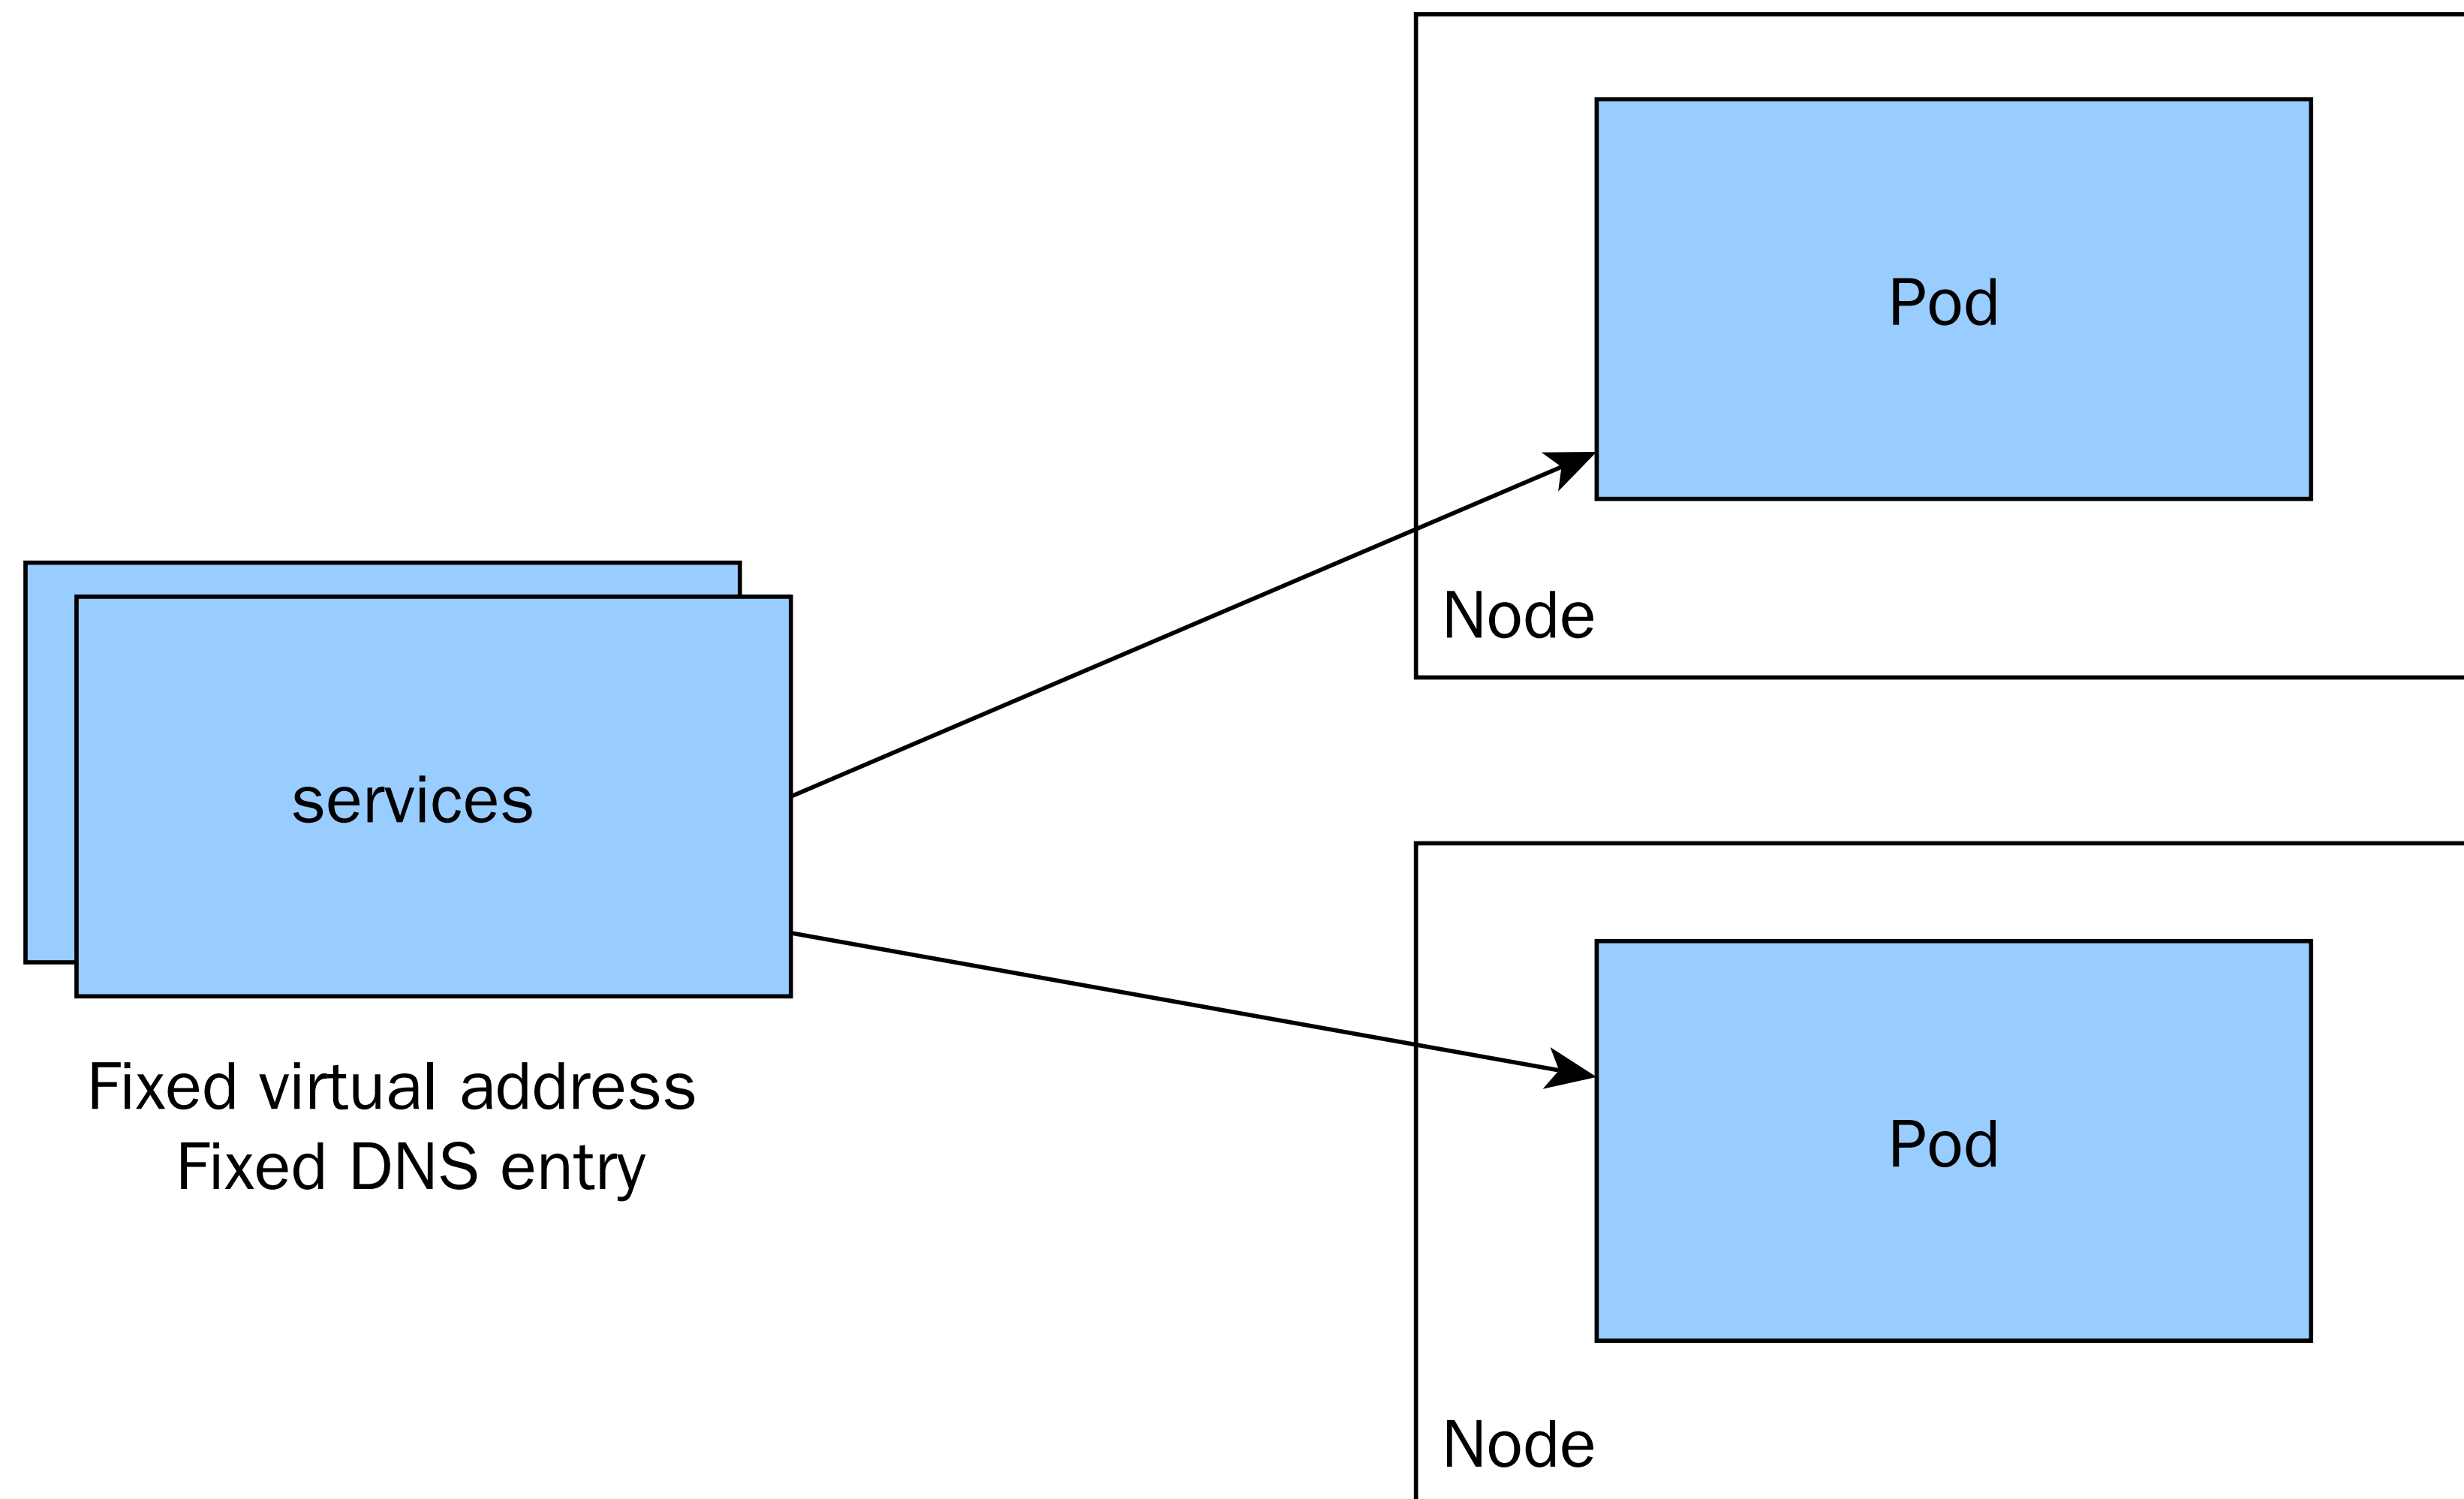
BLACK BOX

Resistance and Migrations

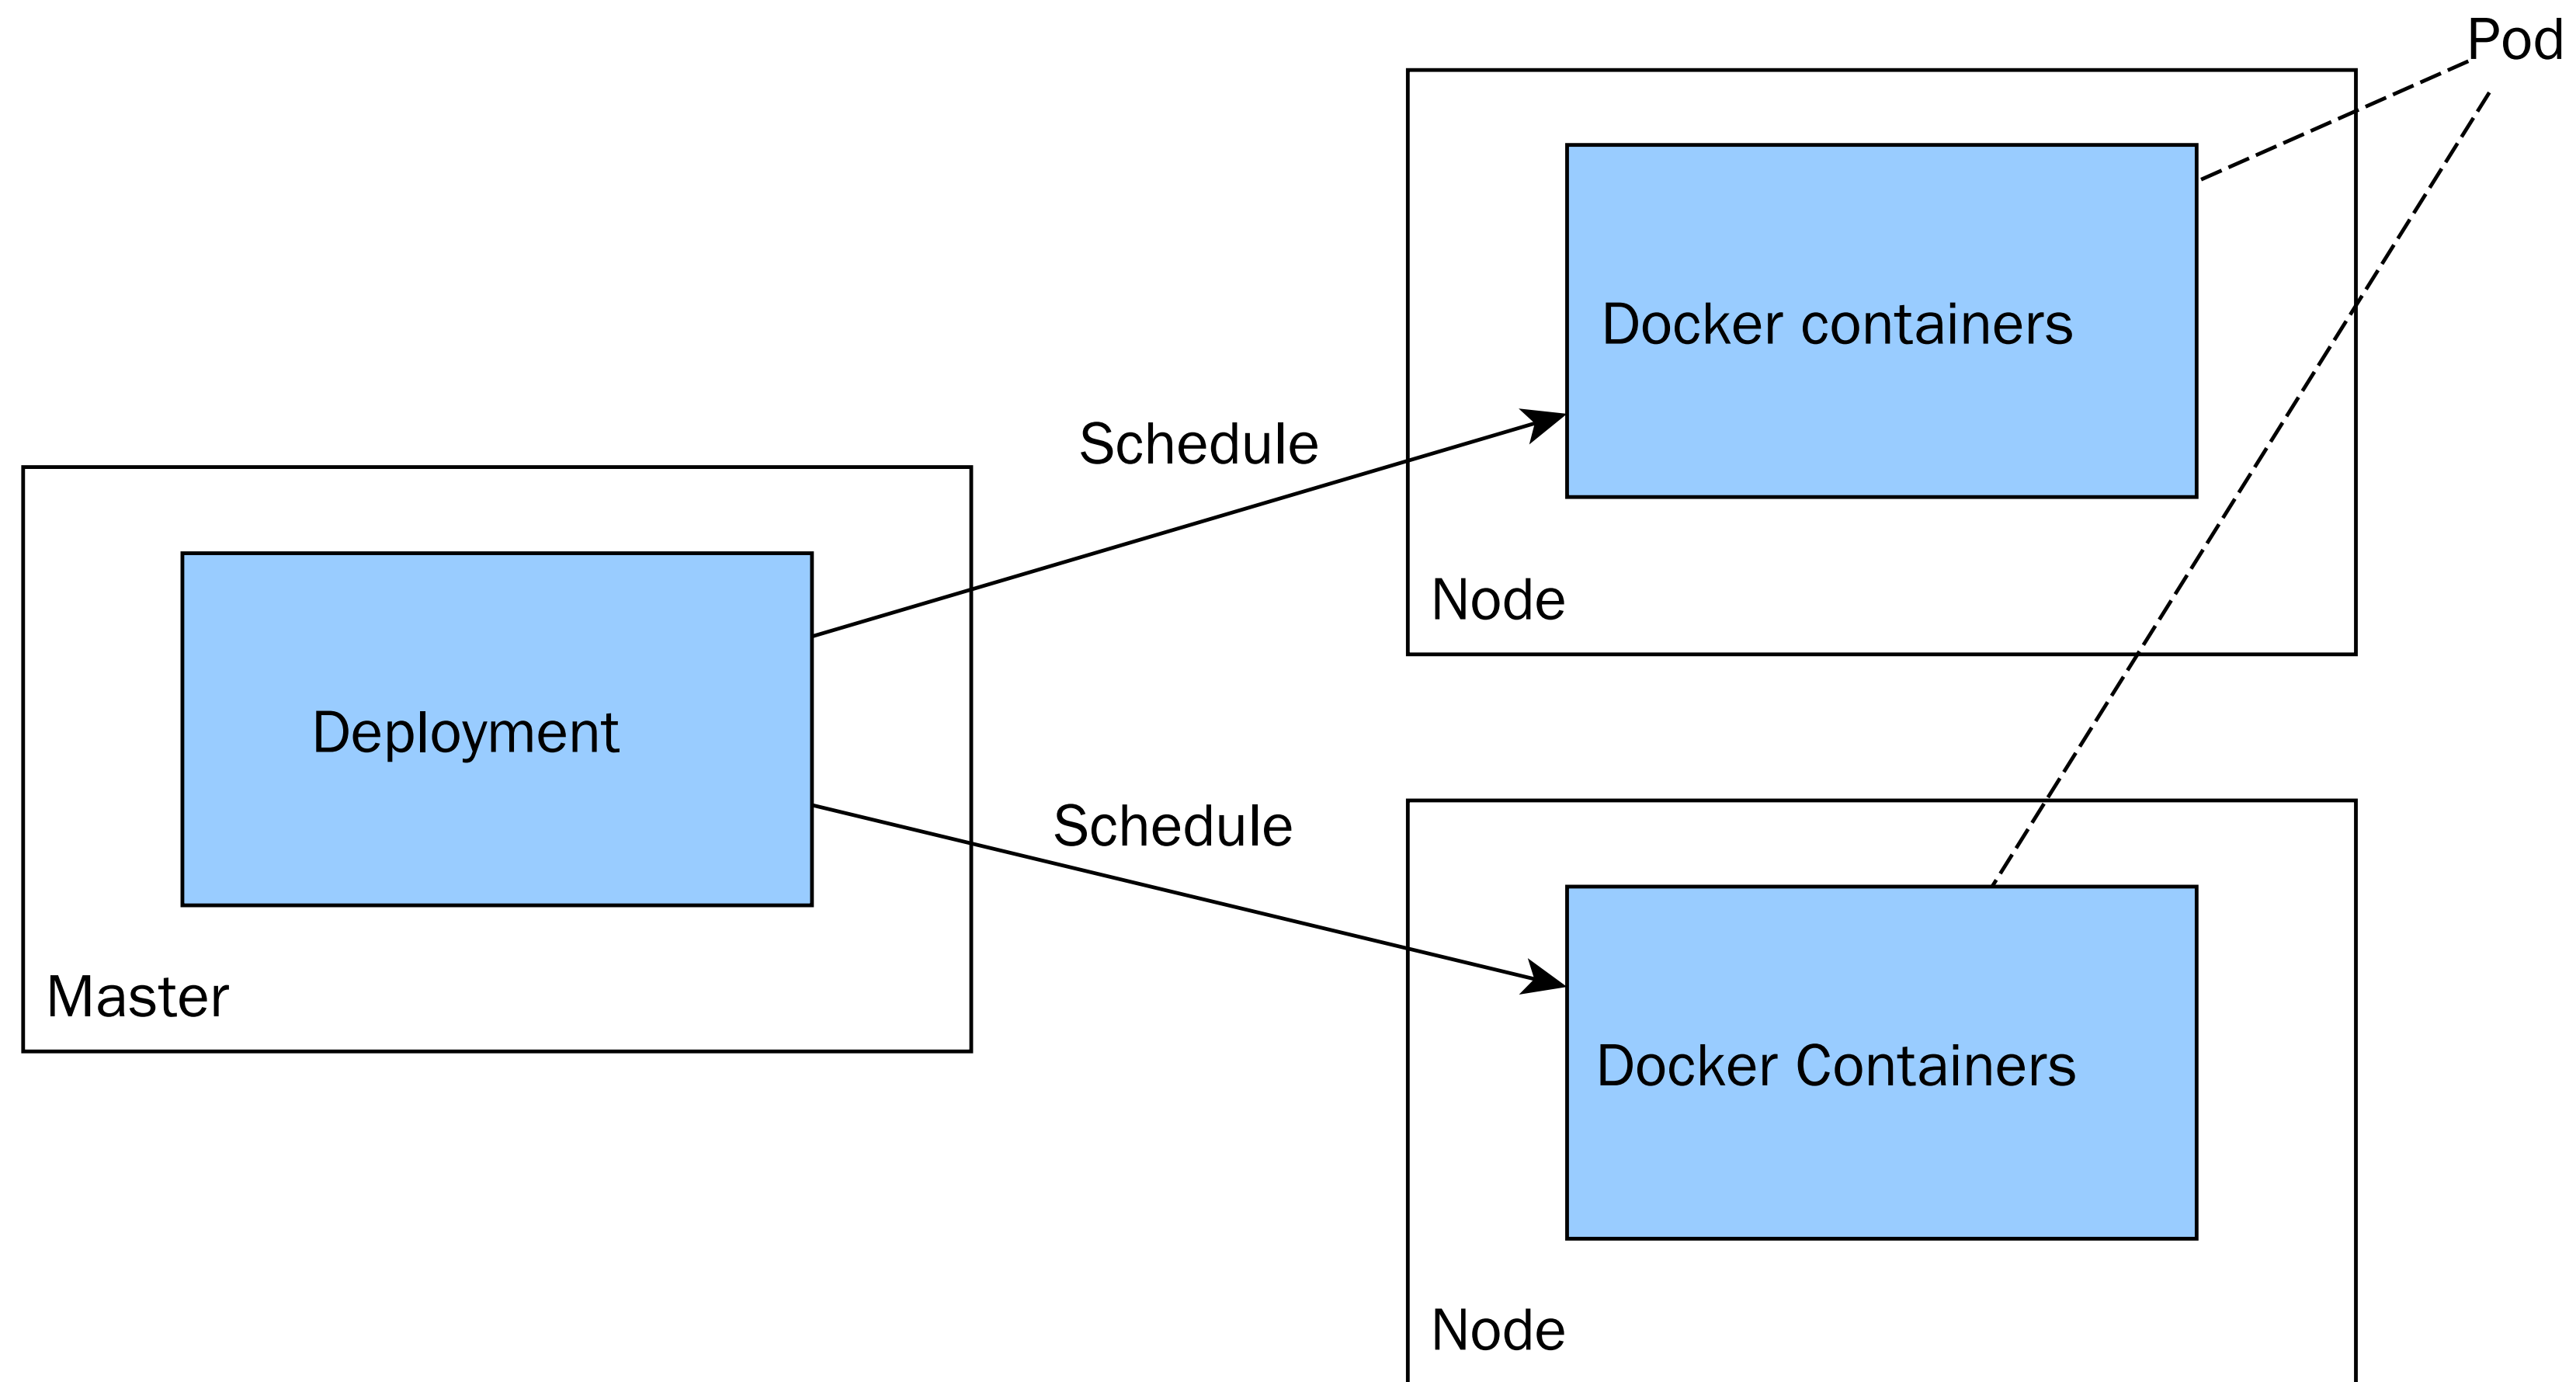


and much more: load balancing

LINGUA FRANCA



LINGUA FRANCA

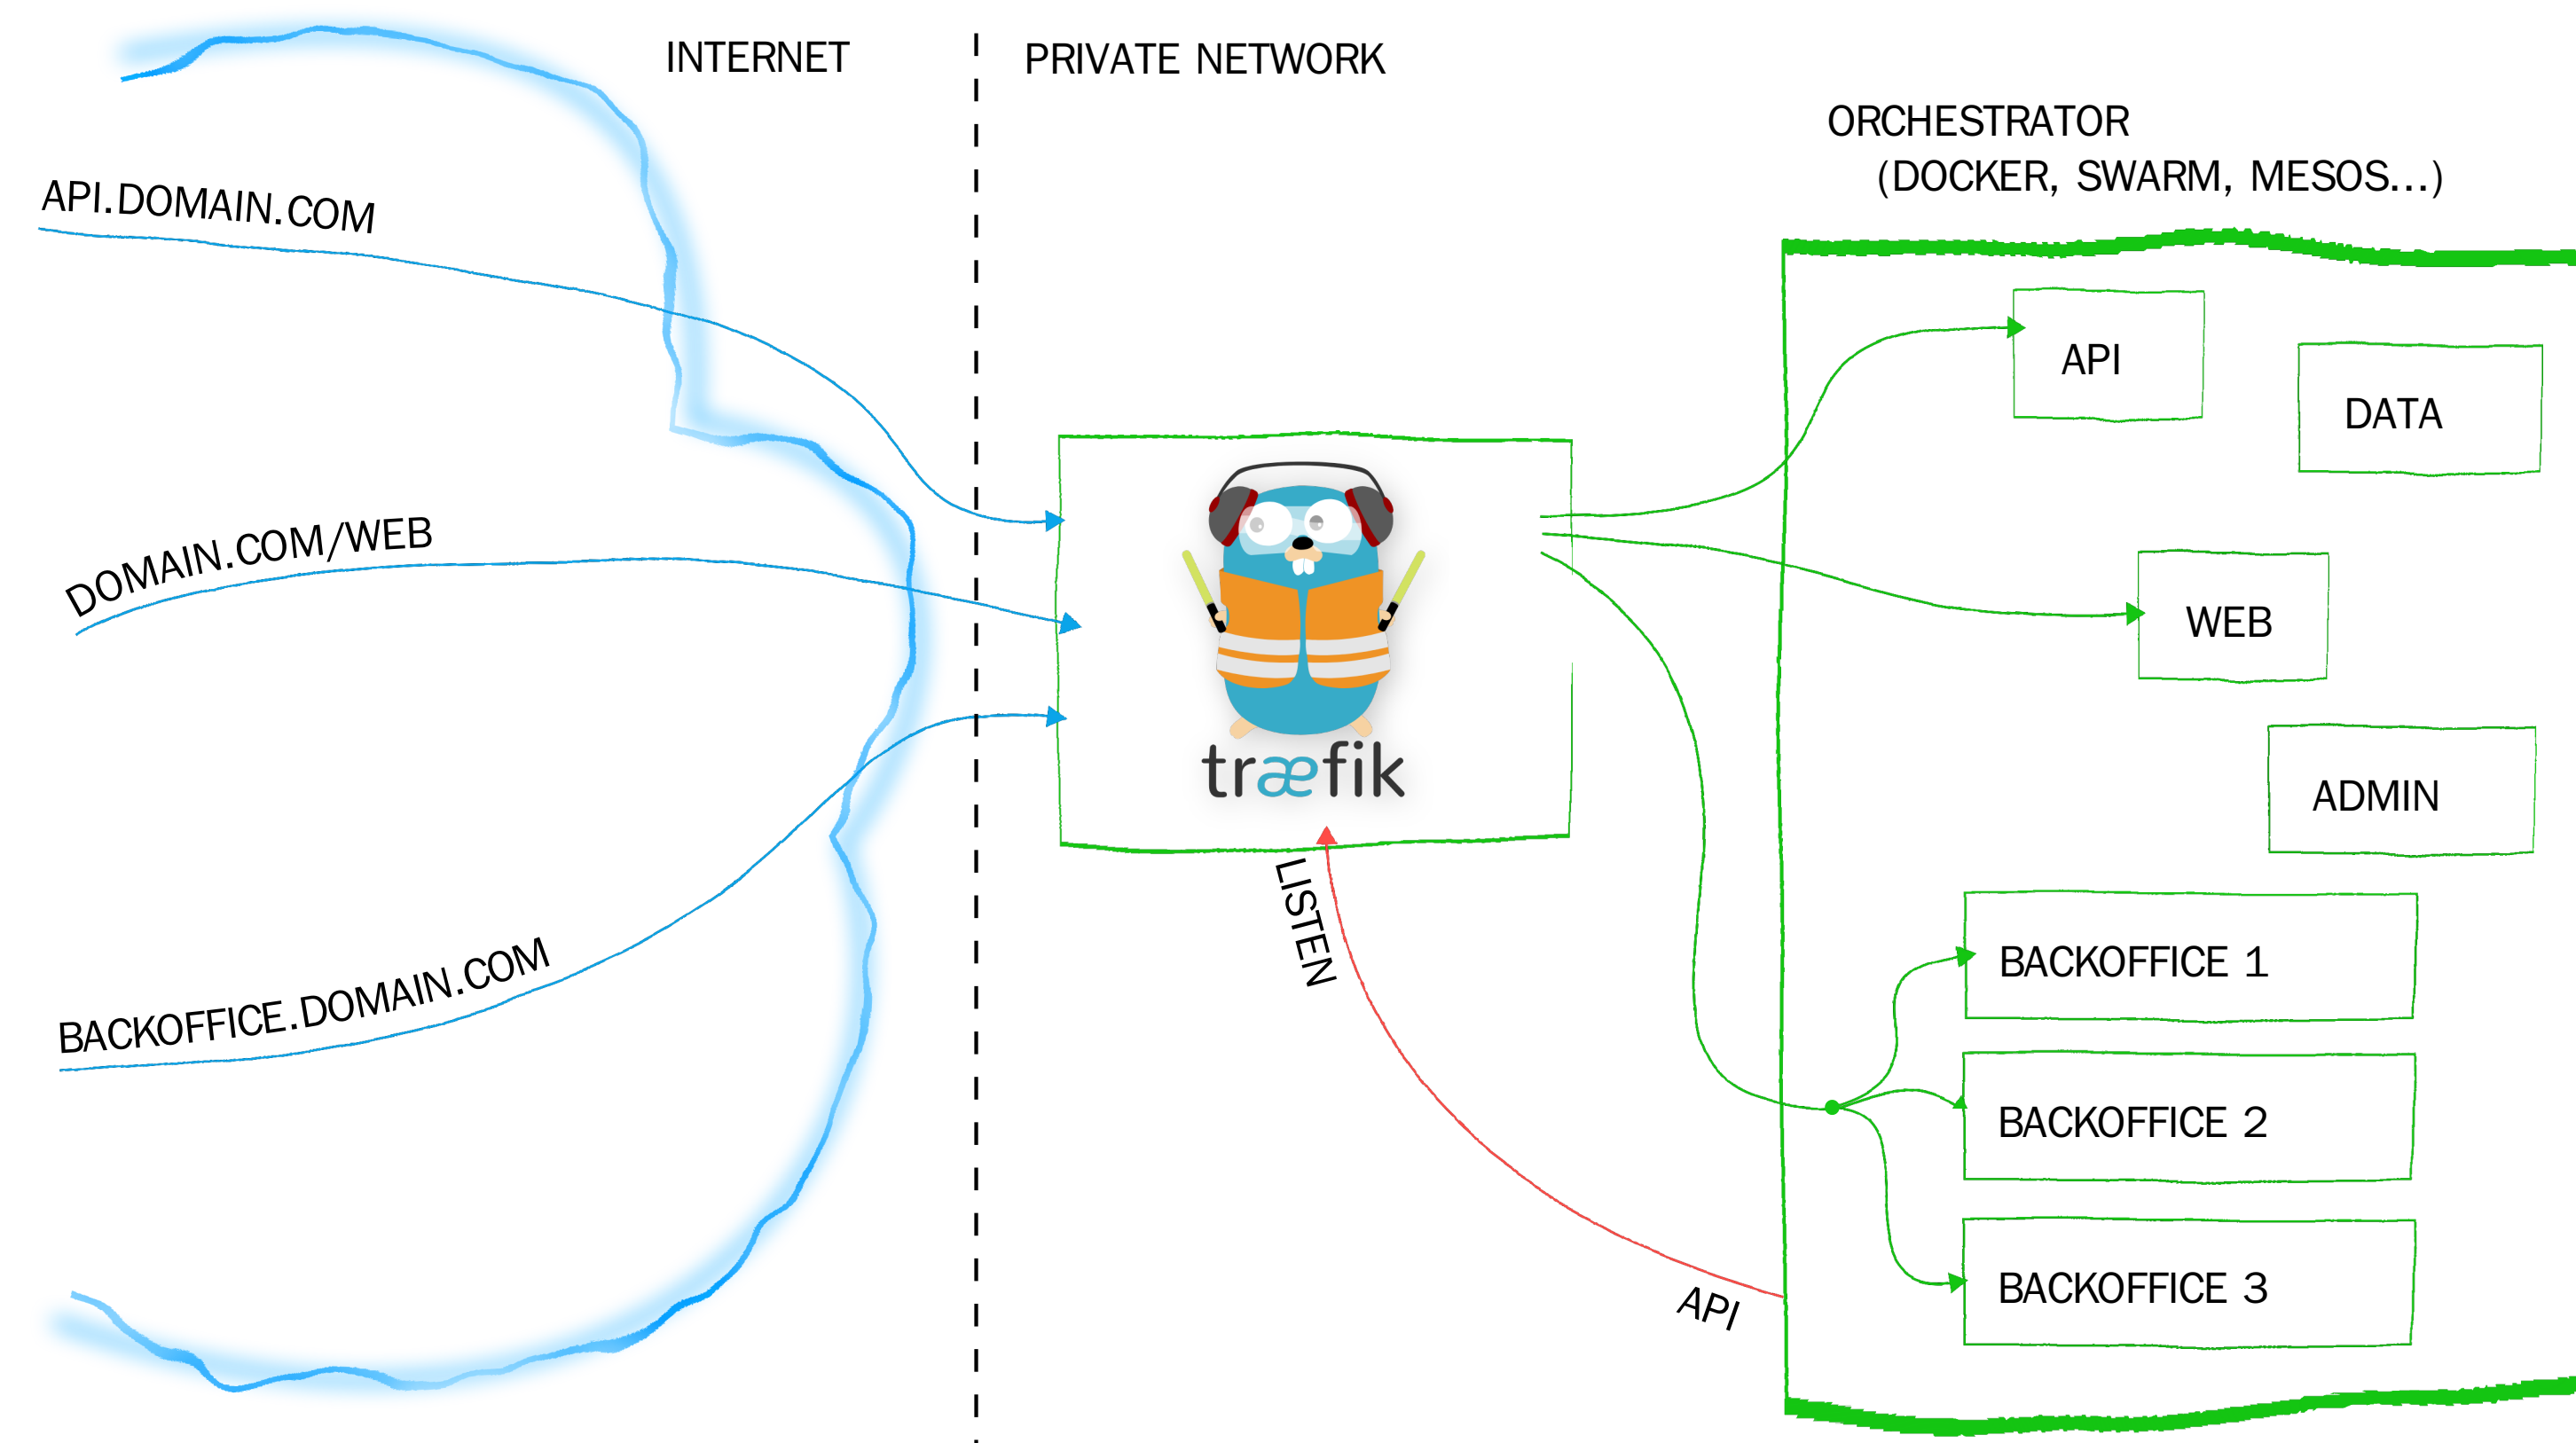


LIGUA FRANCA

Pattern	Target App Service
api.smaacc.io/v1/users	users-v1
api.smaacc.io/v2/users	users-v2
smaacc.io	web

LINGUA FRANCA

Ingress Controller:



Drop-in: traefik, nginx, haproxy, envoy...

BATTERIES

Service Discovery:

- service injected to DNS:

```
curl http://users/list
```

- labels:

```
name=value
```

- annotations:

```
prometheus.io/scrape: "true"
```

SERVICE DISCOVERY

- Loosely couple components
- Auto-wiring with logging and monitoring
- Integration

Learn-as-you-go experience

Monitoring? Ingress?

1. Drop-in: traefik, nginx, ...
2. oooo ruuuuns!
3. Make a hell of mistakes
4. Get the right one or It is OK for now

Learn-as-you-go experience

Where to start?

➡ With service and deployment

Learn-as-you-go experience

memcached in k8s?

➡ statefulset

Learn-as-you-go experience

My pod takes all memory

➡ resource & limits

My ML-based app has slow start.

Users noticed downtime.

LONG LIVE GIT!

- All yaml in your service git
- Integration with e.g.:
 - monitoring, alarming
 - ingress-controller
- Forget about infrastructure... almost


DevOps Culture Dream!



LYKE

LYKE

- E-commerce Indonesia
- Mobile-only
- 50k+ users
- 2M downloads
- Top 10 Fashion Apps
Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,
UI Designer,
Mobile Dev,
and tester - one body



CHALLENGES

- 50+ VMs in Amazon, 1 VM - 1 App
- 65% Idle machine \$\$\$
- Puppet with manual deployment process
- Fear, Forgotten components
- Performance issues

APPROACH

1. Simplify Infrastructure
2. Change the Development Practices
(12factor + kubernetes)
3. Change the Work Organization

see: Conway's law

SIMPLIFY

1. Kubernetes with Google Kubernetes Engine
2. Terraform for all new
3. Database-as-a-service

SIMPLIFY

1. Prometheus, AlertManager, and Grafana
2. Elasticsearch-Fluentd-Kibana
3. Google Identity-Aware-Proxy
to protect all dev dashboards
4. 3rd party SaaS: statuscake and opsgenie

One person effort

CONTINUOUS DEPLOYMENT

- branch-based:
 - master
 - staging
 - production
- repo independent

TRAVIS CI

1. Tests
2. Build docker
3. Deploy to Google Container Registry
4. Deploy to k8s only new docker
5. no config/secrets applied

GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \- kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```

Makefile

```
SERVICE_NAME=v-connector
GCP_DOCKER_REGISTRY=eu.gcr.io
test: test_short test_integration

run_local:

docker_build: docker_push

kube_create_config:

kube_apply:

kube_deploy:
```

Copy&Paste from the project to project

1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

2. GET BACK ALL THE KNOWLEDGE

Extract from:

- Puppet, ... ➡ Dockerfile
- Running Instances ➡ Dockerfile, README.rst
- Nagios, ... ➡ README.rst + *checks/*

3. INTRODUCE RUN_LOCAL

- `make run_local`
- A nice section on how to run in README.rst
- with `docker-compose`

The most crucial point.

4. GET TO KUBERNETES

- `make kube_create_config`
- `make kube_apply`
- Generate the yaml files if your envs differ

secrets from gopass (password manager)

5. CONTINUOUS DEPLOYMENT

Travis:

- the same Makefile as dev use

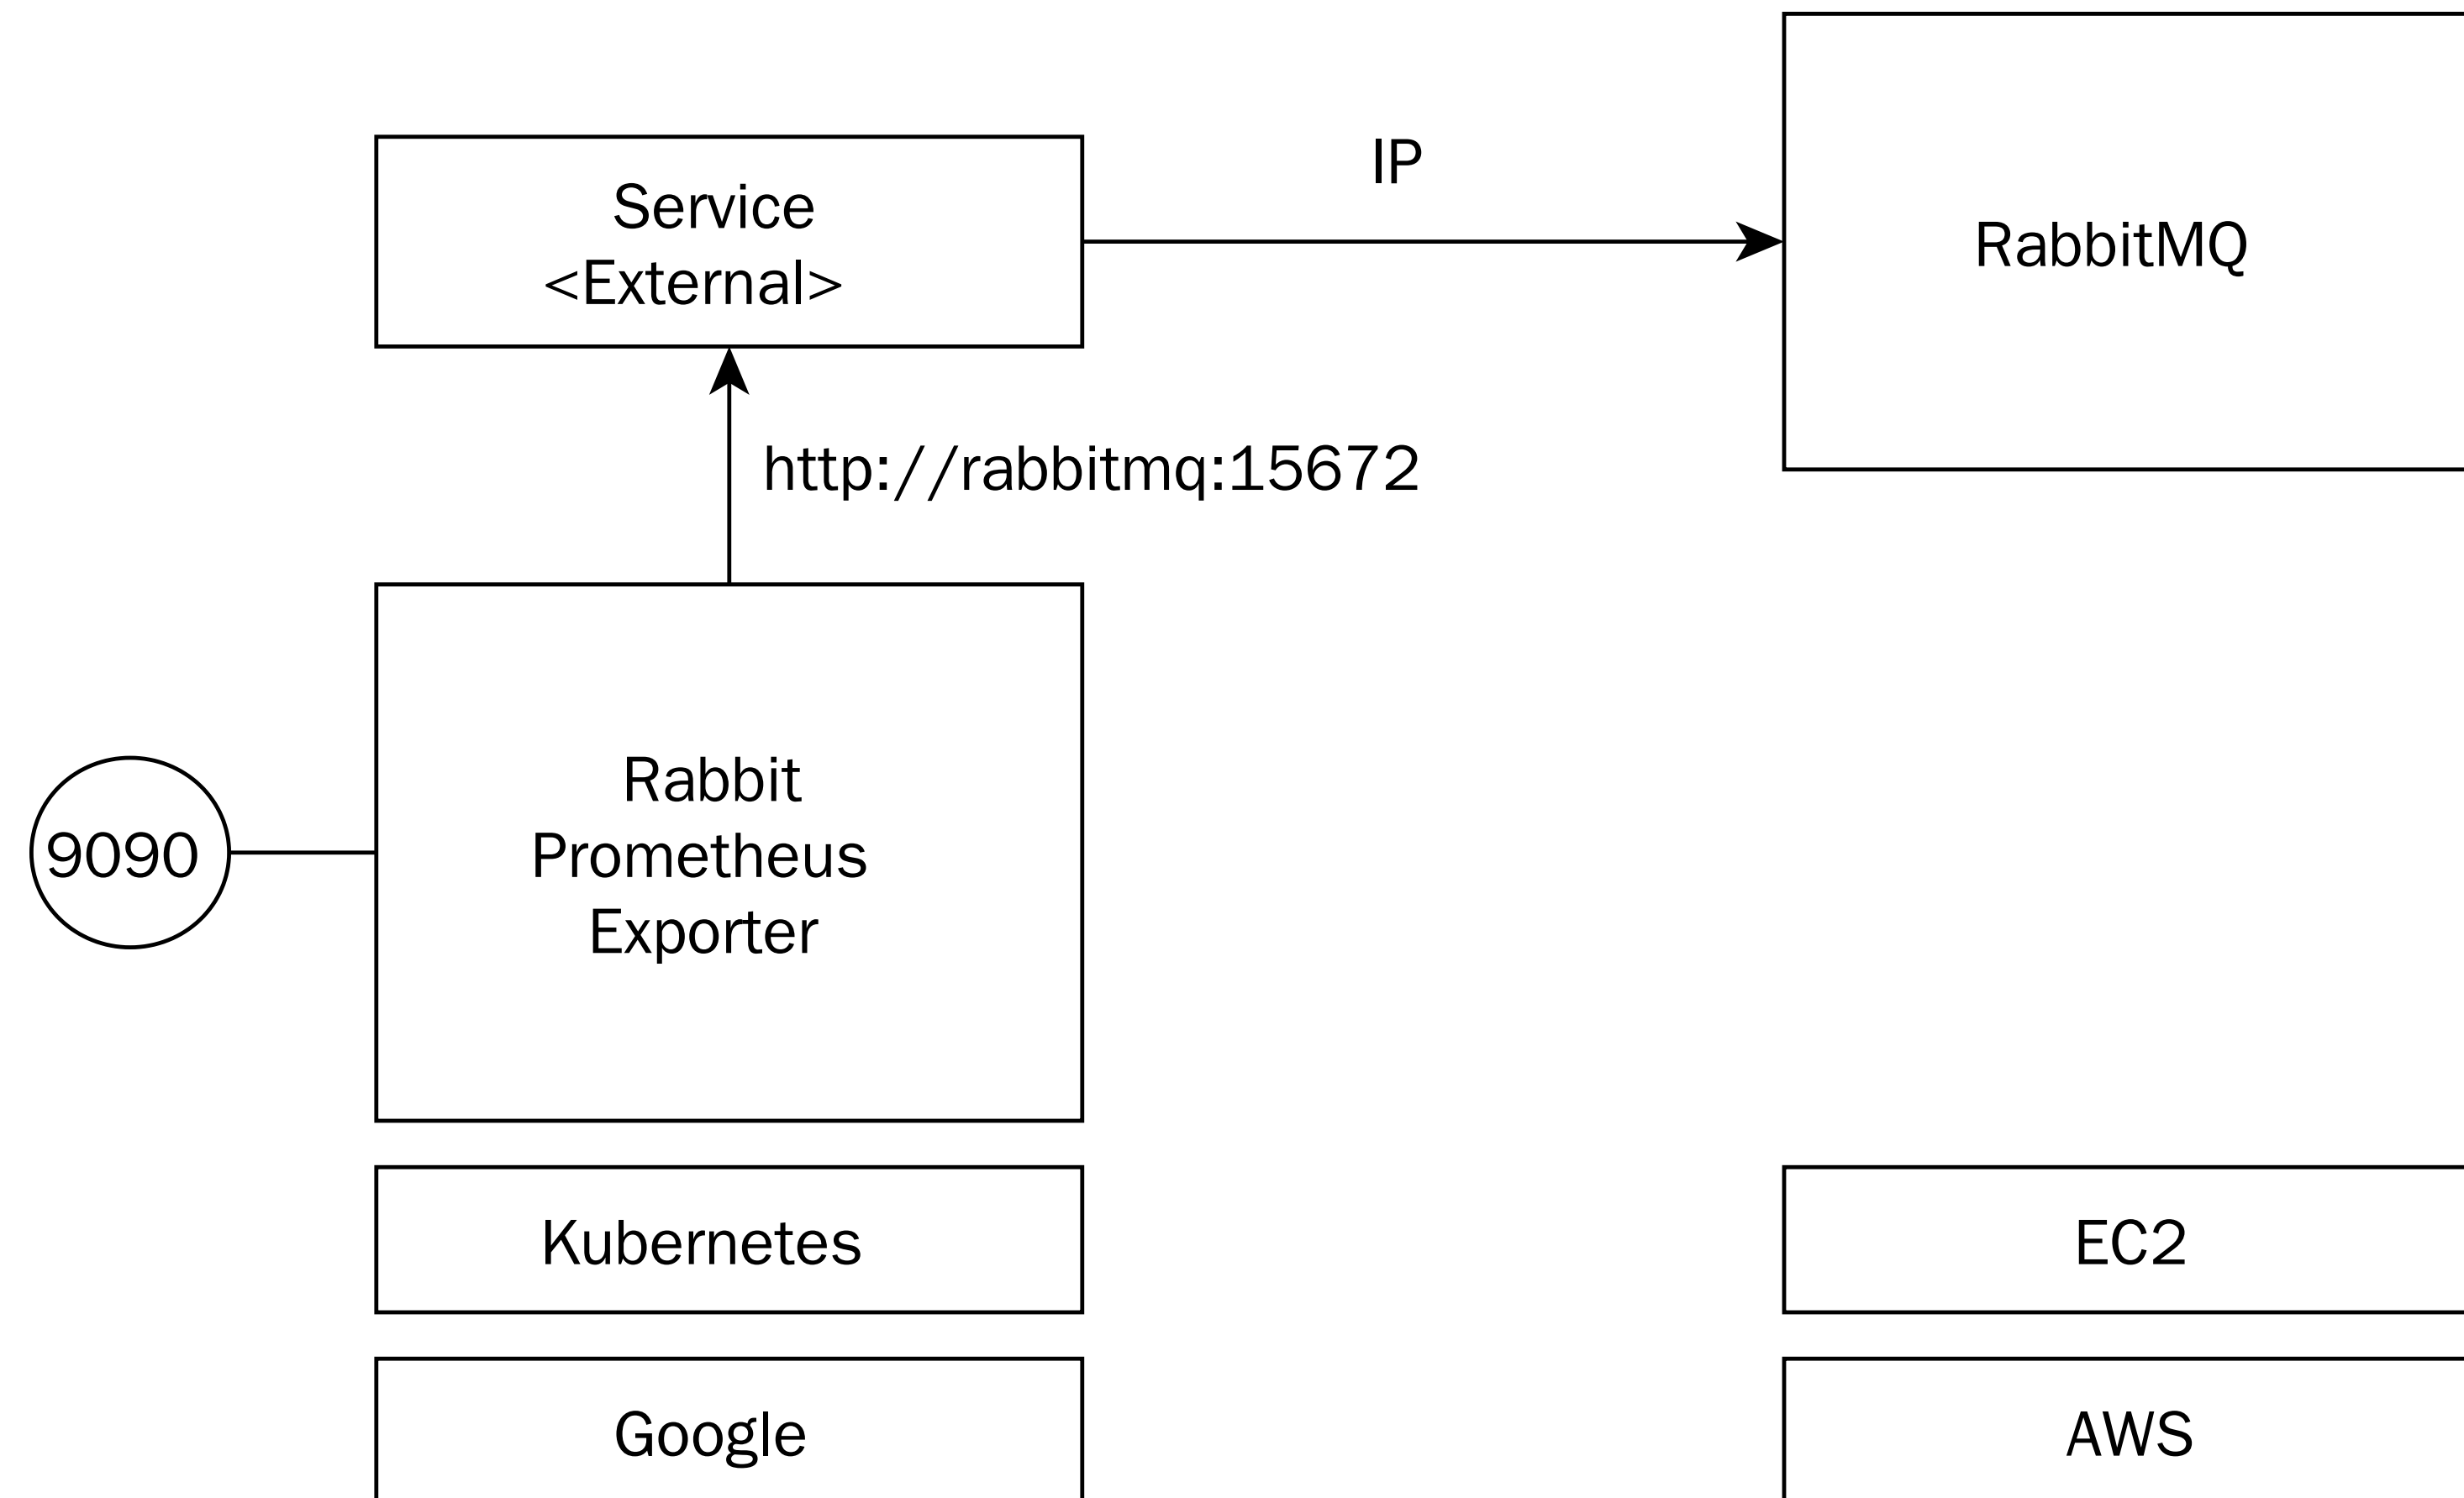
6. KEEP IT RUNNING

Bridge the new with old:

- Use External Services in Kubernetes
- Expose k8s in the Legacy [1]

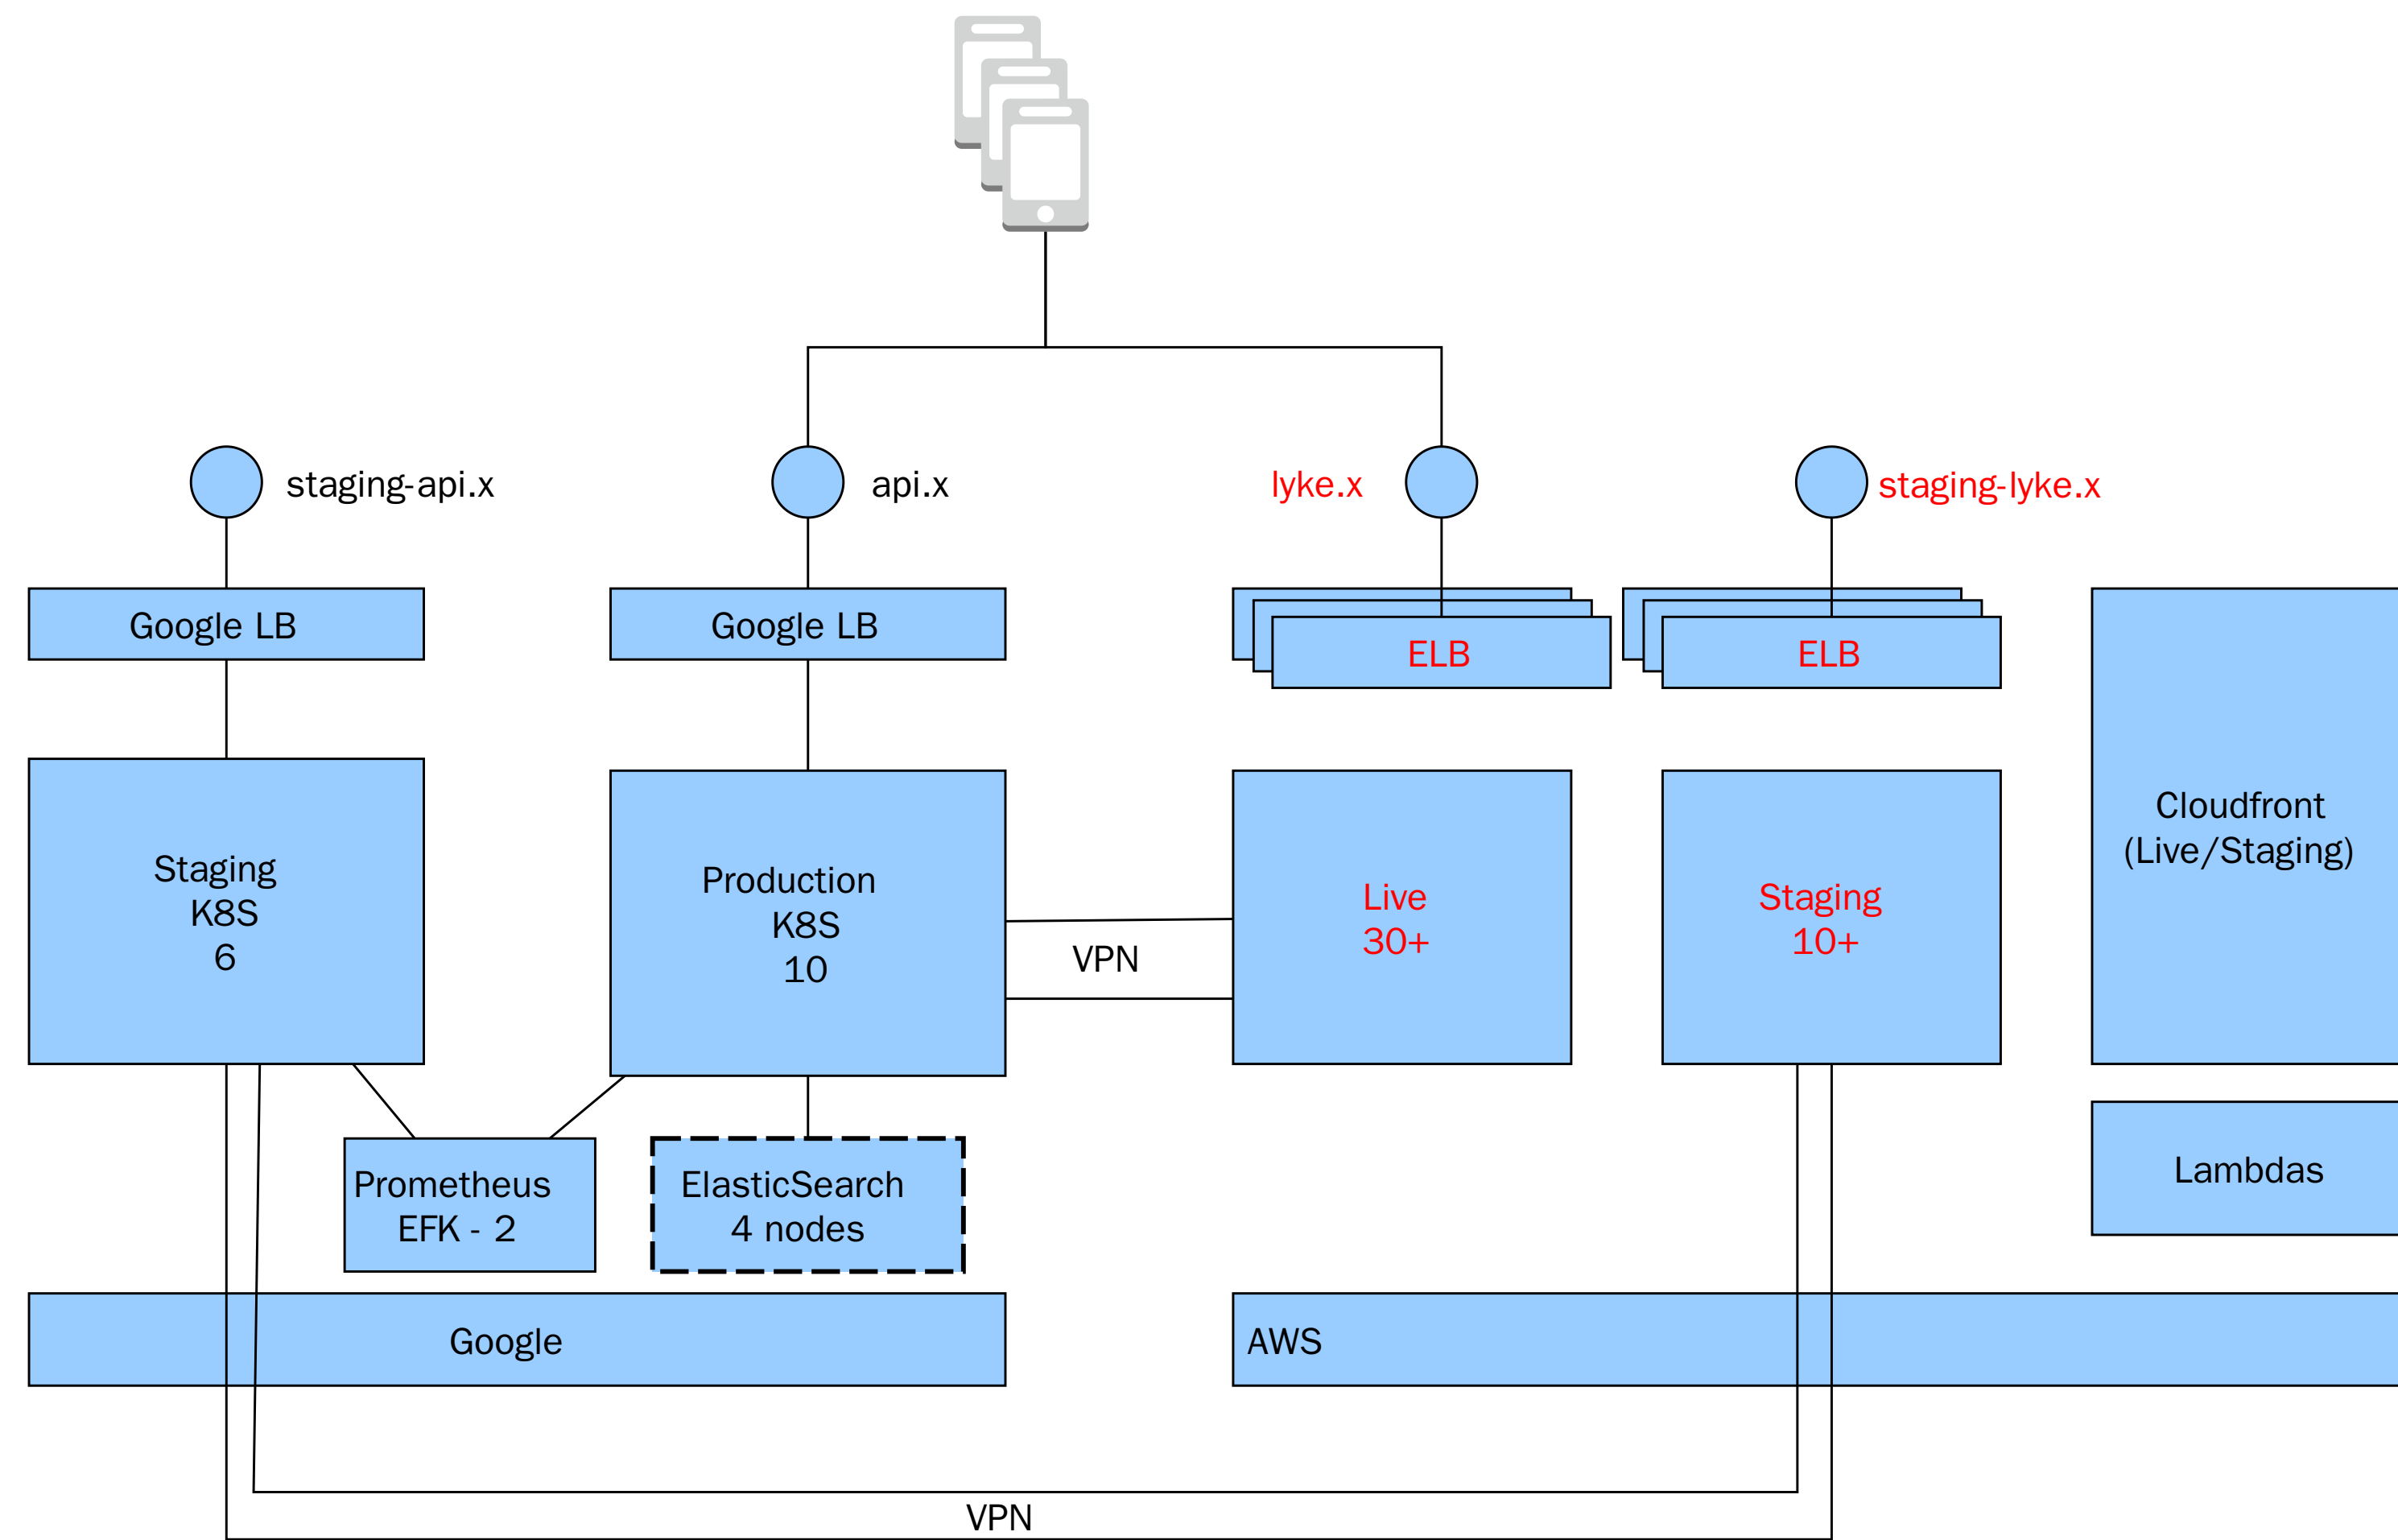
[1] hard coded IP:PORT, considered: K8S events to Consul

Bridge the new with old



Monitor legacy with new stack

Architecture During Migration



7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test  
TARGET_URL=api.example.com/users make smoke_test
```


8. SERVICE SELF-AWARE

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

9. MOVE TO MICRO-SERVICES

Offload Legacy:

- Keep the lights on
- New functionality to micro-services

WHAT WORKED

1. Copy&Paste Makefile and k8s yaml
2. Separate deployments a good transition strategy

WHAT DID NOT WORK

1. Too many PoC, cut to 2 weeks max
2. Doing it with too small steps
3. Push back to k8s yaml [*]
4. Alert rules too hard to write

[*] With coaching, I thought, it is OK

DO DIFFERENT

1. Move data day one
2. Make devs know it is a transition stage
3. Teach earlier about resources
4. EFK could wait
5. All-hands for a paid-XXX% weekend for migration

SMACC



Hypatos

Problem SMACC solves

HARRY OFFICE

Company Name: HARRY OFFICE GmbH & Co. KG
Street & House Number: Alleenstraße 51
City: 12245 Köln
Country: DE
Contact Name: Herr Harry Klein
Phone: Tel. 021 12245 67
Mail: E-Mail: info@harry-office.de
Address: Internet: www.harry-office.de
Invoice Date: 16.07.2017

VAT-ID: USt-ID: DE238953422
Postal Code: 14352 Kleinmachnow
Country: DE
Document Type: Rechnung

Rechnung Nr. HC2017-07-2187
Bitte bei Zahlungen und Schriftverkehr angeben

Kunden-Nr.: 1003
Bestell-Nr.: 2376A12

Pos.	Description	Einheit	Preis	Nr. of Items	Gross Amount
1	Whiteboard Marker, 10x		15,00 EUR	3	45,00 EUR
2	Druckpapier 500 Blatt		8,00 EUR	5	40,00 EUR
3	Kugelschreiber, schwarz, bold		7,00 EUR	2	14,00 EUR
4	Kugelschreiber, schwarz, Medium		6,00 EUR	1	6,00 EUR
5	point, 1,0mm, 12x		3,00 EUR	5	15,00 EUR
6	point, 1,0mm, 12x		4,00 EUR	2	8,00 EUR
7	Ordner, Standard schwarz		8,00 EUR	2	16,00 EUR
8	Prospektstifte, 10x				
9	Permanent Marker, 12x				

Nettobetrag: 144,00 EUR
zzgl. 19 % MwSt: 27,36 EUR
Gesamtbetrag: 171,36 EUR

Bitte überweisen Sie den Rechnungsbetrag innerhalb von 30 Tagen. Zahlbar innerhalb von 7 Tagen abzüglich 3% Skonto.

Harry Office GmbH
Alleenstraße 51 12245 Köln, DE
Geschäftsführer: Herr Klein

Bank Name: KTO
IBAN: DE23 8953 4220 0001 0001 0001 0001
BIC: DEUT33HAN
BLZ: 238953422
KTO-Info: Harry Office GmbH
USt-ID: DE238953422
Telefon: 021 12245 67
E-Mail: info@harry-office.de
Internet: www.harry-office.de

Payment Conditions: KTO

VAT-ID: USt-ID: DE238953422
Tax number: 14352 Kleinmachnow



Global clients



Deloitte.



accenture

McKinsey & Company

Hypatos

SMAC©

STORY

1. Legacy on AWS and AWS ECS :/
2. Self-hosted K8S on ProfitBricks (PB)
3. ooo... K8S can use Azure Active Directory :D

STORY

4. Get to Microsoft ScaleUp
5. Welcome Azure!
6. Luckily. AKS
7. Easy migration from PB to AKS

AZURE KUBERNETES SERVICE

- Independent from IaaS
- Plug and play
- Integration with GPU
- Our OnPrem = Our OnCloud

SIMPLICITY

- az aks CLI for setting k8s - README.rst
- Terraform for everything else
- Secrets: 1Password and gopass.pw

Terraform also sets our AWS

DIFFERENCE 💀

- Two teams in Berlin and Warsaw
- Me in Warsaw

NEW EXPERIENCE

Devs do not:

- like TravisCI, Makefiles, Yaml
- feel it is too much hassle

Transition from PB to AKS was painful

SOLUTION

- make everything lighter
- c&p without modifications
- hide the k8s, remove magic
- deploy on *tag*

Similar to the [Kelsey Hightower approach](#)

Repo .travis.yml

```
language: go
go:
- '1.16'
services:
- docker
install:
- curl -sL https://${GITHUB_TOKEN}@raw.githubusercontent.com
- if [ -f "tools/travis/install.sh" ]; then bash tools/travi
script:
- dep ensure
- make lint
- make test
- if [ -z "${TRAVIS_TAG}" ]; then make snapshot; fi;
deploy:
```


Makefile

```
| - tools
|   | - Makefile
|   | - kube-service.yaml
|   \- kube-deployment.yaml
|
| - Dockerfile
\ - Makefile
```

Makefile only tasks for dev

CONTINUOUS DEPLOYMENT

- Github
- TravisCI
- hub.docker.com
- AKS

CONTINUOUS DEPLOYMENT

1. `git tag` and `push`
2. `smacc-platform.git`
3. Deploy to staging
4. PR to production

KUBERNETES

- Pure, generated, kubernetes config
- 2x kubernetes operators

WHAT WORKED

- Hiding k8s
- Understandable CD process

WOULD DO DIFFERENT

- More sensitive to feedback

NEXT

- Acceptance tests listen on k8s events
- Deployment tool based on [missy](#)
- Keeping an eye on Istio

K8S - Linux / App server?

- Out-of-box integration
- Lingua Franca - AWS Service Operator
- Learn as you go
- onPremise = onCloud = OnLocal (e.g., [kube-desktop](#))

K8S - Linux / App server?

- Do not terrorize your devs with K8S
- No free lunch... app must be smarter
- On VM vs K8S vs Lambdas?

THANK YOU. QUESTIONS?

ps. We are hiring.

```
123 def distance_matrix(regions):~
124     """ Computes a distance matrix against a region list """~
125     tuples = [r.as_tuple() for r in regions]~
126     return cdist(tuples, tuples, region_distance)~
127 ~
128 ~
129 def clusterize(words, **kwargs):~
130     # TODO: write a cool docstring here~
131     db = DBSCAN(metric="precomputed", **kwargs)~
132     X = distance_matrix([Region.from_word(w) for w in words])~
133     labels = [int(l) for l in db.fit_predict(X)]~
```

MAY THE SOURCE
BE WITH YOU.



github.com/wojciech12/talk_cloudnative_and_kubernetes_waw

```
123 def distance_matrix(regions):~
124     """ Computes a distance matrix against a region list """~
125     tuples = [r.as_tuple() for r in regions]~
126     return cdist(tuples, tuples, region_distance)~
127 ~
128 ~
129 def clusterize(words, **kwargs):~
130     # TODO: write a cool docstring here~
131     db = DBSCAN(metric="precomputed", **kwargs)~
132     X = distance_matrix([Region.from_word(w) for w in words])~
133     labels = [int(l) for l in db.fit_predict(X)]~
```

MAY THE SOURCE
BE WITH YOU.



BACKUP SLIDES

```
123 def distance_matrix(regions):~
124     """ Computes a distance matrix against a region list """~
125     tuples = [r.as_tuple() for r in regions]~
126     return cdist(tuples, tuples, region_distance)~
127 ~
128 ~
129 def clusterize(words, **kwargs):~
130     # TODO: write a cool docstring here~
131     db = DBSCAN(metric="precomputed", **kwargs)~
132     X = distance_matrix([Region.from_word(w) for w in words])~
133     labels = [int(l) for l in db.fit_predict(X)]~
```

MAY THE SOURCE
BE WITH YOU.



CHANGE THE WORK ORGANIZATION


- From Scrum
- To Kanban
- from Kanban to Squads

For the next talk

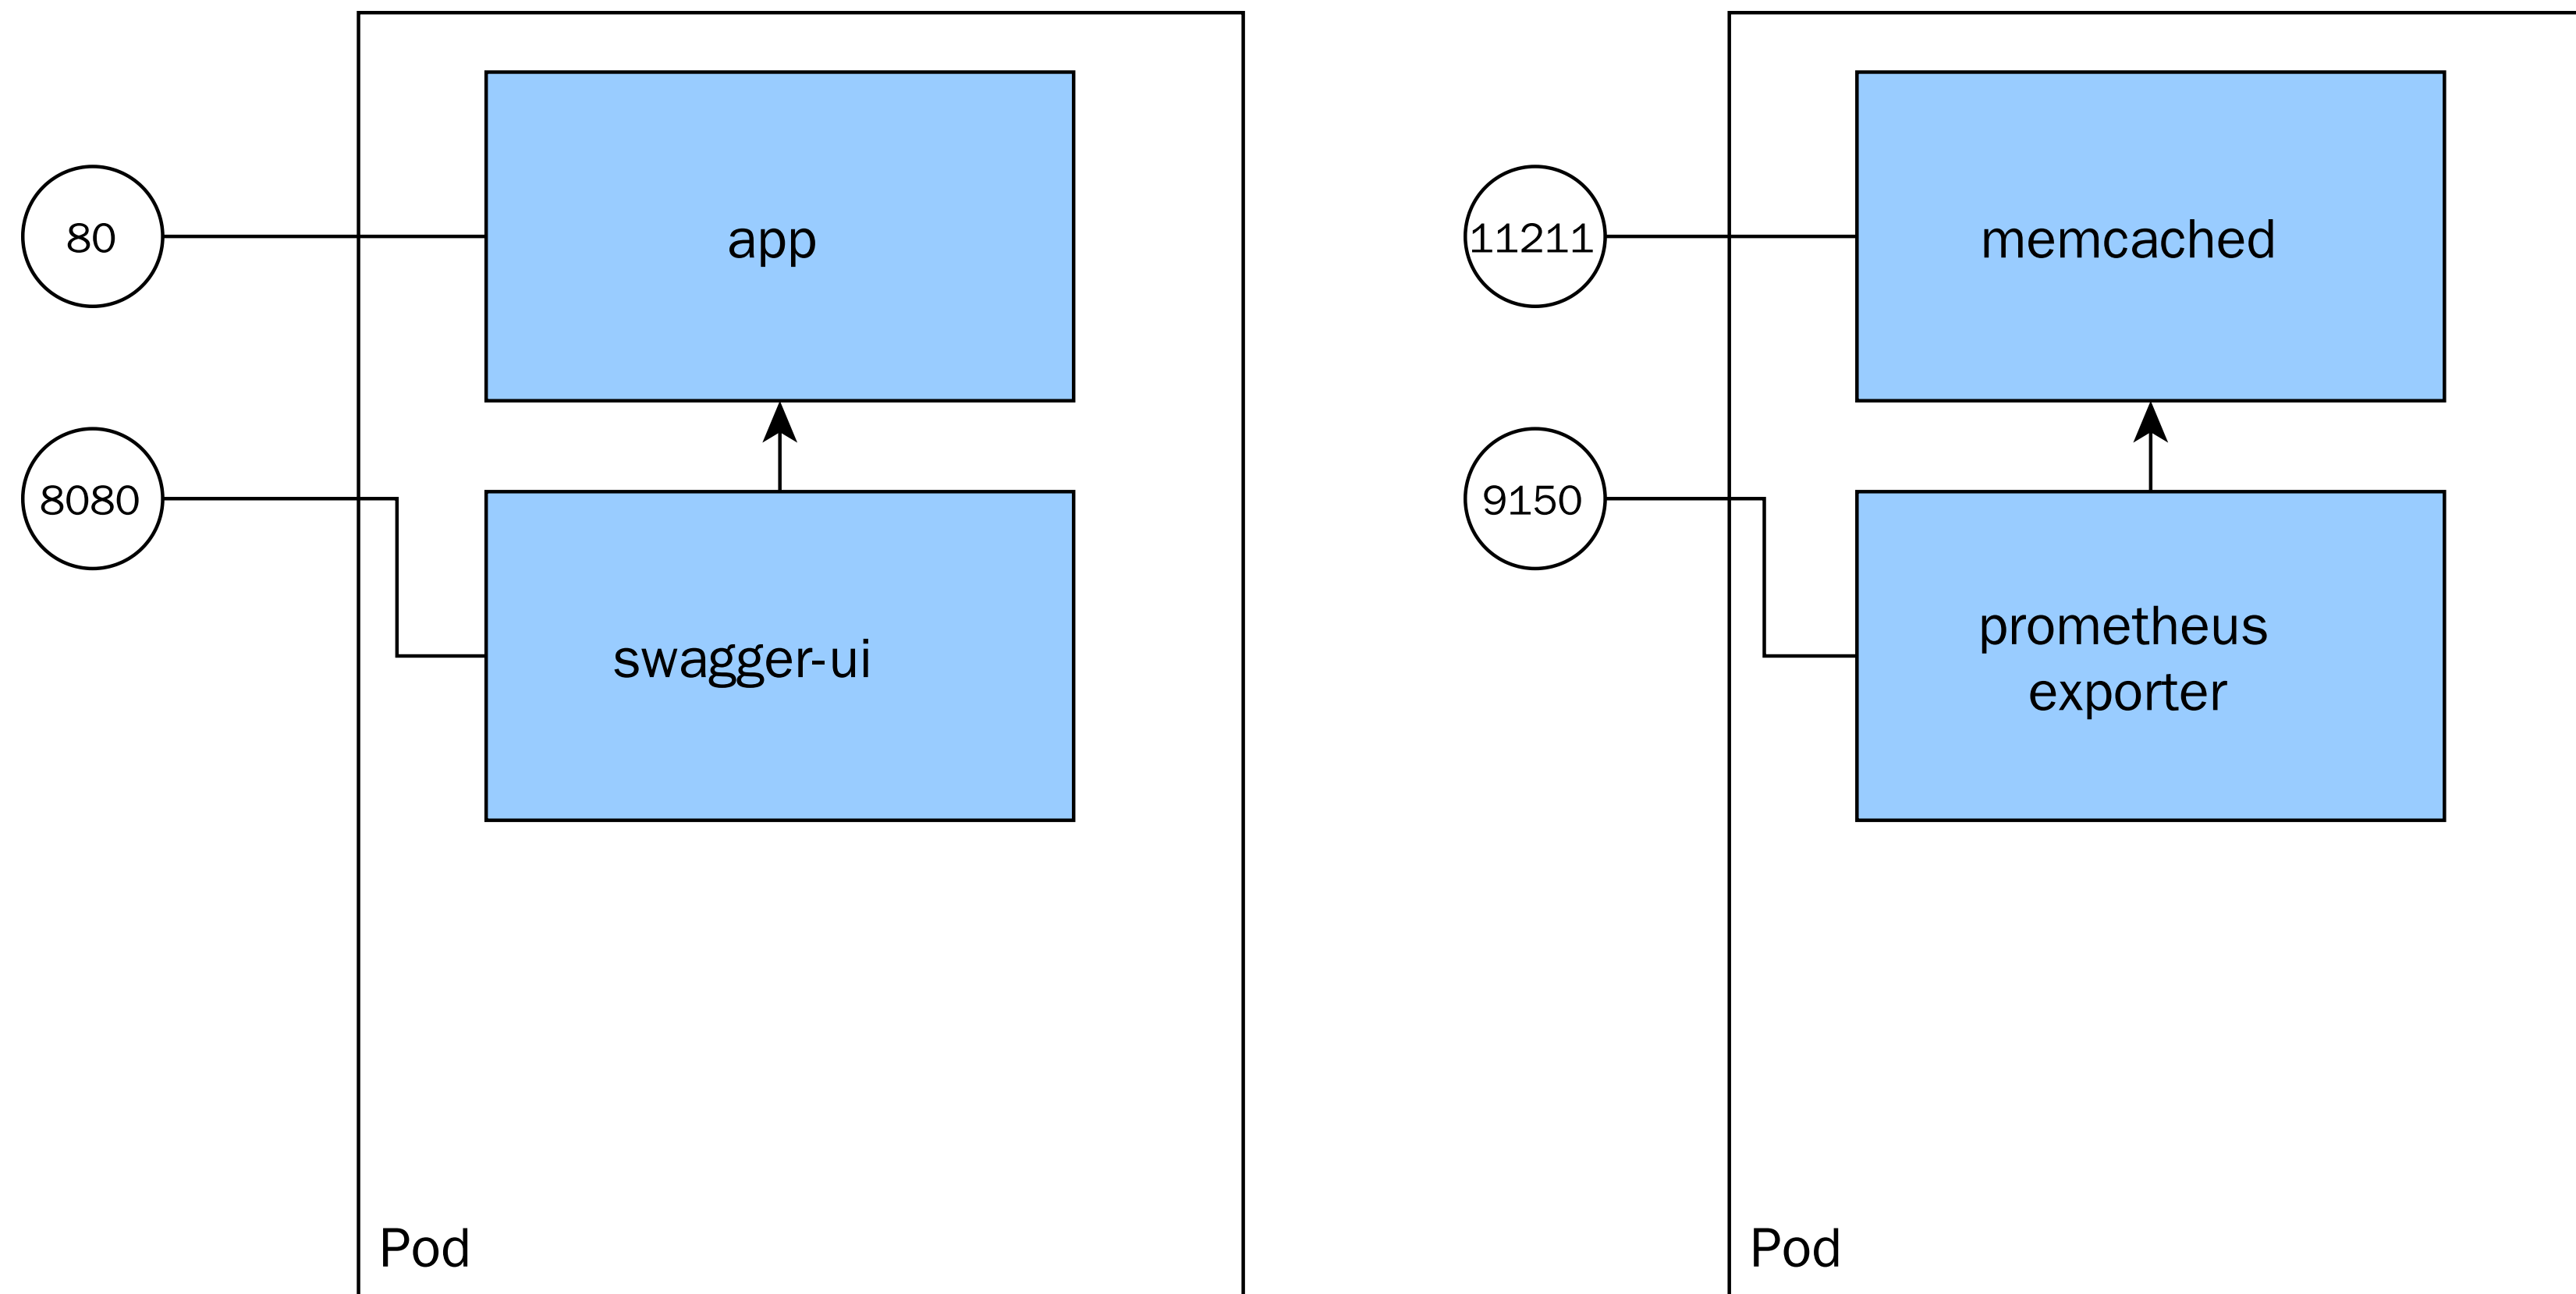
KUBERNETES CONCEPTS

PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



SIDE-CARS



BASIC CONCEPTS

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running